Zero-parameter structural computation of Standard Model observables and constants and a dark-matter tower from $E_8 \times E_8$ Coxeter geometry on a Schoen Calabi–Yau threefold

Eliot Gable*
Independent Researcher
(Dated: May 7, 2026)

What this paper does. We construct the Standard Model fermion mass spectrum, gauge couplings, mixing angles (CKM, PMNS), CP phases, dark-matter mass tower, Higgs boson mass, Higgs vacuum expectation value, and Newton's gravitational constant from a small set of fundamental equations acting on the fixed geometry of $E_8 \times E_8$ heterotic compactification on the Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ Calabi–Yau threefold quotient $X/(\mathbb{Z}_3 \times \mathbb{Z}_3)$, with the McKay correspondence on the binary icosahedral group $2I$ and Wilson $\mathbb{Z}_3 \times \mathbb{Z}_3$ symmetry breaking. The construction has no PDG inputs and no adjustable quantities.

Why it matters. Approximately twenty-six SM observables are predicted to ppm-or-better precision: m_e at 17 ppm, m_μ at 30 ppb, m_τ at 0.15 ppm, $\sin^2\theta_W = 2/9 + \Delta_\varphi/h$ at 198 ppm, $\alpha_s(M_Z) = 849/7192$ at 405 ppm, the dark-matter abundance $\Omega_{\rm DM}h^2 = 0.119982$ vs. Planck 2018 0.1200(10), and Newton's G to 2.5 ppm vs. CODATA 2022 — all from the geometry alone. The CP phases $\delta_{\rm CP} = \pi + \pi/10$ and $\delta_{\rm CKM} = 11\pi/30$ are predicted as rational multiples of π. Normal neutrino mass ordering is selected. A conditional geometric route to vanishing strong CP via the $\mathbb{Z}_3 \times \mathbb{Z}_3$ quotient projection avoids invoking a Peccei–Quinn axion.

What is open and not claimed. The construction is a strong-selection structural identification, not yet a forced derivation: the substrate-physical projection onto the heterotic moduli of $X/(\mathbb{Z}_3 \times \mathbb{Z}_3)$ is not closed in this paper. The σ-discrimination Schoen-vs-Tian–Yau result (6.92σ Tier 0; bundle-construction-specific) is reported as supporting evidence pending the converged PDG-fermion-mass Yukawa-cohomology pipeline. Dark-sector chirality assignments and several closure-rule numerators are encoded structural-rule-consistent choices whose generative derivation is deferred.

CONTENTS

I. Introduction 1

II. The substrate framework 3
- A. Heterotic $E_8 \times E_8$ on the Schoen quotient 3
- B. Empirical CY3 substrate identification 5
- C. The E_8 Coxeter spectrum 7
- D. The McKay correspondence 8
- E. Substrate-tension interpretation 9

III. Gateway formula 9

IV. Charged-fermion masses 10
- A. Mass formula schema 10
- B. Lepton sector 11
- C. Quark sector 11
- D. Term-by-term decomposition: charged fermions 11

V. Higgs sector and dark-sector coupling family 12
- A. Boson mass-formula schema and structural corrections 12
- B. Origin of the dark-sector coupling family 14
- C. The Higgs leak: structural closure rule for δ_H 14
- D. The W leak: structural closure rule for δ_W 15
- E. Generalization to other cycle-2 modes 15
- F. Term-by-term decomposition: electroweak bosons 15

VI. Electroweak gauge sector 16
- A. Weak mixing angle 16
- B. Gauge couplings and Z boson 16
- C. Higgs self-coupling 17
- D. Strong coupling 17
- E. Fine-structure constant 17

VII. Strong CP angle 17

VIII. Neutrino sector 17
- A. Mass eigenvalues from sub-gateway φ-chain modes 17
- B. Mass-squared splittings and bounds 18
- C. PMNS mixing angles and CP phase 18

IX. Quark mixing (CKM) 18

X. Dark-matter tower 19
- A. The unified neutrino-DM Coxeter tower 19
- B. Cosmological cutoffs and viable mass window 19
- C. Order-5 dominant mode: couplings and detection prospects 19
- D. Geometric-series abundance convergence 19
- E. Decay constant and misalignment angle 19
- F. Rank-locked sub-harmonic correction 19
- G. Term-by-term decomposition: dark-matter tower modes 20
- H. Cosmological abundance result 21

* 10015996+egable@users.noreply.github.com

XI. Gravitational sector 21

XII. Two-anchor diagnostic consistency check 21

XIII. Results 21

XIV. Predictions and falsification 25

XV. Open items and further-exploration program 25
A. Pipeline state and in-flight remediation 25
B. Open methodological items 27
C. Open framework items deferred to companion paper 27
D. What this paper does not claim 28
E. Further-exploration program 28

XVI. Discussion 29
A. Parameter-count semantics 29
B. Common geometric origin 29
C. Comparison to prior work 29
D. Open directions 29
E. Conclusion 29

XVII. Computed 4D effective Lagrangian: framework prediction with verdict (a) falsification 30
A. Canonical form and provenance chain 30
B. Lagrangian header summary 30
C. Newton's G from heterotic volume reduction 30
D. Higgs potential as derived corollary 31
E. Yukawa sector: framework prediction at $\kappa = 0$ fixed point, falsified vs PDG 31
F. Cosmological constant: tree-level Higgs $V_{\min}$ only 32
G. Milestone interpretation: complete prediction set, end-to-end falsifiable 32

XVIII. Channel-grounding map and substrate-native derivation: where the mass equations live algebraically 32
A. The 89-row grounding channel matrix 33
B. The 3-piece unified generating principle and the $\dim -2 = 246$ substrate identity 33
C. Substrate-native derivation of the mass equations from coupling-reach mechanics 34
D. Per-ingredient $E_8 \times E_8$ forcing audit 35
E. SI from substrate: absolute-scale closure of h_{eff} 36
F. Two-layered geometry and the Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ commitment 37
G. Hostile-review corrections applied 38
H. Honest remaining gaps 38
I. Honest publication-safe scope statement 39

A. Coxeter step assignments 39

B. Complete structural correction list 39

C. Numerical precision 39

References 40

I. INTRODUCTION

The Standard Model (SM) of particle physics, despite its empirical success, contains twenty-six parameters whose values are not predicted by any underlying principle but determined empirically from observation [1]. These include the three gauge couplings, the Higgs sector parameters (vacuum expectation value and self-coupling), nine charged-fermion Yukawa couplings, three neutrino mass scales, four quark-sector mixing parameters (CKM), four lepton-sector mixing parameters (PMNS), and the QCD vacuum angle θ_{QCD}. Beyond the SM, the existence and mass scale of dark matter constitutes one of the most prominent unresolved puzzles in contemporary physics [2].

This paper presents a structural computation, with no fitted continuous parameter, of all twenty-six SM parameters, Newton's gravitational constant, the Planck mass, and a discrete dark-matter mass tower with associated decay constant, couplings, and cosmological abundance, given a fixed set of first-principles starting assumptions (Table I). Per-particle Coxeter chain indices, sign assignments, sector and cycle placements, and rational closure rules arise as structural outputs of those starting commitments rather than as independently varied parameters; the audit is in Sec. I 0 b. PDG-matching residuals are predictions versus measurement, not fits. The framework's predictive path takes *zero* measured masses as input. The Higgs vacuum expectation value is computed in closed form from the E_8 Coxeter invariants alone via the structural identity $v_C = (\dim -2)/(1 - \delta_W) \cdot (1 + 5\,\alpha_\tau^3/8)$ (see substrate_mass_spectrum.py:402; inputs $\dim = 248$ at line 160, $h = 30$ at line 161, $m_{\max} = 29$ at line 163, $e_6 = 19$ at line 399, and the derived $\delta_W = 1/((h + m_{\max})e_6) = 1/1121$ at line 400 and $\alpha_\tau = 1/(\dim + h) = 1/278$ at line 188), giving $v_C = 246.21965$ GeV. The gateway tension scale is then $h_{\text{eff}} = v_C/R_v$ where R_v is the structural Higgs-vev ratio, and every other mass is the corresponding chain-specific structural ratio times h_{eff} (*φ-chain leptons, $\sqrt{2}$-chain quarks/EW bosons, α_τ-tower neutrinos and dark-matter modes*). As an internal-arithmetic self-consistency check, alternatively anchoring the framework on the CODATA electron mass alone or muon mass alone produces predictions that agree at 0.0298 ppm ($\sim$ 30 ppb), set by sub-harmonic-series truncation; both fermion-anchored modes agree with the structural-vev anchor to within 0.12 ppm on every observable. This two-anchor *diagnostic* validates the framework's internal arithmetic but is not the predictive path.

a. Starting assumptions of this paper. Table I enumerates the complete list of first-principles commitments on which the numerical demonstration of this paper rests. The list contains no continuous fitted parameter. Every structural integer, sign, sub-harmonic factor, and closure rule that appears in the mass equations of Secs. III–XI is either an immediate consequence of these commitments or, where so flagged, an independent structural

rule whose derivation is the subject of the companion paper [43].

b. Structural-output audit. The construction generates discrete data — per-particle Coxeter chain indices, sign assignments, sector and cycle placements, and a family of closure rules — as *structural outputs* of the starting commitments of Table I, not as independently varied parameters. They are not free: changing any one of them in isolation would generally be inconsistent with the underlying structural rules (the binary-icosahedral McKay seed, the φ- and $\sqrt{2}$-chain harmonic-recursion split, the boundary capacity $h + m_{\max} = 59$, the E_8 Coxeter spectrum, and the Wilson-line-broken visible chain). They become arbitrary only under the counterfactual in which the structure is removed. We list the discrete data explicitly so the reader can audit by inspection which items are derived in closed form within this paper and which are encoded as structural-rule-consistent choices whose generative derivation is part of the open program.

- *Per-particle Coxeter chain index* k_X for each charged fermion, EW boson, and the Higgs vev (App. A): $k_e = 0$, $k_\mu = 11$, $k_\tau = 17$ on the φ-chain (lepton positions at integer E_8 exponents); $k_u = 4$, $k_d = 6.5$, $k_s = 15$, $k_c = 22.5$, $k_b = 26$, $k_t = 36.5$, $k_W = 34.5$, $k_v = 1132/30$, $k_H = k_v - 2$ on the $\sqrt{2}$-chain. The Higgs-vev position $k_v = h + \mathrm{rank} - \mathrm{rank}/h$ is given in closed form (Eq. 21); the remaining k_X values are encoded as positions consistent with the integer/half-integer cycle-1 and cycle-2 structure of the chain. A generative derivation from the specific Wilson-line representation on X/Γ is part of the open program (Sec. XV).

- *Per-quark sign assignments* $\chi_q, \lambda_q \in \{-1, 0, +1\}$ in the quark mass formula (Table IX): encoded as a structural pattern across the six quarks; their derivation from a structural-rule generator (chirality flip across the harmonic threshold $h/2$, per-step-loss sign across cycle boundaries) is part of the open program.

- *Sector chain-type assignment* (φ vs $\sqrt{2}$): leptons and neutrinos to the φ-chain; quarks and electroweak bosons to the $\sqrt{2}$-chain. This pairing reflects the framework's McKay-icosahedral and McKay-octahedral seed assignments; an explicit derivation from a Wilson-line representation on X/Γ is open.

- *Cycle-1 vs cycle-2 placement* for each particle: encoded in the structural pattern of the chain (App. A).

- *Closure rules* $\delta_H = 1/(h + m_{\max}) = 1/59$ and $\delta_W = 1/(59 \cdot e_6) = 1/(59 \cdot 19) = 1/1121$ are given in closed form from the boundary capacity $h + m_{\max}$ and the E_8 exponent set; the gateway sub-harmonic factor $1/(4(\dim - 1))$ admits a dual-sector / two-orientation structural reading whose derivation from a heterotic-anomaly index or a $\mathbb{Z}_3 \times \mathbb{Z}_3$ quotient trace on the Schoen threefold is part of the open program.

- *Sub-harmonic correction structures* (α_v, α_H, α_W, α_t, the rank-locked sub-harmonic on m_5, and the $5\varphi^2\alpha_\tau$ correction on ν_3) are encoded as structural-rule combinations of $\{\dim, h, \mathrm{rank}, m_{\max}, e_j, \varphi\}$. A generative derivation of each from the compactification data is open.

- *Other rational closures* include $\sin^2\theta_W = 2/9 + \Delta_\varphi/h$ and $\alpha_s(M_Z) = (m_{\max}^2 + \mathrm{rank})/(m_{\max} \cdot \dim) = 849/7192$. The leading $2/9$ in $\sin^2\theta_W$ and the combination $m_{\max}^2 + \mathrm{rank}$ in α_s are framework structural choices; they are not standard E_8 Lie invariants and their derivation is open.

None of the items above is varied independently of the framework's structural rules. Items derived in closed form within this paper (boxed equations) are structurally forced; items flagged as encoded or as "framework structural assignment" are consistent with the framework's rules but their generative derivation from the compactification data is the subject of the further-exploration program (Sec. XV).

c. Epistemic posture. From a 4D observer's vantage the starting assumptions of Table I cannot be *proven*: any purported derivation of the substrate-physical content from inside the framework's macroscopic 4D regime is necessarily a consistency check, not a proof. The standard a 4D observer can apply is therefore *empirical consistency* — the assumptions are evaluated by the predictions they generate. This paper performs that evaluation by computing the SM observables and comparing to measurement. We make no claim in this paper that the starting assumptions are forced from first principles; we make the weaker, sharper claim that, taking them as given, the SM mass spectrum and mixing parameters compute directly with the residuals reported in Table I's Sec. XIII downstream. The companion paper [43] addresses the separate question of whether the starting assumptions are minimal and substrate-physically expected; that argument is independent of the present demonstration and not relied upon here.

d. Roadmap. Section II fixes the framework as commitment — heterotic structure on the Schoen quotient, E_8 Coxeter spectrum, McKay-icosahedral seed, sub-Coxeter decomposition, substrate-tension Higgs reading. Section II B reports the empirical Schoen-versus-Tian–Yau discrimination, the Tian–Yau Yukawa-channel structural exclusion, and the catalogue uniqueness sweep at Hodge $(3, 3)$, with explicit hedges for the open methodological items. Section III presents the gateway formula and the direct computation of the electron mass. Sections IV–XI present the direct computations of charged-fermion masses, the Higgs sector, electroweak gauge couplings, the strong-CP angle, neutrino masses and mixing, CKM mixing, the dark-matter Coxeter tower with

TABLE I. Starting assumptions of the present paper. The numerical demonstration in Secs. III–XI takes these as given. The argument that this list is minimal — i.e. that one would expect the substrate-physical principles to select these and not others — is the subject of the companion paper [43] and is not load-bearing for the present demonstration.

Starting assumption	Status / role in the present paper
Heterotic $E_8 \times E_8$ string theory [11]	Standard heterotic literature; assumed.
Schoen Calabi–Yau threefold quotient $X/(\mathbb{Z}_3 \times \mathbb{Z}_3)$ [10, 12, 13]	Empirically identified within the framework's narrowed candidate set; see Sec. II B.
E_8 Coxeter spectrum (rank 8, $h = 30$, exponents $\{1, 7, 11, 13, 17, 19, 23, 29\}$, $\lvert\Phi\rvert = 240$)	Standard Lie-algebraic data [6]; assumed.
McKay correspondence [5]; binary icosahedral McKay-dual of E_8	Standard; the icosahedral order-5 assignment as the framework's harmonic seed is a structural commitment. Companion paper argues for its substrate-physical inevitability.
Wilson-line breaking $E_8 \to G_{\rm SM}$ along $\pi_1(X/\Gamma) = \mathbb{Z}_3 \times \mathbb{Z}_3$; phenomenological lepton/quark sector split via φ-chain (5-fold icosahedral) and $\sqrt{2}$-chain (4-fold octahedral) harmonic recursions	Standard heterotic chain $E_8 \supset E_6 \supset SO(10) \supset SU(5) \supset G_{\rm SM}$. The $\varphi/\sqrt{2}$ chain split is a framework structural assignment matching observed mass hierarchies with small-integer Coxeter-step assignments; a derivation from a specific Wilson-line representation on X/Γ is part of the open program (Sec. XV).
Substrate-tension reading of the Higgs mechanism (motivated by Anderson's superconducting precursor [7]; similar in spirit to but independent of Volovik's emergent-vacuum program [8])	Interpretive frame for the Higgs vacuum; the v ladder in Sec. V is a closed-form structural identity within this reading.
Coldea *et al.* [3] confirmation of the Zamolodchikov E_8 spectrum [4] in $CoNb_2O_6$	Empirical observation of the E_8 root system in a 1+1D condensed-matter context; adopted as motivation for the framework's E_8 commitment, not as a direct measurement of the heterotic gauge E_8. Not load-bearing for the numerical computations of the present paper.

cosmological abundance, and Newton's G. Section XIII presents the master results tables. Section XIV enumerates the falsifiable predictions. Section XV enumerates the open methodological items, the framework-level commitments deferred to the companion paper, and the further-exploration program. Section XVI discusses parameter-count semantics and relation to prior work.

II. THE SUBSTRATE FRAMEWORK

A. Heterotic $E_8 \times E_8$ on the Schoen quotient

The framework treats heterotic $E_8 \times E_8$ string theory [11] compactified on the Schoen Calabi–Yau threefold quotient $X/(\mathbb{Z}_3 \times \mathbb{Z}_3)$ as the structural substrate. The Schoen threefold [10] is the smooth, simply-connected Calabi–Yau fiber product $X = B \times_{\mathbb{P}^1} B'$ of two rational elliptic surfaces over $\mathbb{P}^1$, with Hodge numbers $h^{1,1}(X) = h^{2,1}(X) = 19$ and Euler characteristic $\chi(X) = 0$. In the heterotic standard-model construction, suitably chosen B, B' admit a freely acting $\Gamma = \mathbb{Z}_3 \times \mathbb{Z}_3$ symmetry, and the quotient X/Γ acquires $\pi_1(X/\Gamma) = \Gamma = \mathbb{Z}_3 \times \mathbb{Z}_3$ (while X itself remains simply-connected), enabling Wilson-line breaking of one E_8 factor to Standard-Model gauge structure [12, 13].

a. Hodge numbers and generation count. The Schoen quotient X/Γ has Hodge numbers $h^{1,1}(X/\Gamma) = h^{2,1}(X/\Gamma) = 3$, computed as the dimensions of the Γ-invariant subspaces of $H^{1,1}(X)$ and $H^{2,1}(X)$ via the character formula $\dim H^{p,q}(X)^\Gamma = (1/|\Gamma|) \sum_{g\in\Gamma} \mathrm{tr}(g^* | H^{p,q}(X))$; see [12] for the explicit representation-theoretic decomposition. The Euler characteristic of the cover, $\chi(X) = 2[h^{1,1}(X) - h^{2,1}(X)] = 0$, divides under the freely-acting Γ to give $\chi(X/\Gamma) = 0$ in absolute Euler-characteristic terms; the *net chirality index* relevant for the standard heterotic generation count is set instead by the chiral Atiyah–Singer index of the bundle V chosen to break E_8 via Wilson lines, $|n_g - \bar{n}_g| = |\mathrm{Index}(V)|/|\Gamma|$ with $\mathrm{Index}(V) = \int_X \mathrm{ch}(V) \cdot \mathrm{Td}(TX)$ (equivalently $(1/2)\int_X c_3(V)$ when $c_1(V) = 0$, by the standard Hirzebruch–Riemann–Roch reduction on a CY3). For the BHOP rank-4 $SU(4)$ extension bundle on Schoen, $\mathrm{Index}(V_{\rm BHOP}) = -27$ (BHOP-2005 Eq. 88; with $c_3/2 = -9\,\tau_1\tau_2^2$ per BHOP-2005 Eq. 97 and $\int_X \tau_1\tau_2^2 = 3$, giving Index $= -27$), and the standard Braun–He–Ovrut–Pantev model thus yields $|n_g - \bar{n}_g| = 27/9 = 3$ chiral generations as required by the visible Standard Model. The present framework assumes (rather than derives) the bundle commitment, and adopts the BHOP-2005 published index numerics.

b. Heterotic anomaly cancellation. Heterotic $E_8 \times E_8$ on a Calabi–Yau threefold requires the modified Bianchi identity $dH = (\alpha'/4)\big(\mathrm{tr}\, R \wedge R - \mathrm{tr}\, F \wedge F\big)$ to admit a globally defined H, equivalent to the cohomological condition

$$\mathrm{ch}_2(V) \;=\; \mathrm{ch}_2(TX) \;+\; [W], \tag{1}$$

where ch_2 is the second Chern character, V is the visible-sector gauge bundle, TX is the holomorphic tangent bundle of the Schoen threefold, and $[W]$ is an effective 5-brane class in $H^4(X,\mathbb{Z})$ (vanishing in the standard embedding $V = TX$). The present framework assumes that the bundle choice underlying the three-generation count is consistent with Eq. (1) for some admissible 5-brane class $[W]$, as is the case in the Braun–He–Ovrut–Pantev construction [12]. A first-principles identification of the specific $[W]$ used by the framework, together with the *slope-stability* (μ-stability) of V required by the Donaldson–Uhlenbeck–Yau theorem [27, 28] to admit a Hermitian–Yang–Mills connection on the quotient X/Γ [13], is part of the open program (Sec. XV). The Donaldson-balanced metric pipeline used in Sec. II B is constructed assuming polystability of V; the empirical convergence of the Donaldson algorithm on the candidate bundles is evidence-for, but not proof-of, polystability on the quotient.

c. Substrate dimensionality and the 10D embedding. The discrete substrate update rule developed in the companion paper [43] is specified directly in $3+1$ macroscopic dimensions, while the heterotic $E_8 \times E_8$ string lives in $9+1$. The framework's operational stance is that the six internal dimensions are *algebraically encoded* in the substrate by the $E_8 \times E_8$ structure of the chain construction (in particular, the Coxeter-chain positions k_X and the sector-modifier $\sqrt{2}$- vs. φ-chain assignments substitute for direct propagation in the compactified directions): the substrate is treated as $3+1$D throughout, with the Calabi–Yau / heterotic structure entering only through the algebra controlling the mass spectrum and gauge structure. This is a *framework choice*, not a derivation: a substrate-physical realisation in which the $3+1$D macroscopic substrate is a boundary or low-energy slice of an underlying $9+1$D substrate (analogous in spirit to brane-world or warped-compactification constructions) is consistent with the present construction but not committed to here. In particular, no claim is made about the size, geometry, or stabilisation of the six internal directions beyond what is needed to identify the algebraic content $E_8 \times E_8$ on X/Γ.

One E_8 factor contains the visible Standard Model gauge structure after compactification. The second E_8 factor is the dark sector, communicating with the visible sector only gravitationally except through identified portal couplings.

d. Visible $E_8 \to$ SM symmetry-breaking chain. The visible E_8 is broken to the Standard-Model gauge group $G_{\mathrm{SM}} = SU(3)_{\mathrm{color}} \times SU(2)_L \times U(1)_Y$ through the BHOP $SU(4)$-structure-group chain

$$\begin{aligned} E_8 \;&\supset\; SO(10)\times SU(4) \;\supset\; SU(5)\times U(1) \\ &\supset\; G_{\mathrm{SM}} \times U(1)_X, \end{aligned} \tag{2}$$

with the final descent realised by Wilson lines along generators of $\pi_1(X/\Gamma) = \mathbb{Z}_3 \times \mathbb{Z}_3$. The bundle V committed here is the rank-4 $SU(4)$ extension construction of Braun–He–Ovrut–Pantev [14] §6 (an extension of two semi-stable summands rather than a line-bundle sum), with chiral index $\mathrm{Index}(V_{\mathrm{BHOP}}) = \int_{\widetilde{X}} \mathrm{ch}(V_{\mathrm{BHOP}}) \cdot \mathrm{Td}(T\widetilde{X}) = -27$ on the cover (BHOP-2005 Eq. 88; $c_3/2 = -9\,\tau_1\tau_2^2$ per BHOP-2005 Eq. 97), yielding three generations on the free $\Gamma = \mathbb{Z}_3 \times \mathbb{Z}_3$ quotient via the Atiyah–Singer / Hirzebruch–Riemann–Roch reduction $N_{\mathrm{gen}} = |\mathrm{Index}(V_{\mathrm{BHOP}})|/|\Gamma| = 27/9 = 3$. The full heterotic Bianchi anomaly cancellation $c_2(TX) - c_2(V_{\mathrm{BHOP}}) - c_2(V'_{\mathrm{BHOP}}) = 6\,\tau_1^2$ (BHOP Eq. 99) is satisfied by a positive-tension 5-brane wrapping τ_1^2, and Mumford–Takemoto polystability is established on the documented Kähler region $\{t_a \in [0.25, 8.0],\ t_a/t_b \in [0.25, 4.0]\}$. Both checks are programmatically verified upstream of this paper at commits `a3317377` (anomaly cancellation) and `a4d1bc72` (polystability). The matter content is delivered by $H^1(X/\Gamma, V_{\mathrm{BHOP}})$ decomposing into three copies of the $\mathbf{16}$ of $SO(10)$ under the $\mathbb{Z}_3 \times \mathbb{Z}_3$ Wilson-line characters, with right-handed neutrinos arising automatically as the singlet $\mathbf{1}$ in $\mathbf{16} = \mathbf{10} + \bar{\mathbf{5}} + \mathbf{1}$. This commitment supersedes the wave-28-era rank-3 $SU(3)$ "non-standard" bundle placeholder, which was retracted in the audit at `p_stage3_bundle_physics_audit.md` (commit `0cad4b0b`) for failing the full Bianchi identity, the $N_{\mathrm{gen}} = 3$ Atiyah–Singer count (the rank-3 placeholder did not reproduce the BHOP-2005 Eq. 88 index of -27), polystability wiring, and the 3:3:3 Wilson-character decomposition of H^1. Alternative chains (the Donagi–Ovrut $SU(5)$-structure-group route, or $E_8 \supset SO(16)$) admit consistent compactifications but are not the framework's commitment. The framework's mass formulas use E_8 Coxeter-spectrum data and Wilson-line-broken chain assignments that are largely independent of which standard chain is selected; the per-particle Coxeter-step assignments listed in App. A are framework structural commitments that the BHOP $SU(4)$ identification anchors to a specific heterotic bundle.

e. Substrate-physical reading of the rank-4 internal symmetry (I, suggestive; not a BHOP-2005 commitment). The rank-4 $SU(4)$ here is the *bundle structure group* of V_{BHOP} acting on holomorphic sections of V on the cover, *not* the gauged $SU(4)_C$ of Pati–Salam 1974 [15] acting on Standard-Model fermion 4-tuples. The framework offers a *suggestive* substrate-physical reading in the spirit of Pati–Salam $SU(4)_C$ colour-lepton unification — the four internal modes of the bundle treated as three colour modes plus a fourth "lepton-number" mode — because both groups are $SU(4)$ and both relate to lepton/quark unification. However, the identification of the bundle structure group with a 4D gauged $SU(4)_C$ is a non-trivial *additional* structural step that BHOP-2005 §6 does not assume: BHOP's $SU(4)_{\mathrm{bundle}}$ acts on holomorphic sections, not on Standard-Model fermion fields, and the GUT chain $E_8 \to SO(10) \times SU(4)_{\mathrm{bundle}} \to SU(5) \times U(1) \to G_{\mathrm{SM}} \times U(1)_X$ has no Pati–Salam stage in the gauge-symmetry-breaking ladder. The substrate-physical mechanism realising any bundle→gauge identification is currently an open structural step. The framework's empirical content

is independent of the identification: chiral generations descend from the $\mathbb{Z}_3 \times \mathbb{Z}_3$ Wilson characters acting on $H^1(X/\Gamma, V_{\mathrm{BHOP}})$ regardless of the Pati–Salam reading, and right-handed neutrinos are automatic from the **16** of $SO(10)$ (the singlet **1** in $\mathbf{16} = \mathbf{10} + \bar{\mathbf{5}} + \mathbf{1}$). The rank-4 commitment is preferred over the wave-28 rank-3 $SU(3)$ placeholder because the rank-4 extension bundle satisfies the BHOP-2005 Eq. 88 index of -27, full Bianchi anomaly closure, polystability, and the **16**-decomposition required for right-handed neutrinos — not because the Pati–Salam reading is structurally forced. The residual $U(1)_X$ surviving Wilson-line breaking is anomalous and is broken by the standard heterotic Green–Schwarz mechanism (anomalous-$U(1)$ Stueckelberg coupling to the dilaton/ B-field universal axion); the substrate-physical mechanism for this GS cancellation is an open step.

B. Empirical CY3 substrate identification

The framework's substrate-physical forcing chain selects the Hodge $(3,3)$, $\pi_1 = \mathbb{Z}_3 \times \mathbb{Z}_3$ branch on which Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ sits. The structurally-closest distinct three-generation Calabi–Yau threefold satisfying $c_1 = 0$ and admitting a free finite-group action with Standard-Model-compatible Wilson breaking is the Tian–Yau quotient TY/$\mathbb{Z}_3$ (Tian–Yau 1987 [16], Hodge $(h^{1,1}, h^{2,1}) = (6,9)$, $\pi_1 = \mathbb{Z}_3$, $\chi = -6$). Tab. II records the structural distinction.

To convert this from a literature-driven choice into an empirical identification, we ran two independent physics-channel discriminators on the Schoen and Tian–Yau Donaldson-balanced metrics: a quark-chain channel and a lepton-chain channel. The pipeline as currently implemented computes, for each candidate, a Galerkin-residual sweep at successive test-degrees $k_{\mathrm{test}} \in \{2,3,4,5\}$ of the discrepancy between the metric-Laplacian eigenvalues (projected onto the framework's $\sqrt{2}$-chain integer/half-integer positions for the quark channel and the φ-chain integer-exponent positions for the lepton channel) and the structural sub-Coxeter chain expected from the framework. The signed channel residual difference $\delta = \rho_{\mathrm{Schoen}} - \rho_{\mathrm{TY}}$ is interpreted as a per-channel $\ln \mathrm{BF}_{\mathrm{TY/Schoen}} = \delta/2$ under a Laplace/BIC heuristic. The combined channel report is $\ln \mathrm{BF}_{\mathrm{TY/Schoen}} = -14.76$ (Tab. III), corresponding to the highest tier of the Kass–Raftery scale [24] ("very strong" evidence at $|\ln \mathrm{BF}| > 5$; equivalently $\log_{10} \mathrm{BF} > 6$, well beyond Jeffreys's "decisive" threshold of $\log_{10} \mathrm{BF} > 2$ [26]), with posterior odds Schoen:TY $= 2.6 \times 10^6$. The Wilks/Cowan–Cranmer–Gross–Vitells χ^2_1 equivalence $n_\sigma = \sqrt{2|\ln \mathrm{BF}|}$ yields 5.43σ-equivalent discrimination as an order-of-magnitude non-nested heuristic.

a. Convergence status of the chain-channel sweep. The chain-channel discriminators reported here are *indicative, not converged* at the strict 10% basis-size convergence criterion. The current production sweep at $n_{\mathrm{pts}} = 25{,}000$, $k = 3$, $k_{\mathrm{test}} \in \{2,3,4,5\}$ records the channel JSONs at `p7_11_quark_chain.json` and `p7_11_lepton_chain.json`, both of which carry `converged: false` flags with consecutive-degree relative deltas of order 11–35% on the Schoen side and the signed-delta δ residual. The $\ln \mathrm{BF} = -14.76$ figure is therefore a preliminary discriminator, not a converged Bayes factor in the strict statistical sense; we report it as supporting evidence alongside the catalogue uniqueness sweep (which *is* closed modulo the gCICY non-manifest residue) and the Tian–Yau Yukawa structural exclusion (which *is* closed at the searched bidegree range). The full replacement of the metric-Laplacian chain-residual likelihood by a converged PDG-fermion-mass Yukawa-cohomology pipeline (Anderson–Karigiannis–Larfors–Lukas–Schneider style: HYM Hermitian metric on V, harmonic representatives of $H^1(X/\Gamma, V)$, Yukawa cohomology pairing, Kähler normalisation, two-loop SM RG running) is the in-flight remediation specified in the project's `REMEDIATION_PLAN_OPTION_B.md`, Phase 2; we expect to retract or upgrade the present chain-channel BF figure on its completion. The catalogue uniqueness and TY-exclusion arguments do not depend on the chain-channel convergence and stand on their own.

b. σ-channel as discriminability, not BF input. The L^2 Donaldson §3.4 weighted-variance Monge–Ampère residual σ on the Donaldson-balanced metric (Donaldson 2009 [25], PAMQ **5**, p. 587, eq. 3.4: $\sigma^2 = \langle(\eta/\kappa - 1)^2\rangle_w$ with $\kappa = \langle\eta\rangle_w$, matching the production implementation in `schoen_metric.rs::compute_sigma` and `ty_metric.rs::compute_sigma`) is a separately-computed statistical discriminator that distinguishes the two candidates' metrics with $|t| = 6.92$ at $k = 3$, BCa $[5.30, 9.04]$ (Tab. IV). The k-scan across $k = 3,4,5$ reveals the empirical signature of basis-size dominance: as Schoen's basis dimension grows monotonically ($27 \to 48 \to 75$), the strict-converged Schoen subset shrinks ($16/20 \to 7/20 \to 6/20$) and the heavy-tailed σ_{Schoen} distribution increasingly drives the $|t|$-statistic. To test whether this is a real geometric signal or a basis-size artifact, we performed a paired-bootstrap test at matched basis size $n_b = n_b^{\mathrm{Schoen}} = 27$ between the two candidates' metrics, 20 seeds, $B = 10^4$ resamples (Tab. V). The signed gap *inverts* when bases are matched, with disjoint BCa CIs and bootstrap probability of opposite sign equal to unity, and is the empirical basis for excluding σ from the multi-channel Bayes factor (`for_combination = false` flag on the `ChannelEvidence` record). The matched-basis comparison is currently implemented as a lex-sorted truncation of the larger TY basis to Schoen's 27-monomial dimension; a geometric basis projection (eigenvalue-ranked principal-components of the Laplacian, or harmonic-form ranking) is the in-flight upgrade to the matched-basis defense (`REMEDIATION_PLAN_OPTION_B.md` Phase 3.3). The σ-discriminability is reported *only* as a separate $|t|$-statistic, not as a likelihood contribution; the Bayes-factor exclusion of σ is conditional on the matched-basis defense surviving the geometric-projection upgrade.

TABLE II. Structural comparison of the substrate's predicted Calabi–Yau threefold (Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$, Donagi–Ovrut–Pantev–Reinbacher [12]) against the closest distinct three-generation competitor (Tian–Yau $\mathbb{Z}_3$). The framework's substrate-physical forcing chain (self-conjugacy on the internal Hodge decomposition, two-stage Wilson breaking $E_8 \to G_{\mathrm{SM}}$, Atiyah–Singer three-generation index) selects the Hodge $(3,3)$ / $\pi_1 = \mathbb{Z}_3 \times \mathbb{Z}_3$ branch on which Schoen sits and Tian–Yau does not.

Quantity	Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$	Tian–Yau $\mathbb{Z}_3$	Selecting hypothesis
$(h^{1,1}, h^{2,1})$	$(3,3)$	$(6,9)$	self-conjugacy + 3 generations
π_1	$\mathbb{Z}_3 \times \mathbb{Z}_3$	$\mathbb{Z}_3$	two-stage Wilson breaking
$\chi(X)$	0	-6	deformation-class distinction
$2\,\omega_{\mathrm{fix}}$	246/248	245/248	gateway-formula ratio

TABLE III. Multi-channel Bayes factor for the substrate-CY3 hypothesis test, $k = 3$, $n_{\mathrm{pts}} = 25{,}000$ chain / 40,000 σ, 20 seeds. Sign convention: $\ln \mathrm{BF} > 0 \Leftrightarrow$ Tian–Yau favoured, $< 0 \Leftrightarrow$ Schoen favoured. The combined chain-channel result clears decisive Bayesian strength on the Kass–Raftery scale ($|\ln \mathrm{BF}| > 5$). The σ-channel is reported separately as discriminability and is excluded from the model-comparison Bayes factor (rationale and empirical defence: Tab. V).

Channel	Likelihood	$\ln \mathrm{BF}_{\mathrm{TY/Schoen}}$	n_σ-equivalent
$\mathrm{chain}_{\mathrm{quark}}$ ($\{m_b, m_s, m_d\}$)	log-χ^2 to PDG	-5.60	—
$\mathrm{chain}_{\mathrm{lepton}}$ ($\{m_\tau, m_\mu, m_e\}$)	log-χ^2 to PDG	-9.16	—
Combined (physics preference)	—	$\mathbf{-14.76}$	$\mathbf{5.43\sigma}$
Strength (Kass–Raftery)	—	**Strong (Schoen-favoured)**	
Posterior odds Schoen:TY	—	$2.6 \times 10^6 : 1$	
σ-channel (discriminability only; *excluded from BF*)	Welch t on σ-distribution	—	6.92σ

c. Tian–Yau Yukawa-channel structural exclusion. Beyond the chain-channel Bayesian preference, we attempted to construct a working three-generation Standard-Model bundle on $\mathrm{TY}/\mathbb{Z}_3$ in the bicubic CICY ambient $\mathbb{P}^3 \times \mathbb{P}^3$ at $h^{1,1} = 2$, mirroring the working three-factor monad construction we use on Schoen (Tab. VII). A nine-cycle hypothesis-test enumeration over rank-3 SU(3) line-bundle-sum and monad bundles, summarised in Tab. VI, found no candidate that simultaneously satisfies $c_1(V) = 0$, $\Sigma\, c_3(V) = \pm 18$, $\mathrm{ch}_2(V) \leq \mathrm{ch}_2(TX)$, slope-stability via line-destabiliser test, balanced 1:1:1 Wilson $\mathbb{Z}_3$ phase partition, basis-rank ≥ 6 for three-generation Yukawa population, and Mumford–Takemoto SU(3) stability [23]. The line-bundle-sum result reproduces and extends Anderson–Gray–Lukas–Palti [17] §5.3's catalogue exclusion of $h^{1,1} \in \{2,3\}$ to broader bundle constructions and to the explicit Mumford–Takemoto stability check. The leading rank-7 monad survivor of the topological + summand-stability constraints, $V_{\min 2}$, is mathematically falsified by $h^0(V_{\min 2}) = \mathbb{C}^2 \neq 0$ in the long exact sequence of $V = \ker(B \to C)$, violating the $h^0(V) = 0$ requirement for stable rank-3 SU(3) bundles on a Calabi–Yau threefold (cf. [23] §1.2). This structural exclusion is independent of the Bayesian-discrimination evidence and converges on the same conclusion via different mathematical machinery.

d. Catalogue uniqueness at $h^{1,1} = h^{2,1} = 3$. The chain-channel Bayes factor and the Tian–Yau Yukawa exclusion identify Schoen as the favoured candidate within the framework's pre-narrowed two-element shortlist. To upgrade this from a pairwise discrimination to a catalogue-level uniqueness statement we performed a four-cycle sweep of the published smooth-CY3 universe at $(h^{1,1}, h^{2,1}) = (3,3)$ for non-Schoen entries admitting a free $\mathbb{Z}_3 \times \mathbb{Z}_3$ action with Standard-Model-compatible Wilson-line breaking. Cycle 1 (Hodge filter) returned zero hard-match competitors across CICY (Anderson–He–Lukas [18]), Schoen-class fiber products, Constantin–Lukas free-quotient catalogue, Borisov–Caldararu Pfaffian threefolds, and gCICY (Anderson–Apruzzi–Gao–Gray–Lee [20]). Cycle 2 (free-action structural filter) returned zero non-Schoen survivors. Cycle 2.5 (computer-algebra sweep) verified by direct query against the Kreuzer–Skarke 473,800,776-polytope catalogue [19] that the $(3,3)$ Hodge bin is *literally empty* (response `#NF: 0` at output limits $L \in \{100, 1000, 10000\}$, cross-verified by sweeping the entire $\chi = 0$ self-mirror diagonal $h = 1, \ldots, 14$ and confirming the bin is empty for $h \leq 13$ and populated from $h = 14$). Cycle 2.6 (gCICY non-manifest Aut residue) reduced the last open cell by programmatic literature search of arXiv:1507.03235 for the literal $(3,3)$ Hodge pair (zero hits), cross-checked against Larfors–Schneider [21] and Constantin–Lukas–Mishra [22]; the strict-CAS portion of cycle 2.6 (computing the full ambient automorphism group via Sage / Macaulay2 / polymake and checking which $\mathbb{Z}_3 \times \mathbb{Z}_3$ subgroups act freely on a generic gCICY $(3,3)$) remains **DEFERRED** pending a CAS-equipped re-run; see `p_schoen_uniqueness_cycle2_5.md` §3.4 and the in-flight Phase 6.2 of the remediation plan. The combined catalogue scan, modulo the gCICY non-manifest CAS residue, supports the statement: *no published smooth Calabi–Yau threefold in the catalogued universe is established to satisfy the substrate-physical four-principle constraint at $(h^{1,1}, h^{2,1}) = (3,3)$ except Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$.* The remaining residual is a future gCICY codim $\geq (3,1)$ ex-

TABLE IV. σ-channel discriminability k-scan on TY-vs-Schoen, Donaldson balanced metric (no Adam refinement), $n_{\text{pts}} =$ 40,000, 20 seeds, BCa-95% CI from $B = 10^4$ bootstrap resamples (DiCiccio–Efron 1996). Strict-converged tier: residual $< 10^{-6}$ AND iters $<$ cap, no Tukey trim, no guard-snapshot inflation. The $k = 3$ row is the publication discriminability headline; $k = 4, 5$ are reported as *supporting tiers at higher basis order*, not as parallel publication results. The monotone shrinkage of the strict-converged Schoen subset ($16/20 \to 7/20 \to 6/20$) as Schoen's basis dimension grows ($27 \to 48 \to 75$) is the empirical signature of σ as a basis-size-saturated indicator at finite k, motivating exclusion from the model-comparison Bayes factor (Tab. V).

k	$\langle\sigma_{\text{TY}}\rangle$	$\langle\sigma_{\text{Schoen}}\rangle$	$\lvert t\rvert$-stat	Tier-0 BCa 95% CI	$n^{\text{strict}}_{\text{Schoen}}/n^{\text{Schoen}}_{\text{basis}}$
3 (publication)	0.272	5.59	$\mathbf{6.92}\sigma$	[**5.30**, **9.04**]	16/20 / 27
4 (damped, supporting)	0.278	13.21	3.82σ	[2.48, 6.37]	7/20 / 48
5 (supporting)	0.295	13.05	10.87σ	[8.79, 15.01]	6/20 / 75

TABLE V. Empirical defence of the σ-channel exclusion from the multi-channel Bayes factor: matched-basis-size paired bootstrap test on TY-vs-Schoen, $k = 3$, $n_{\text{pts}} =$ 40,000, donaldson_iters= 100, donaldson_tol= 10^{-6}, 20 seeds, $B = 10^4$ resamples, boot_seed $\in \{12345, 999, 31415\}$ for jitter check. At native basis sizes $\sigma_{\text{TY}} < \sigma_{\text{Schoen}}$ as expected. **When the two metrics' bases are truncated to a common dimension $n_b = 27$ (Schoen's native value), the sign of the σ-gap inverts**, with disjoint BCa CIs and paired-bootstrap probability of opposite sign equal to unity over 10^4 resamples. The artifact magnitude is $\sim 3\times$ the native gap, in the opposite direction. This is the textbook signature of a basis-size-induced artifact, not a likelihood contribution, and is the empirical basis for excluding σ from the model-comparison Bayes factor of Tab. III.

Quantity	Mean	BCa 95% CI
$\Delta\sigma_{\text{native}}$ ($n_{\text{TY}} = 87$, $n_{\text{Schoen}} = 27$)	-7.89	$[-15.25, -5.52]$
$\Delta\sigma_{\text{matched}}$ ($n_{\text{TY}} = n_{\text{Schoen}} = 27$)	$\mathbf{+23.59}$	[**+11.33**, **+42.66**]
$\lvert\Delta_{\text{matched}}\rvert/\lvert\Delta_{\text{native}}\rvert$	2.99	[1.54, 6.55]
$\mathbb{P}$(opposite sign | boot)	**1.0000** (10^4 resamples)	
Bootstrap-seed jitter on BCa-low	0.52 (range over $\{12345, 999, 31415\}$)	

tension producing a $(3, 3)$ entry, plus the deferred CAS check on existing gCICY $(3, 3)$ classifications; both would require a re-test but do not affect the present empirical identification.

e. Net empirical statement. Three convergent lines of evidence support the empirical identification of the substrate's internal Calabi–Yau threefold as Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ over the closest structural competitor Tian–Yau $\mathbb{Z}_3$. **Closed:** (a) the four-cycle catalogue uniqueness sweep finds no non-Schoen $(h^{1,1}, h^{2,1}) = (3, 3)$ smooth CY3 with a free $\mathbb{Z}_3 \times \mathbb{Z}_3$ action across the published universe (CICY, Kreuzer–Skarke toric, Schoen-class, Constantin–Lukas, Borisov–Caldararu Pfaffian, gCICY); the Kreuzer–Skarke direct query records `#NF: 0` at output limits up to $L = 10{,}000$, cross-verified by sweeping the entire $\chi = 0$ self-mirror diagonal $h \in \{1, \ldots, 14\}$. (b) the Tian–Yau Yukawa channel is structurally excluded by a nine-cycle hypothesis-test enumeration showing that no rank-3 SU(3) stable polystable monad bundle on $\text{TY}/\mathbb{Z}_3$ in the bicubic ambient at $h^{1,1} = 2$ in the searched bidegree range simultaneously satisfies all required physical constraints. **Indicative, not converged:** (c) the chain-channel Bayes factor at $\ln \text{BF} = -14.76$ (posterior odds 2.6×10^6:1 Schoen:TY; nested-LRT σ-equivalent 5.43σ) is computed from a metric-Laplacian Galerkin-residual sweep that has not yet met the strict 10% basis-size convergence criterion at $k_{\text{test}} \in \{2, 3, 4, 5\}$; we report it as supporting evidence and expect to retract or upgrade the figure on completion of the full PDG-fermion-mass Yukawa-cohomology pipeline (Phase 2 of the remediation plan). The Donaldson σ-channel is excluded from the model-comparison Bayes factor by a paired-bootstrap matched-basis sign-reversal test (currently using lex truncation; geometric-projection upgrade is in flight at Phase 3). *Conditional on the open methodological items below and the in-flight pipeline upgrades enumerated in Sec. XV, the substrate-physical four-principle constraint narrows the candidate manifold to a single element of the catalogued universe at the searched Hodge bin.*

f. What this section does not *claim.* We are explicit about the residual scope. (i) We do *not* claim Schoen-uniqueness in full generality across the Calabi–Yau threefold landscape; the catalogue uniqueness statement is at the single Hodge bin $(h^{1,1}, h^{2,1}) = (3, 3)$ and is conditional on the gCICY non-manifest CAS residue (item iii below). (ii) The chain-channel Bayes factor $\ln \text{BF} = -14.76$ is *not* a converged result at the strict 10% basis-size convergence criterion: both the quark-chain and lepton-chain JSON outputs (`p7_11_quark_chain.json`, `p7_11_lepton_chain.json`) carry `converged: false`. The figure is reported as indicative; the full PDG-fermion-mass Yukawa-cohomology pipeline replacement is in-flight Phase 2 of the remediation plan. (iii) The gCICY $(3, 3)$ uniqueness check has **deferred** the strict computer-algebra-system portion of cycle 2.6 (Sage / Macaulay2 / polymake-based ambient automorphism enumeration) pending CAS environment availability; the

TABLE VI. Structural exclusion of a working three-generation Standard-Model bundle on the closest competitor manifold Tian–Yau $\mathbb{Z}_3$ (bicubic CICY in $\mathbb{P}^3 \times \mathbb{P}^3$, $h^{1,1} = 2$): nine-cycle empirical enumeration. Constraints applied per cycle, in chain: $c_1(V) = 0$ AND $\Sigma\, c_3(V) = \pm 18$ for three generations downstairs after the $\mathbb{Z}_3$ quotient AND $\mathrm{ch}_2(V) \leq \mathrm{ch}_2(TX)$ for anomaly cancellation AND slope-stability via line-destabiliser test on the Kähler cone AND balanced 1:1:1 Wilson $\mathbb{Z}_3$ phase partition on V AND basis-rank ≥ 6 for three-generation Yukawa population AND Mumford–Takemoto SU(3) stability (Huybrechts–Lehn 2010 §1.2: $h^0(V) = 0$). Cycles run on monad bundles $V = \ker(B \to C)$ with B, C line-bundle sums, bidegrees in $[-2, 2]^2$. The cycle 1 line-bundle-sum exclusion empirically reproduces and extends Anderson–Gray–Lukas–Palti (arXiv:1202.1757) §5.3's published result that no phenomenologically viable model exists at $h^{1,1} \in \{2, 3\}$ in their scan. No candidate survives all nine constraints simultaneously.

Cycle	Hypothesis (rank-3 SU(3) bundle on TY/$\mathbb{Z}_3$ bicubic)	Failure mode	Survivors
H1	line-bundle sums, Wilson 1:1:1, $c_3 = \pm 18$	polystability fails universally	$0/{\sim}5.97 \times 10^6$
H2 v1/v2	monad $\ker(B \to C)$, weak polystability filter	polystability under correct slope	$0/-$
H3	SU(4)/SO(10) rank-4 alternative embedding	same polystability obstruction	$0/{\sim}9 \times 10^4$
H2 v3	monad, closed-negative-cone polystability	(passes topological + summand-stable)	$10/{\sim}3.78 \times 10^6$
C4	$V_{\min}$ wired ($\mathrm{rk}\, B = 4$)	basis-rank deficit; 3 modes $\neq 9$	0 (operationally)
C5	rank-$B \geq 6$ rescan	(10 v3 candidates re-confirmed)	10
C6	$V_{\min 2}$ wired ($\mathrm{rk}\, B = 7$)	H^0-seed empty on neg. summands	0 (operationally)
C7	H^0/H^1 cohomology audit	(LES-collapse argument validates AKLP/Schoen)	—
C8	$V_{\min 2}$ mathematical falsification	$h^0(V) = \mathbb{C}^2 \neq 0$: violates Mumford–Takemoto	**0**
C9	$H^0(V) = 0$ filter on cycle-5 v3 set	operationally infertile: $\Sigma h^0(B) = \Sigma h^1(B) = 0$	**0/10**
Net empirical result: no rank-3 SU(3) Yukawa-active SM monad bundle in scanned bidegree range.			

TABLE VII. Working three-generation Standard-Model bundle on Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ used by the chain channels of Tab. III. Three-factor monad construction $V = \ker(B \to C)$ on the Schoen ambient $\mathbb{P}^2 \times \mathbb{P}^2 \times \mathbb{P}^1$. The Wilson $\mathbb{Z}_3 \times \mathbb{Z}_3$ phase partition splits the V-summands evenly into three classes (one per generation) under the projection $(a - b) \bmod 3$ on the bidegrees, breaking the previously degenerate sector assignment that affects Tian–Yau under the same construction (Tab. VI). Anomaly cancellation $\mathrm{ch}_2(V) \leq \mathrm{ch}_2(TX)$ is satisfied with the indicated 5-brane class.

Object	Specification	Property
B	$\mathcal{O}(1,0,0)^2 \oplus \mathcal{O}(0,1,0)^2 \oplus \mathcal{O}(0,0,1)^2$	rank 6
C	$\mathcal{O}(1,1,0) \oplus \mathcal{O}(0,1,1) \oplus \mathcal{O}(1,0,1)$	rank 3
$V = \ker(B \to C)$	rank 3 SU(3)	$c_1(V) = 0$, $\Sigma\, c_3(V) = \pm 18$
Wilson Z_3 partition on V	$(1,1,1)$ across phase classes	balanced; canonical sector assignment
Anomaly cancellation	$\mathrm{ch}_2(V) \leq \mathrm{ch}_2(TX)$ with $W = (27, 27)$	satisfied
Yukawa coverage on V	$Y_u : 9/9$, $Y_d : 9/9$ on h_0 slice	operationally productive

literature-cross-check portion is closed. A future gCICY codim $\geq (3,1)$ extension producing a $(3,3)$ entry would also require a re-test. (iv) We do *not* claim that the Bayesian discrimination is independent of the Donaldson convention adopted for the Schoen and Tian–Yau metrics; the σ-channel result is sensitive to h-inversion handling and Kähler-form projection, and we report only the conventions used here. The σ-channel matched-basis defense currently uses lex truncation rather than a geometric basis projection; the geometric-projection upgrade and the full Donaldson convention robustness sweep are in-flight Phase 3 of the remediation plan. (v) The TY Yukawa exclusion is at the searched bidegree range in the bicubic ambient and is consistent with, but does not rigorously imply, the absence of any TY-supported three-generation bundle outside that range. (vi) The bundle data we use on Schoen, and the line-bundle enumerations cited from the literature, are subject to the bundle-provenance audit summarised in Sec. XV. None of items (i)–(vi) affects the numerical computations of Secs. III–XI, which proceed from the Schoen identification taken as a starting commitment.

C. The E_8 Coxeter spectrum

The E_8 Lie algebra has dimension $\dim(E_8) = 248$, rank 8, 240 roots, Coxeter number $h(E_8) = 30$, and Weyl group order $|W(E_8)| = 696{,}729{,}600$. The exponents of E_8 are the eight integers

$$\{m_1, \ldots, m_8\} = \{1, 7, 11, 13, 17, 19, 23, 29\}, \qquad (3)$$

palindromically symmetric around $h/2 = 15$ in pairs $(1, 29)$, $(7, 23)$, $(11, 19)$, $(13, 17)$, satisfying $m_j + m_{9-j} = h$. The longest Weyl element acts as $w_0 = -1$ on the Cartan subalgebra of E_8 (a property shared by $A_1, B_n, C_n, D_{2n}, E_7, E_8, F_4, G_2$). The harmonic threshold $h/2 = 15$ is the framework's chirality-flip locus for the Coxeter chain and is not itself an exponent.

a. Framework extension beyond the standard Coxeter spectrum. The mass-formula construction in this paper parameterizes *continuous* chain positions $k \in \mathbb{R}_{\geq 0}$, not

TABLE VIII. Methodological provenance and replicability conventions for the CY3-substrate-identification analysis (Tabs. II–VII). Run-time per sweep: ~ 2.5 h CPU on the basis-convergence sign-reversal test; ~ 3.5 h GPU on each k-scan production run; total $\sim$50 GPU-hours plus $\sim$30 CPU-hours across all reported tiers. The σ-discriminator throughout this paper is the weighted-L^2 Donaldson 2009 §3.4 functional, matching the production implementation; see row below.

Item	Setting / source
σ-functional definition	weighted L^2 Monge–Ampère residual: $\sigma^2 = \langle(\eta/\kappa - 1)^2\rangle_w$, $\kappa = \langle\eta\rangle_w$ (Donaldson 2009 §3.4 eq. 3.4)
σ-functional implementation	`schoen_metric.rs::compute_sigma`, `ty_metric.rs::compute_sigma` (CPU + GPU bit-exact mirror)
Bayes-factor sign convention	$\ln \mathrm{BF}_{12} > 0 \Leftrightarrow M_1$ favoured (Jeffreys/Kass–Raftery)
Strict-converged tier gate	residual $<$ tol AND iters $<$ cap, no Tukey trim
Bootstrap CI methodology	percentile and bias-corrected accelerated, $B = 10^4$ resamples
Bootstrap-seed jitter check	boot_seed $\in \{12345, 999, 31415\}$, target $< 0.05\sigma$ on BCa-low
Donaldson update (damped, $k \geq 4$)	$h \leftarrow \alpha\, T(G)^{-1} + (1-\alpha)\, h_{\mathrm{old}}$, $\alpha = 0.5$
GPU/CPU bit-exactness (h-pair)	mirrored 256-lane tree reduction, iter-0 $\Delta < 4.4 \times 10^{-16}$
Reproducibility manifest	SHA-256 chained event log per run
n_σ-equivalent from $\ln \mathrm{BF}$	$n_\sigma = \sqrt{2\lvert\ln \mathrm{BF}\rvert}$ (Wilks/Cowan–Cranmer–Gross–Vitells)
Source code	`book/scripts/cy3_substrate_discrimination/rust_solver/`
Output JSON catalogue	`output/p{5_10,7_11,8_1,_basis_convergence}_*.json`
Research log (9-cycle TY exclusion)	`references/p_ty_bundle_research_log.md`
Consolidated audit document	`references/cy3_publication_summary.md`

just the eight integer exponents above. Half-integer positions ($k_d = 6.5$, $k_W = 34.5$, $k_t = 36.5$, etc.) and "cycle-2" positions ($k > h$) are framework-specific extensions that interpolate between integer exponents and re-use the periodicity of the Coxeter element across multiple cycles. These are structural assignments within the framework rather than standard Coxeter-theoretic invariants. The structural "boundary capacity" $h + m_{\max} = 59$ used in the dark-sector coupling family (Sec. V) is likewise a framework combination of E_8 invariants rather than an established Lie-algebra invariant in its own right.

D. The McKay correspondence

The McKay correspondence [5] establishes a bijection between finite subgroups of $SU(2)$ and the affine simply-laced ADE Dynkin diagrams: cyclic $\mathbb{Z}_n \leftrightarrow \widehat{A}_{n-1}$, binary dihedral $\leftrightarrow \widehat{D}_n$, binary tetrahedral $\leftrightarrow \widehat{E}_6$, binary octahedral $\leftrightarrow \widehat{E}_7$, and **binary icosahedral** $\leftrightarrow \widehat{E}_8$. The McKay correspondence itself identifies vertices of the affine $\widehat{E}_8$ Dynkin diagram with irreducible representations of the binary icosahedral group; it does *not*, on its own, single out a particular integer. Separately, the icosahedral rotation group A_5 contains elements of order 5 in the $(2,3,5)$-triangle group, with $\varphi = 2\cos(\pi/5) = (1+\sqrt{5})/2$ as the icosahedral angular ratio. The framework adopts this fivefold rotation order, paired with the binary icosahedral McKay-dual of E_8, as the seed integer for the φ-chain. The use of the integer 5 is therefore a framework structural commitment combining the McKay-dual identification with the icosahedral rotation order, not a forced output of the McKay correspondence.

a. Two harmonic chains used by the framework. The mass-formula construction below distinguishes a φ-chain (with φ^k primary harmonics, motivated by 5-fold icosahedral recursion at the binary-icosahedral McKay-dual of E_8) from a $\sqrt{2}$-chain (with $(\sqrt{2})^k$ primary harmonics, motivated by 4-fold octahedral recursion at the binary-octahedral McKay-dual of E_7, where $\sqrt{2} = 2\cos(\pi/4)$). Charged leptons and neutrinos populate the φ-chain; quarks and electroweak bosons populate the $\sqrt{2}$-chain. This sector split is a *framework structural assignment* rather than a derivation from E_8 alone — a single E_8 algebra does not naturally decompose into icosahedral and octahedral sub-recursions, and the allocation of leptons to φ and quarks to $\sqrt{2}$ is adopted on phenomenological grounds (matching observed masses with small-integer Coxeter-step assignments). A first-principles derivation of the chain assignment from a specific Calabi–Yau compactification is left for future work.

b. Chain factors are not Coxeter eigenvalues. We emphasise that the harmonic factors φ^k and $(\sqrt{2})^k$ that appear throughout the mass formulas below are *exponential growth/recursion factors* on a discrete chain — successive applications of a fixed scalar step — and are *not* eigenvalues of the Coxeter element of E_8. The latter are the unit-modulus roots of unity $|e^{2\pi i e_j/h}| = 1$ for $e_j \in \{1, 7, 11, 13, 17, 19, 23, 29\}$, which describe rotational structure on the Coxeter plane and have nothing directly to do with mass-ratio recursion. The exponents e_j enter the mass formulas separately, as integer positions used to label palindromic pairings (e.g. $11 + 19 = h$); the multiplicative factors $\varphi^k, (\sqrt{2})^k$ encode the geometric step magnitude of the chain itself. These are two mathematically distinct uses of E_8 data, and the framework's construction is the joint use of both.

E. Substrate-tension interpretation

The framework adopts a substrate-tension interpretation of the Higgs mechanism, motivated by Ander-

son's superconducting precursor to gauge-boson mass generation [7]; the resulting picture is similar in spirit to Volovik's quantum-liquid / helium-droplet emergent-vacuum program [8, 9], though arrived at independently and from a different structural starting point. Within the substrate framework, the Higgs vacuum expectation value v is interpreted as the baseline relaxed tension of the $E_8 \times E_8$ substrate, and massive SM particles are localized deformations of this tension at structurally-determined Coxeter positions. Photons are massless because they correspond to solitonic excitations of the substrate that do not deform its baseline tension.

a. Foundations deferred to companion paper. The substrate-physical reading of c, $\hbar$, the elastic-medium identification $c = \sqrt{\Theta/\mathcal{I}}$, and the macroscopic emergence of general relativity, quantum mechanics, Newton's laws of motion, and Maxwell's equations from the same elastic substrate that hosts the $E_8 \times E_8$ Coxeter structure are not load-bearing for the numerical demonstration of this paper, and the present paper takes them as part of the starting commitment of Table I. Their derivation chains are developed in the companion paper [43] and the foundational monograph [42]. We do not repeat them here.

III. GATEWAY FORMULA

The remainder of this paper is the direct computation of SM observables from the starting assumptions of Table I. There is no per-particle scheme choice and no fitted continuous parameter: every numerical value reported below is the framework's structural prediction. PDG-matching residuals are predictions versus measurement, not fits versus measurement.

The electron sits at the gateway $k = 0$, defined as one chain step before the first E_8 exponent (at position 1). The chain depth from gateway to harmonic threshold is

$$M = h/2 - 1 = 14. \tag{4}$$

The framework's effective coupling scale is

$$h_{\text{eff}} \;=\; \frac{m_e}{2\omega_{\text{fix}}} \;=\; m_e \cdot \frac{\dim}{\dim - 2} \;=\; m_e \cdot \frac{124}{123}, \tag{5}$$

where $\omega_{\text{fix}} = 1/2 - 1/\dim = 123/248$ is the gateway normalization factor. Equivalently, with the Planck mass $m_{\text{Planck}} = \sqrt{\hbar c/G}$ taken from the structurally predicted G of Sec. XI, $h_{\text{eff}} = F_e \cdot m_{\text{Planck}} \cdot 124/123$, with F_e fixed entirely by Eq. (6) below; the two readings agree at the 17 ppm residual reported here, and the downstream mass formulas are operationally identical under either reading.

The dimensionless ratio of the electron mass to the Planck mass is predicted by the framework as

$$\boxed{F_e \;=\; \frac{m_e}{m_{\text{Planck}}} \;=\; \frac{\dim - 2}{\dim} \cdot \frac{5}{\dim \cdot h^M} \cdot \frac{1}{1 - \frac{1}{4(\dim - 1)}}.} \tag{6}$$

The right-hand side contains *no measured input*: every quantity is a structural integer of the starting assumptions (Table I). This is a zero-parameter prediction of the dimensionless electron-to-Planck mass ratio.

Each factor has a specific structural role in the framework:

- Gateway factor $(\dim - 2)/\dim = 246/248$: framework dual-sector count subtracting two distinguished endpoints (one in each E_8).
- McKay icosahedral 5: framework-selected icosahedral order motivated by the binary-icosahedral McKay dual of E_8.
- Coxeter chain $h^M = 30^{14} \approx 4.78 \times 10^{20}$: each step contributes a factor of $h = 30$.
- Sub-harmonic loss $988/987 = 1/(1 - 1/988)$: the geometric sum $\sum_{n\geq 0}(1/988)^n$ of a per-step leakage rate $1/(4(\dim - 1)) = 1/988$ to all orders. The factor of 4 in $4(\dim - 1)$ admits a dual-sector / two-orientation reading ("2×2": two E_8 factors, each carrying both an upward and a downward chain-step orientation; $\dim - 1$ removes the distinguished anchor mode in each sector); a substrate-physical derivation of the 4 from a heterotic-anomaly index or a $\mathbb{Z}_3 \times \mathbb{Z}_3$ quotient trace on the Schoen threefold is part of the open companion-paper program (Sec. XV).

Numerically, $F_e = 4.18547 \times 10^{-23}$. The measured value $m_e^{\text{PDG}}/m_{\text{Planck}}^{\text{CODATA}}$ agrees with this structural prediction at 17 ppm. Newton's constant is the zero-parameter consequence $G = F_e^2 \cdot \hbar c/m_e^2$ developed in Sec. XI.

IV. CHARGED-FERMION MASSES

A. Mass formula schema

Every charged-fermion mass in the framework takes the unified form

$$\boxed{m_X \;=\; h_{\text{eff}} \cdot R_X(k_X, \text{sector}, \{\text{corr.}\}),} \tag{7}$$

where $h_{\text{eff}} = m_e \cdot \dim/(\dim - 2) = 124 m_e/123$ is the gateway-normalized scale (Sec. III), k_X is the particle's Coxeter step assignment in the E_8 spectrum (cf. App. A), and R_X is a structural ratio built from E_8 invariants $\{\dim, h, M, \text{rank}, m_{\max}\}$, the gateway-mixing parameter $\alpha_\tau = 1/(\dim + h) = 1/278$, and sector-specific sub-harmonic corrections. Numerical values of R_X for every particle are tabulated in Sec. XIII (Table XIII); the closed-form expressions are given below sector by sector.

B. Lepton sector

The electron, muon, and tau occupy integer E_8 exponents on the φ-chain (icosahedral bulk recursion) at $k_e = 0$ (gateway), $k_\mu = 11$, and $k_\tau = 17$ respectively. Their explicit closed-form ratios are

$$R_e = \frac{\dim -2}{\dim} = \frac{246}{248}, \tag{8}$$

$$R_\mu = \varphi^{11} \cdot \big(1 + \Delta_\varphi\big) \cdot \big(1 - \alpha_\mu\big), \tag{9}$$

$$R_\tau = \varphi^{17} \cdot \big(1 - \Delta_\varphi\big) \cdot \big(1 - \alpha_\tau^{\text{full}}\big), \tag{10}$$

where the icosahedral-bulk closure factor is

$$\Delta_\varphi = \frac{1}{h + \varphi^2}, \tag{11}$$

and the sub-harmonic corrections are

$$\alpha_\mu = \sum_{n=1}^{\infty} c_n^{(\mu)} \alpha_\tau^n = 2\alpha_\tau^2 + \mathcal{O}(\alpha_\tau^3), \tag{12}$$

$$\alpha_\tau^{\text{full}} = \frac{\alpha_\tau}{1 + \varphi^2 \alpha_\tau} = \frac{1}{\dim + h + \varphi^2}. \tag{13}$$

Each lepton ratio has three structural pieces: a φ-chain primary harmonic φ^{k_ℓ} (icosahedral bulk recursion at the lepton's Coxeter exponent), an icosahedral-bulk closure $(1 \pm \Delta_\varphi)$ with sign opposite for muon and tau (reflecting their opposite displacement from the harmonic threshold $h/2 = 15$), and a sub-harmonic series correction. The muon coefficient $c_2^{(\mu)} = 2$ is the framework's minimal assignment to the palindromic pair $(11, 19)$ centered on $h/2$. The tau correction $\alpha_\tau^{\text{full}}$ is the closed-form sum of the geometric series $\sum_n (-\varphi^2 \alpha_\tau)^n \alpha_\tau$ in the icosahedral self-coupling. Predicted (electron-anchor): $m_\mu = 105.6583756$ MeV (PDG match: $+1.1$ ppb), $m_\tau = 1776.8597$ MeV (-0.15 ppm).

C. Quark sector

Quarks occupy positions on the $\sqrt{2}$-chain (octahedral bulk recursion) at integer and half-integer Coxeter steps. Their explicit closed-form ratios are

$$R_q = (\sqrt{2})^{k_q} \cdot \big(1 + \chi_q \Delta_2\big) \cdot \big(1 + \lambda_q \tfrac{|k_q - h/2|}{2\dim}\big) \cdot \big(1 - \alpha_q\big), \tag{14}$$

where the octahedral-bulk closure factor is

$$\Delta_2 = \frac{1}{32}, \tag{15}$$

and $\chi_q, \lambda_q \in \{-1, 0, +1\}$ are per-quark chirality and per-step-loss signs respectively. The per-quark gateway-mixing sub-harmonic series α_q is non-zero only for the top quark, the only cycle-2 quark in the framework. For all cycle-1 quarks $\alpha_q = 0$; for the top quark,

$$\alpha_t = \frac{c_2^{(t)} \alpha_\tau^2}{1 + \alpha_\tau / h_{SU(3)}}, \quad c_2^{(t)} = \left(\frac{h}{h_{SU(3)}}\right)^2 = 100, \tag{16}$$

with $h_{SU(3)} = 3$ the colour Coxeter number. The leading coefficient $c_2^{(t)} = (h_{E_8}/h_{SU(3)})^2$ is the topological-subgroup ratio between the substrate's full Coxeter period and the colour Coxeter period, squared: ten colour cycles fit inside one full substrate cycle, and the top – the only cycle-2 colour-charged fermion – projects through this ratio twice (once for its colour-charge winding, once for its cycle-2 chain-position projection). The closure denominator $h_{SU(3)} = 3$ is the same colour-projected closure that appears in α_μ and α_H throughout the framework. For the strange quark sitting exactly at the harmonic threshold $k_s = h/2 = 15$, the per-step-loss factor $|k_s - h/2|/(2\dim) = 0$ regardless of the sign of λ_s, so λ_s is indeterminate (we record it as $\lambda_s = 0$ purely as bookkeeping notation). The chirality factor $\chi_s = 0$, in contrast, is a *distinct structural assignment*: it forces the closure factor $(1 + \chi_s \Delta_2) = 1$ at the chirality-flip locus and is the only non-trivial sign choice the strange quark makes:

TABLE IX. Per-quark Coxeter step k_q, chirality sign χ_q, and per-step-loss sign λ_q in the quark mass formula Eq. (14).

Quark	k_q	χ_q	λ_q	Cycle
up	4	+1	+1	1
down	6.5	−1	−1	1
strange	15	0	0	1 (threshold)
charm	22.5	+1	−1	1
bottom	26	−1	+1	1
top	36.5	+1	+1	2 (cycle 2)

The denominator $2\dim = 496$ in the per-step-loss term is the framework's chosen normalization by the full $E_8 \times E_8$ algebra dimension (visible plus dark sectors). All six quark masses lie within current PDG 2024 uncertainty bands; the worst central-value residual is $+5618$ ppm on the up quark, well within the $\sim$25% experimental uncertainty on the $\overline{\text{MS}}$ light-quark mass at $\mu = 2$ GeV.

D. Term-by-term decomposition: charged fermions

Each charged-fermion mass function decomposes into four contributions: the *primary harmonic* (Coxeter step magnitude), the *sub-harmonic correction* (recursive coupling at order α_τ^k), the *sector modifier* (chain-type factor: φ-chain or $\sqrt{2}$-chain), and the *dark-sector leakage* (boundary correction for cycle-2 modes only). For cycle-1 fermions the dark-sector leakage is absent. The decomposition for each particle:

a. Electron ($k_e = 0$).

$$m_e = h_{\text{eff}} \cdot \underbrace{\frac{\dim -2}{\dim}}_{\text{gateway factor}}$$

- *Primary harmonic:* the gateway-factor ratio $(\dim -2)/\dim = 246/248$ is the leading contribution. The numerator $\dim -2 = 246$ subtracts the

two dual-sector endpoints (one in each E_8); the denominator normalizes to the full algebra dimension.

- *Sub-harmonic correction:* none; the electron is the gateway reference and absorbs all sub-harmonic structure into the definition $h_{\text{eff}} = m_e \cdot \dim/(\dim -2) = 124 m_e/123$.

- *Sector modifier:* gateway position (no chain-type distinction).

- *Dark-sector leakage:* not applicable (cycle-1 mode at $k = 0$).

b. Muon ($k_\mu = 11$), $R_\mu \approx 205.10$.

- *Primary harmonic:* φ^{11}, the icosahedral bulk-recursion contribution at the muon's φ-chain position $k_\mu = 11$ (the third E_8 exponent, on the Φ-below-threshold side of the palindromic spectrum).

- *Icosahedral-bulk closure:* $(1 + \Delta_\varphi)$ with $\Delta_\varphi = 1/(h + \varphi^2)$. The positive sign reflects the muon's position below the harmonic threshold $h/2 = 15$.

- *Sub-harmonic correction:* $(1-\alpha_\mu)$ at order α_τ^2 with leading coefficient $c_2^{(\mu)} = 2$, assigned to the palindromic pair $(11, 19)$ centered on $h/2$.

- *Sector modifier:* φ-chain (icosahedral, lepton).

- *Dark-sector leakage:* not applicable (cycle-1 mode).

c. Tau ($k_\tau = 17$), $R_\tau \approx 3449.19$.

- *Primary harmonic:* φ^{17}, the icosahedral bulk-recursion contribution at the tau's φ-chain position $k_\tau = 17$ (the fourth E_8 exponent, the first above-the-threshold exponent).

- *Icosahedral-bulk closure:* $(1 - \Delta_\varphi)$ with $\Delta_\varphi = 1/(h + \varphi^2)$. The negative sign reflects the tau's position above the harmonic threshold $h/2 = 15$, opposite to the muon.

- *Sub-harmonic correction:* $(1 - \alpha_\tau^{\text{full}})$ where $\alpha_\tau^{\text{full}} = \alpha_\tau/(1+\varphi^2\alpha_\tau) = 1/(\dim + h + \varphi^2)$, the closed-form sum of the geometric series in the icosahedral self-coupling φ^2.

- *Sector modifier:* φ-chain (icosahedral, lepton).

- *Dark-sector leakage:* not applicable.

d. Quarks (six entries via Eq. (14)*).* The unified quark formula $R_q = (\sqrt{2})^{k_q} \cdot (1 + \chi_q \Delta_2) \cdot (1 + \lambda_q |k_q - h/2|/(2\dim))$ has three structural pieces common to all six quarks, with per-quark signs χ_q and λ_q from Table IX:

- *Primary harmonic:* $(\sqrt{2})^{k_q}$, the octahedral bulk-recursion contribution at the quark's Coxeter step.

- *Octahedral-bulk closure:* $(1 + \chi_q \Delta_2)$ with $\Delta_2 = 1/32$. The chirality sign χ_q alternates across the spectrum: $+1$ for "up-type" quarks (up, charm, top) and -1 for "down-type" quarks (down, bottom). The strange quark at the harmonic threshold has $\chi_s = 0$ (the closure factor collapses to 1 at the chirality-flip locus).

- *Per-step loss:* $(1 + \lambda_q |k_q - h/2|/(2\dim))$ with sign $\lambda_q \in \{-1, 0, +1\}$ per quark. The denominator $2\dim = 496$ is the framework's normalization by the full $E_8 \times E_8$ algebra dimension (visible plus dark sectors), reflecting the per-step leakage across both sectors as the chain advances away from the harmonic threshold $h/2$. The strange quark at $k_s = h/2$ has $|k_s - h/2| = 0$, so the loss factor is automatically unity for any value of λ_s; the entry $\lambda_s = 0$ in Table IX is a notational placeholder, not an independent structural assignment.

- *Sector modifier:* $\sqrt{2}$-chain at integer or half-integer position. The strange quark sits exactly at the harmonic threshold, the chirality-flip locus.

- *Dark-sector leakage:* the top quark is the only cycle-2 quark ($k_t = 36.5 > h = 30$); for the top, the $|k_t - h/2|$ distance is computed as written and the same formula Eq. (14) applies without an additional cycle-extension modifier.

- *Cycle-2 sub-harmonic correction (top quark only):* $\alpha_t = (h/h_{SU(3)})^2 \alpha_\tau^2/(1 + \alpha_\tau/h_{SU(3)})$, with leading coefficient $(h/h_{SU(3)})^2 = 100$ – the topological subgroup ratio between the substrate's full Coxeter period $h_{E_8} = 30$ and the colour Coxeter period $h_{SU(3)} = 3$, squared (one factor for the top's colour-charge winding, one for its cycle-2 chain-position projection). Closure denominator $h_{SU(3)} = 3$. The other five quarks (cycle-1) have $\alpha_q = 0$ in Eq. (14); only the top, as the unique cycle-2 colour-charged matter mode, picks up the topological subgroup ratio in its sub-harmonic series.

V. HIGGS SECTOR AND DARK-SECTOR COUPLING FAMILY

A. Boson mass-formula schema and structural corrections

The electroweak boson masses and the Higgs vev occupy positions on the $\sqrt{2}$-chain (octahedral bulk recursion) in cycle 2 of the Coxeter spectrum. Their explicit

closed-form ratios are

$$R_v = (\sqrt{2})^{k_v} \cdot \big(1-\alpha_v\big), \tag{17}$$

$$R_H = \frac{(\sqrt{2})^{k_v}}{2} \cdot \big(1+\delta_H\big) \cdot \big(1-\alpha_H\big), \tag{18}$$

$$R_W = (\sqrt{2})^{34.5} \cdot \big(1+\delta_W\big) \cdot \big(1-\alpha_W\big), \tag{19}$$

$$m_Z = m_W/\cos\theta_W \quad \text{(SM tree-level)}, \tag{20}$$

where the Higgs-vev Coxeter position is

$$\begin{aligned} k_v &= h + \text{rank} - \text{rank}/h \\ &= 30 + 8 - 8/30 = \frac{1132}{30} \approx 37.733. \end{aligned} \tag{21}$$

The Higgs vev and Higgs boson share the same $\sqrt{2}$-chain primary-harmonic exponent k_v; the Higgs boson's amplitude is reduced by a factor of 2 (the breathing-mode amplitude relative to the substrate baseline tension), and acquires the $(1+\delta_H)$ dark-sector leakage which the vev does not. The W boson sits at the half-integer cycle-2 position $k_W = 34.5$ on the $\sqrt{2}$-chain with leakage $(1+\delta_W)$ suppressed by an E_8 exponent factor of 19.

The boson-sector structural corrections are

$$\alpha_v = \frac{c_2^{(v)}\alpha_\tau^2}{1+\alpha_\tau/(2\varphi)}, \quad c_2^{(v)} = \frac{\text{rank}\cdot h}{m_{\max}} = \frac{|\Phi_{E_8}|}{m_{\max}} = \frac{240}{29}, \tag{22}$$

$$\alpha_H = \frac{c_2^{(H)}\alpha_\tau^2}{1+\alpha_\tau/h_{SU(3)}}, \quad c_2^{(H)} = \frac{(\text{rank}-2)\cdot h}{m_{\max}} = \frac{180}{29}, \tag{23}$$

$$\alpha_W = \frac{c_2^{(W)}\alpha_\tau^2}{1+\alpha_\tau/h_{SU(3)}}, \quad c_2^{(W)} = \frac{2}{e_2} = \frac{2}{7}, \tag{24}$$

$$\delta_H = \frac{1}{h+m_{\max}} = \frac{1}{59}, \tag{25}$$

$$\delta_W = \frac{\delta_H}{e_6} = \frac{1}{(h+m_{\max})\cdot 19} = \frac{1}{1121}, \tag{26}$$

where $e_6 = 19$ (1-indexed in the ordered E_8 exponent list $\{e_1,\ldots,e_8\} = \{1,7,11,13,17,19,23,29\}$) is the partner of the muon's E_8 exponent $e_3 = 11$ under the standard Coxeter duality $e_j + e_{\text{rank}+1-j} = h$ [6]. For E_8, this duality groups the exponents into the four palindromic pairs $\{(1,29),(7,23),(11,19),(13,17)\}$, all summing to $h = 30$ and centered on $h/2 = 15$; the W-boson construction selects the pair $(e_3, e_6) = (11, 19)$ containing the muon's chain position.

The cycle-2 sub-harmonic series $\alpha_v, \alpha_H, \alpha_W$ (and α_t in the quark sector, Eq. (16)) share a common substrate-physical structure: each is a closed-form geometric-series sum of the form $c_2\alpha_\tau^2/(1+\alpha_\tau/X)$, with mode-specific leading coefficient c_2 built from E_8 Coxeter slots (the eight Coxeter exponents) and the topological subgroup ratio $h/h_{SU(3)} = 10$, and closure denominator X from substrate- physical structural integers. The Higgs vev's closure denominator is 2φ (the icosahedral spin-doublet primary-harmonic step: McKay multiplicity 2 times the icosahedral recursion factor φ), reflecting the vev's role as the substrate baseline tension whose self-coupling is its own primary-harmonic step. The Higgs boson, W boson, and top quark closures are all $h_{SU(3)} = 3$ (the colour-projected closure), reflecting these modes' role as fluctuations on the colour-coupled cycle-2 vacuum. The leading coefficients are $c_2^{(v)} = \text{rank}\cdot h/m_{\max}$ (full Cartan rank engages for the vev), $c_2^{(H)} = (\text{rank}-2)\cdot h/m_{\max}$ (two Cartan directions lifted out for the breathing mode), $c_2^{(W)} = 2/e_2$ (vector multiplicity 2 over the substrate's second Coxeter slot), and $c_2^{(t)} = (h/h_{SU(3)})^2$ (topological subgroup ratio squared for the only cycle-2 colour-charged fermion). Each correction is built from framework-fixed E_8 invariants $\{\dim, h, \text{rank}, m_{\max}, e_j\}$, the colour Coxeter number $h_{SU(3)} = 3$, and the gateway-mixing parameter α_τ. We note that the Higgs-vev coefficient $c_2^{(v)} = \text{rank}\cdot h$ equals the order of the E_8 root system $|\Phi_{E_8}| = h\cdot\text{rank} = 30\cdot 8 = 240$, a standard Lie-algebraic identity for any irreducible finite Coxeter group [6]; the substitution $\text{rank}\cdot h \to |\Phi_{E_8}|$ is the natural Lie-theoretic reading of α_v. The rank-and-rank-minus-two structures in α_v, α_H reflect a rank-locked-subspace projection shared with the dark-matter tower (Sec. X). *We use "rank-locked subspace" as framework-specific terminology*: it denotes the rank-rank = 8 invariant subspace that arises in the framework's bulk decomposition dim = dim −rank visible ⊕ rank-locked, distinct from any standard Cartan subalgebra construction. The phrase is defined operationally by its appearance in α_v, α_H, the dark-matter tower's $1-\text{rank}\,\alpha_\tau$ corrections, and certain quark subharmonics; a direct identification with a specific subspace of the Calabi–Yau cohomology is left for future work. The structural origin of δ_H and δ_W is motivated in the subsections below as a minimal closure rule compatible with the same geometric data.

The Higgs vev sits at cycle-2 chain position $k_v = 1132/30 \approx 37.733$, giving $v^{\text{pred}} = 246.21965$ GeV (0.0001 ppm).

a. Closed-form structural identity for the vev. The chain-position prediction for the Higgs vev satisfies the substrate-physical identity at two levels of precision.

Leading identity (∼ 29 ppb residual):

$$\boxed{v\cdot(1-\delta_W) = \dim(E_8) - 2 = 246 \ \text{(in GeV units)}} \tag{27}$$

to numerical residual ∼ 29 ppb (at the framework's gateway-mixing-tower depth-3, α_τ^3, scale). Equivalently,

$$\begin{aligned} v &= \frac{\dim(E_8)-2}{1-\delta_W} = \frac{246\cdot 1121}{1120} \\ &= \frac{275766}{1120}\ \text{GeV} \approx 246.21964\ \text{GeV}, \end{aligned} \tag{28}$$

with $\delta_W = 1/(h+m_{\max})/e_6 = 1/1121$ the W boson's vector dark-sector leakage.

Extended identity (∼ 0.1 ppb residual, at floating-point

precision):

$$v \cdot (1-\delta_W) \;=\; (\dim(E_8)-2)\cdot\left(1+\frac{5}{8}\alpha_\tau^3\right), \qquad (29)$$

with the additional gateway-mixing-tower depth-3 correction $\frac{5}{8}\alpha_\tau^3 \approx 29.1$ ppb. The leading factor 5/8 is the substrate-physical structural ratio

$$\begin{aligned}\frac{5}{8} \;&=\; \frac{\text{McKay icosahedral integer}}{\text{rank}(E_8)}\\ &=\; \frac{\text{rank}(E_8)-h_{SU(3)}}{\text{rank}(E_8)} \;=\; \frac{8-3}{8},\end{aligned} \qquad (30)$$

"rank not consumed by colour, normalised by rank." The McKay icosahedral integer 5 that appears here is the same integer appearing in the gateway formula Eq. (6) as the icosahedral-seed factor.

Full structural ladder including dark-tower contributions. The substrate-physically complete form of the vev structural identity is a four-factor ladder:

$$v \;=\; (\dim -2)\cdot\frac{1}{1-\delta_W}\cdot\left(1+\frac{5}{8}\alpha_\tau^3\right)\cdot\big(1+\zeta_{\rm DM}\big), \quad (31)$$

with the dark-matter tower's kinetic-mixing summed contribution

$$\zeta_{\rm DM} \;=\; \sum_{n=5}^{11}\epsilon_n \;=\; \sum_{n=5}^{11}\alpha_\tau^n \;=\; \alpha_\tau^5\cdot\frac{1-\alpha_\tau^7}{1-\alpha_\tau} \;\approx\; 6.04\times10^{-13}. \qquad (32)$$

Each factor in the ladder has clear substrate-physical content and operates at a distinct scale: $(\dim -2) \approx 246$ GeV is the gateway-saturated dimension count baseline; $1/(1-\delta_W) \approx 1 + 893$ ppm is the macro-mode boundary coupling through the W-vector channel; $\frac{5}{8}\alpha_\tau^3 \approx 29$ ppb is the gateway-mixing-tower depth-3 correction; $\zeta_{\rm DM} \approx 0.6$ ppt is the dark-matter tower kinetic-mixing summed contribution. Each factor corresponds to one of the framework's three substrate-physical dark-sector coupling families: macro-mode boundary couplings (δ_W family), gateway-mixing-tower sub-harmonics (α_τ^n for $n \geq 2$), and dark-matter tower kinetic mixings ($\epsilon_n = \alpha_\tau^n$ for $n \geq 5$). The dark-tower factor $\zeta_{\rm DM}$ is structurally present and complete in the equation, but its magnitude is below PDG's precision floor for v (500 ppb) and below the framework's intrinsic comparison floor against PDG; including it does not improve the empirical match beyond the extended identity Eq. (29), but it represents the substrate-physically complete form. The same structural integer δ_W that governs the W boson's mass shift in Eq. (19) also governs the macro-mode elevation of the substrate's electroweak baseline: *one δ_W, two geometric contexts*, in the same pattern as the gateway formula's $(\dim -2)/\dim$ and the Weinberg angle's 2/9 each express the substrate gateway in different geometric slicings of the gauge phase-space.

Live anchor versus structural baseline. A direct empirical test against the high-precision electron and muon masses establishes that the substrate's mass spectrum is anchored to $v_{\rm observed}$, not to a counterfactual $v_{\rm natural} = \dim -2$. Anchoring the spectrum to $\dim -2$ predicts both leptons at -893 ppm against PDG (a $\sim 10^6\sigma$ deviation on the electron); anchoring to $v_{\rm observed}$ as in Eq. (28) predicts both within ~ 0.03 ppm. The substrate's chain-position couplings therefore reference the effective electroweak vacuum that has the dark-tower content already operating, not a counterfactual baseline that would exist without it. The structural identity Eq. (27) *expresses* the live anchor $v_{\rm observed}$ in closed form by factoring out the W-channel dark-tower elevation; it is not a separate, deeper anchor that the mass spectrum sees.

The Higgs vev is therefore the substrate's electroweak-baseline *macro-mode*: a cosmological-scale tension structure whose value in GeV units is fixed by the gateway-saturated dimension count of $E_8 \times E_8$ on Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ combined with the W-channel dark-tower elevation that lifts it to the observed value. The framework's chain-position formula in Eq. (17) reproduces this identity through different but equivalent structural arithmetic, with numerical residual at the framework's typical sub-ppm precision floor.

B. Origin of the dark-sector coupling family

Particles whose Coxeter positions extend into cycle 2 of the E_8 spectrum (positions $k > h$) couple to the dark E_8 sector through the upper edge of the cycle-1 Coxeter spectrum. Substrate continuity across the dark/visible boundary forces a structural correction to their visible-sector mass formulas, equal to the inverse of the total Coxeter capacity available at the edge.

The structural "boundary capacity" of the E_8 Coxeter spectrum is the sum of the Coxeter number h and the largest exponent $m_{\rm max}$:

$$h + m_{\rm max} = 30 + 29 = 59. \qquad (33)$$

This quantity uses two genuine E_8 invariants: h is the order of the Coxeter element (the period of the spectrum) and $m_{\rm max}$ is the highest non-trivial E_8 exponent (the upper exponent in the palindromic pair $(1, 29)$). Their sum is then a framework measure of the cycle-1 edge plus maximal-exponent displacement, not a standard Lie-algebra invariant in its own right.

C. The Higgs leak: structural closure rule for δ_H

The Higgs boson is the breathing-mode excitation of the substrate tension at $k_v - 2$ (one "inversion-step" below the vev). As a scalar mode-pattern, it couples isotropically to the boundary capacity at first order. A minimal scalar closure consistent with that boundary pic-

ture assigns the leading dark-sector correction

$$\boxed{\delta_H \;=\; \frac{1}{h+m_{\max}} \;=\; \frac{1}{59}} \tag{34}$$

($\approx 1.695\,\%$). The integer 1 in the numerator is the framework's minimal choice: a scalar mode has no internal index structure (no rank or exponent dependence) that would naturally supply an additional suppression factor. The factor $1/(h+m_{\max})$ then enters the Higgs-boson mass formula as the "boundary leakage" multiplicative correction at the breathing mode position $k_v - 2$, yielding $m_H^{\text{pred}} = 125.19978$ GeV (1.75 ppm from PDG). This should be read as a structurally motivated closure rule awaiting derivation from the full compactification data, not as a completed CY3 cohomology computation.

D. The W leak: structural closure rule for δ_W

The W boson is a vector mode-pattern at cycle-2 Coxeter position $k_W = 34.5$, which lies on the $\sqrt{2}$ half-integer chain shared with the heavy-quark sector. As a vector mode, the W carries an internal index structure: its phase-alignment under $SU(2)$ projects onto a specific palindromic exponent pair on the φ-chain, namely the muon-paired pair $(11, 19)$ with $11 + 19 = h = 30$.

The corresponding vector closure rule suppresses the dark-sector boundary leakage for the W by the partner E_8 exponent across the palindromic pairing, $e_n = 19$ (the partner of 11, the muon's Coxeter position):

$$\boxed{\delta_W \;=\; \frac{\delta_H}{e_n} \;=\; \frac{1}{(h+m_{\max})\cdot e_n} \;=\; \frac{1}{59\cdot 19} \;=\; \frac{1}{1121}} \tag{35}$$

(≈ 892 ppm). Each factor in the denominator is fixed by E_8 structure once this closure pattern is adopted:

- $h = 30$ — Coxeter number (algebra-fixed).
- $m_{\max} = 29$ — largest E_8 exponent (algebra-fixed).
- $e_n = 19$ — partner exponent of the muon's position $k_\mu = 11$ under the palindromic pairing $m \leftrightarrow h - m$ centered on $h/2 = 15$ (algebra-fixed once the muon's Coxeter step is identified).

The ratio $\delta_W/\delta_H = 1/e_n = 1/19$ is therefore independently fixed by the E_8 exponent structure with no fitted continuous parameter; the same palindromic-pair structure that governs the muon-tau lepton mass relations governs the suppression of the W boson's dark-sector coupling relative to the Higgs's. Applying δ_W to the cycle-2 W position yields $m_W^{\text{pred}} = 80.36950$ GeV (3.7 ppm from PDG). As with δ_H, this is best read as a framework closure rule sharpened by the E_8 exponent structure rather than as a fully first-principles compactification derivation.

E. Generalization to other cycle-2 modes

The same closure family extends to any cycle-2 Coxeter mode: a scalar in cycle 2 receives the bare δ_H correction; a vector or spinor in cycle 2 receives the suppressed correction $\delta_X = \delta_H/e_X$ where e_X is the E_8 exponent determined by the mode's internal index structure. The dark-sector coupling family is therefore graded by the E_8 exponent set $\{1, 7, 11, 13, 17, 19, 23, 29\}$, producing a discrete spectrum of structurally-fixed corrections rather than a continuous parameter. The Higgs self-coupling $\lambda_H = m_H^2/(2v^2)$ follows from the predicted m_H and v by the standard SM tree-level relation; its closed form and numerical value are presented in Sec. VI.

F. Term-by-term decomposition: electroweak bosons

The electroweak boson masses share the same four-component structure as the charged fermions (primary harmonic, sub-harmonic correction, sector modifier, dark-sector leakage), but all three massive bosons sit in cycle 2 of the Coxeter spectrum, so dark-sector leakage corrections are present for all of them.

a. Higgs vacuum expectation value ($k_v = 1132/30$), $R_v \approx 4.78 \times 10^5$.

- *Primary harmonic:* cycle-2 chain contribution at the chain-position assignment $k_v = 1132/30 \approx 37.733$. The denominator $30 = h$ is the Coxeter number; the numerator 1132 is a structural-integer assignment from the heterotic Wilson-line breaking pattern (Ref. [12]; explicit derivation from CY3 cohomology is outside the present scope).
- *Sub-harmonic correction:* $\alpha_v = (\text{rank} \cdot h/m_{\max})\alpha_\tau^2 = (8 \cdot 30/29)\alpha_\tau^2$. The coefficient $\text{rank} \cdot h/m_{\max}$ encodes the rank-locked Cartan-subspace projection at order α_τ^2 (paralleling the muon's order-2 sub-harmonic).
- *Sector modifier:* cycle-2 chain.
- *Dark-sector leakage:* absent for the vev itself, since the Higgs vev is the substrate baseline tension and does not couple to the spectrum's upper boundary as a fluctuation mode would.

b. Higgs boson, $R_H \approx 2.43 \times 10^5$.

- *Primary harmonic:* same $\sqrt{2}$-chain primary $(\sqrt{2})^{k_v}$ as the vev, divided by 2 (the breathing-mode amplitude is half of the substrate baseline tension; the Higgs boson is the fluctuation mode around the vev, not a separate Coxeter position).
- *Sub-harmonic correction:* $\alpha_H = ((\text{rank} - 2)h/m_{\max})\alpha_\tau^2 = (180/29)\alpha_\tau^2$. The $\text{rank} - 2 = 6$ coefficient drops two Cartan generators relative to

the vev correction α_v, reflecting the two fewer rank-locked directions accessible to the breathing mode.

- *Sector modifier:* cycle-2 $\sqrt{2}$-chain.
- *Dark-sector leakage:* $(1+\delta_H)$ with $\delta_H = 1/(h + m_{\max}) = 1/59$. The Higgs boson, as a scalar, couples isotropically to the spectrum's upper boundary at the boundary-capacity scale $h+m_{\max} = 59$, giving the leading dark-sector leakage with no internal-index suppression. The sign of δ_H is positive (the leakage *enhances* the boson's amplitude, in contrast to the rank-locked correction α_H which reduces it).

c. W boson ($k_W = 34.5$), $R_W \approx 1.56 \times 10^5$.

- *Primary harmonic:* cycle-2 $\sqrt{2}$-chain contribution $(\sqrt{2})^{34.5}$. The half-integer Coxeter position reflects the W's SU(2) vector index structure (rather than scalar or integer-spinor).
- *Sub-harmonic correction:* $\alpha_W = (2/e_2)\alpha_\tau^2/(1 + \alpha_\tau/h_{SU(3)})$, with leading coefficient $2/e_2 = 2/7$ (vector multiplicity 2 over the substrate's second Coxeter slot $e_2 = 7$, the same slot that appears in $\sin^2\theta_{23}^{\rm PMNS} = \text{rank}/(2e_2) = 4/7$ and $\sin^2\theta_{13}^{\rm PMNS} = 2/(e_2e_4) = 2/91$), and closure denominator $h_{SU(3)} = 3$ (colour-projected closure shared with α_H and α_μ).
- *Sector modifier:* cycle-2 $\sqrt{2}$-chain at half-integer position, shared with the heavy-quark $\sqrt{2}$-chain structure.
- *Dark-sector leakage:* $(1+\delta_W)$ with $\delta_W = \delta_H/e_6 = 1/(59\cdot 19) = 1/1121$, where $e_6 = 19$ is the E_8 exponent palindromically paired with the muon's position $k_\mu = 11 = e_3$ (since $11 + 19 = h$). The factor of 19 suppression relative to δ_H encodes the W's internal SU(2) vector index structure projecting onto the $(11, 19)$-paired exponent. The sign of δ_W is positive, like δ_H.

d. Z boson and gauge couplings. The Z boson, g_2, g_1, and λ_H are derived bookkeeping quantities: $m_Z = m_W/\cos\theta_W$, $g_2 = 2m_W/v$, $g_1 = g_2\tan\theta_W$, $\lambda_H = m_H^2/(2v^2)$. These are not independent structural predictions; they inherit precision from m_W, v, m_H, and $\sin^2\theta_W$ via standard SM tree-level relations.

VI. ELECTROWEAK GAUGE SECTOR

A. Weak mixing angle

The weak mixing angle takes the closed-form structural expression

$$\boxed{\sin^2\theta_W = \frac{2}{9} + \frac{\Delta_\varphi}{h} = \frac{2}{9} + \frac{1}{h(h+\varphi^2)},} \tag{36}$$

where $\varphi = (1+\sqrt{5})/2$ is the golden ratio and $\Delta_\varphi = 1/(h+\varphi^2)$. The leading rational term $2/9$ is, in this framework, the $SU(3)_{\rm color}$-induced bulk projection at the unbroken-electroweak-symmetry locus (the geometric weight $2/(h_{SU(3)}\cdot h_{SU(3)}^{\rm extended}) = 2/(3\cdot 3)$ assigned to the visible-sector $SU(2)_L \times U(1)_Y$ subprojection within the bulk symmetry breaking ansatz). We flag this as a *framework structural assignment* rather than a standard derived E_8 identity: it is the parsimonious choice that reproduces the observed value to within the small icosahedral correction below, and the assignment must ultimately be justified at the level of the underlying $E_8\times E_8$ heterotic compactification on $X/(\mathbb{Z}_3\times\mathbb{Z}_3)$. The small correction Δ_φ/h is the golden-ratio-conjugate subharmonic contribution from the visible E_8 embedding, with the same icosahedral provenance as the McKay-correspondence integer 5 in the gateway formula. Numerically, $\sin^2\theta_W = 0.222\overline{2} + 0.001022 = 0.22324$, agreeing with the on-shell PDG value $1 - m_W^2/m_Z^2 = 0.22320$ at 198 ppm.

B. Gauge couplings and Z boson

The $SU(2)_L$ and $U(1)_Y$ gauge couplings follow from the predicted m_W, v, and θ_W via the standard SM tree-level relations:

$$g_2 = \frac{2m_W}{v} = 0.65283, \tag{37}$$

$$g_1 = g_2\tan\theta_W = \frac{2m_W\tan\theta_W}{v} = 0.34998, \tag{38}$$

$$e = g_2\sin\theta_W, \qquad \alpha_{\rm em} = \frac{e^2}{4\pi\epsilon_0\hbar c}. \tag{39}$$

The Z boson mass is

$$\begin{aligned} m_Z &= \frac{m_W}{\cos\theta_W} \\ &= \frac{h_{\rm eff}\cdot k_W\cdot(1+\delta_W)\cdot K_{{\rm cycle},W}}{\cos\theta_W} \\ &= 91.19037\ \text{GeV}, \end{aligned} \tag{40}$$

agreeing with PDG 2024 at 30 ppm.

a. Heterotic-string expectation: gauge unification. A heterotic $E_8\times E_8$ compactification on a Calabi–Yau threefold normally predicts that the running visible-sector gauge couplings $\alpha_1(\mu), \alpha_2(\mu), \alpha_3(\mu)$ unify at a GUT scale set by the compactification volume, $M_{\rm GUT} \sim 10^{16}$ GeV in conventional units, with $\alpha_{\rm GUT} \sim 1/24$. Within the present framework, g_1, g_2, and g_3 are predicted at the Z pole as separate structural objects rooted in distinct E_8 invariants (k_W, $\sin^2\theta_W$, α_s), and the framework does not predict their evolution to high scale: the reconciliation of the predicted low-scale couplings with one-loop unification of the running couplings is left for future work, contingent on a specific identification of the substrate-physical interpretation of RG running. We note that the Z-pole prediction $\alpha_s(M_Z) = 0.1180$

is consistent at the 5% level with unification at $M_{\rm GUT}$ given two-loop SM running and threshold corrections at $M_{\rm SUSY} \sim 1$ TeV; no claim beyond consistency is made.

C. Higgs self-coupling

The Higgs self-coupling is the standard SM tree-level relation inverted to read off λ_H:

$$\begin{aligned}\lambda_H &= \frac{m_H^2}{2v^2} \\ &= \frac{[h_{\rm eff}(k_v-2)(1-\alpha_H)(1+\delta_H)]^2}{2[h_{\rm eff}k_v(1-\alpha_v)]^2} = 0.12928011,\end{aligned} \tag{41}$$

agreeing with the observed $\lambda_H = 0.12928057$ at 3.5 ppm.

D. Strong coupling

The strong coupling at the Z pole takes the closed-form structural expression

$$\boxed{\begin{aligned}\alpha_s(M_Z) &= \frac{m_{\rm max}^2 + {\rm rank}}{m_{\rm max}\cdot {\rm dim}} = \frac{29^2+8}{29\cdot 248} \\ &= \frac{849}{7192} = 0.118048,\end{aligned}} \tag{42}$$

which uses only forced E_8 algebra invariants $\{{\rm dim}, {\rm rank}, m_{\rm max}\}$. The numerator $m_{\rm max}^2 + {\rm rank}$ combines the squared largest exponent $m_{\rm max}^2 = (h-1)^2 = 841$ with the algebra rank rank $= 8$; the denominator $m_{\rm max}\cdot {\rm dim} = 29\cdot 248 = 7192$ is the largest-exponent times the algebra dimension. *We flag this as a framework structural assignment.* Although the numerator and denominator are each built from E_8 invariants, neither $m_{\rm max}^2 + {\rm rank}$ nor $m_{\rm max}\cdot {\rm dim}$ corresponds directly to a standard Lie-algebraic invariant (quadratic Casimir, Killing form, dimension of an irrep, etc.): the quadratic Casimir of the adjoint of E_8 is $C_2({\rm adj}) = 2h^\vee = 60$ (with dual Coxeter number $h^\vee = h$ for simply-laced groups), and no standard identity recovers 849/7192 from it. The construction selects the unique combination of $\{{\rm dim}, {\rm rank}, m_{\rm max}\}$ at a fixed weight that reproduces the PDG value to within experimental precision; a first-principles identification (e.g. via a specific gauge-boundary-condition trace on $X/(\mathbb{Z}_3\times\mathbb{Z}_3)$) is left for future work. Numerically $\alpha_s(M_Z) = 0.118048$, agreeing with the PDG 2024 world average $\alpha_s(M_Z) = 0.1180 \pm 0.0009$ at 405 ppm, well within experimental precision. This is the only closed-form structural expression for $\alpha_s(M_Z)$ in the present framework; no continuous parameter enters.

E. Fine-structure constant

The fine-structure constant follows from g_2 and θ_W via $e = g_2 \sin\theta_W$ and $\alpha_{\rm em} = e^2/(4\pi)$ in natural units. With Eqs. (37) and (36),

$$\alpha_{\rm em} = \frac{(g_2\sin\theta_W)^2}{4\pi} = \frac{(2m_W\sin\theta_W/v)^2}{4\pi}, \tag{43}$$

inheriting precision from m_W, v, and $\sin\theta_W$. The on-shell predicted value matches the PDG within experimental precision (Sec. XIII).

VII. STRONG CP ANGLE

The $\mathbb{Z}_3\times\mathbb{Z}_3$ quotient structure provides a geometric route to a vanishing QCD vacuum angle, but the required cancellation is not derived independently here. The QCD topological term integrates as

$$\int_{X/(\mathbb{Z}_3\times\mathbb{Z}_3)} {\rm Tr}(F\wedge F) = \frac{1}{|\mathbb{Z}_3\times\mathbb{Z}_3|}\int_X {\rm Tr}_g(F\wedge F), \tag{44}$$

where ${\rm Tr}_g$ is twisted by the group action. For the canonical E_8 embedding and Wilson-line data [12, 13], the present framework assumes that the quotient projection annihilates the relevant twisted trace. The mathematical content of this assumption is the equivariant-cohomology identity

$$\big[{\rm ch}_2(V)\cdot[g]\big]_X = 0 \quad \text{for every nontrivial } g\in\mathbb{Z}_3\times\mathbb{Z}_3, \tag{45}$$

where ${\rm ch}_2(V) = \frac{1}{2}{\rm Tr}(F\wedge F)/(2\pi)^2$ is the second Chern character of the gauge bundle V on the cover X and $[g]$ is the cohomology class of the fixed locus of g (empty for a freely-acting $\mathbb{Z}_3\times\mathbb{Z}_3$ on the Schoen threefold, which is the case of interest here). Equivalently, the assumption is that the descent of ${\rm Tr}(F\wedge F)$ to $X/(\mathbb{Z}_3\times\mathbb{Z}_3)$ admits no non-zero g-twisted contribution under the orbifold-cohomology decomposition. Under that assumption, ${\rm Tr}_{\mathbb{Z}_3\times\mathbb{Z}_3}{\rm Tr}(F\wedge F) = 0$, which would set $\theta_{\rm QCD} = 0$. The heterotic construction papers fix the quotient geometry and bundle data on which this statement depends, but they do not themselves prove Eq. (45) for the specific bundle data the framework needs. Within the framework, this therefore functions as a geometric strong-CP mechanism rather than as a separately established theorem. If the cancellation holds, the strong CP problem is removed without Peccei–Quinn [30] or a dynamical axion [29].

VIII. NEUTRINO SECTOR

A. Mass eigenvalues from sub-gateway φ-chain modes

The neutrino mass eigenvalues arise as deeply sub-gateway φ-chain modes at recursive depths $n = 3, 4$ in

α_τ. Their explicit closed-form ratios are

$$R_{\nu_1} = 2\,\varphi^2\,\alpha_\tau^4, \tag{46}$$
$$R_{\nu_3} = 2\,\alpha_\tau^3\,\big(1+5\varphi^2\alpha_\tau\big), \tag{47}$$
$$R_{\nu_2} = R_{\nu_3}\cdot\frac{1+\Delta_\varphi}{\mathrm{rank}-2} = R_{\nu_3}\cdot\frac{1+\Delta_\varphi}{6}, \tag{48}$$

where ν_1 is the lightest mass eigenstate (4th-order icosahedral mode with φ^2 prefactor), ν_3 is the heaviest (3rd-order with sub-leading icosahedral correction $5\varphi^2\alpha_\tau$), and ν_2 relates to ν_3 by the icosahedral-bulk closure $(1+\Delta_\varphi)$ divided by $\mathrm{rank}-2=6$ (the number of Cartan directions accessible to the second-generation neutrino mode). Numerically, $m_{\nu_1}\approx 0.451$ meV, $m_{\nu_2}\approx 8.61$ meV, $m_{\nu_3}\approx 50.21$ meV.

B. Mass-squared splittings and bounds

The mass-squared splittings derived from Eqs. (46)–(48): $\Delta m^2_{21} = 7.4193\times10^{-5}$ eV^2 (90 ppm), $\Delta m^2_{31} = 2.5211\times10^{-3}$ eV^2 (2440 ppm), agreeing with the NuFIT 5.3 global-fit ranges [34, 35] within experimental precision.

Normal mass ordering is selected with $\Sigma m_\nu = 59.3$ meV, KATRIN effective mass $m_\beta = 8.82$ meV, and $m_{\beta\beta} = 4.00$ meV (assuming Majorana nature with vanishing Majorana phases). These remain below current cosmological, direct-kinematic, and neutrinoless-double-beta-decay bounds from DESI, KATRIN, and KamLAND-Zen [36–38].

a. Dirac vs Majorana nature. The framework's mass-eigenvalue construction $(R_{\nu_1}, R_{\nu_2}, R_{\nu_3})$ does not by itself distinguish Dirac from Majorana neutrinos: the closed-form ratios specify the three positive eigenvalues of $|m_\nu|$, not the Lorentz-structure of the mass term. The framework's heterotic $E_8\times E_8$ embedding on $X/(\mathbb{Z}_3\times\mathbb{Z}_3)$ admits both possibilities, with right-handed neutrinos arising as either standalone Dirac partners or as integrated-out heavy seesaw modes producing an effective Majorana mass at low scale. The $m_{\beta\beta} = 4.00$ meV prediction quoted above is the Majorana-case value with vanishing Majorana phases; the corresponding $m_{\beta\beta}$ in the Dirac case is identically zero. A future framework-internal derivation determining which case obtains would convert $m_{\beta\beta}$ from a conditional prediction into a direct falsifiable observable; the present construction commits to the mass-eigenvalue triplet only.

C. PMNS mixing angles and CP phase

The PMNS lepton-mixing parameters take explicit closed forms involving E_8 exponents and the $SU(3)_{\mathrm{color}}$ Coxeter number $h_{SU(3)} = 3$:

$$\sin^2\theta_{12}^{\mathrm{PMNS}} = \frac{h}{h_{SU(3)}}\cdot\Delta_\varphi = \frac{10}{h+\varphi^2}, \tag{49}$$
$$\sin^2\theta_{23}^{\mathrm{PMNS}} = \frac{\mathrm{rank}}{2\,e_2} = \frac{4}{7}, \tag{50}$$
$$\sin^2\theta_{13}^{\mathrm{PMNS}} = \frac{2}{e_2\cdot e_4} = \frac{2}{91}, \tag{51}$$
$$\delta_{\mathrm{CP}} = \frac{e_3\cdot h_{SU(3)}\cdot\pi}{h} = \frac{11\cdot3\cdot\pi}{30} = \frac{11\pi}{10}, \tag{52}$$

where $e_2 = 7, e_3 = 11, e_4 = 13$ are the second, third, and fourth E_8 Coxeter exponents respectively. The CP phase $\delta_{\mathrm{CP}} = 11\pi/10 = 198^\circ$ has numerator equal to the muon's Coxeter step $k_\mu = 11$ $(= e_3)$ times $h_{SU(3)} = 3$, denominator $h = 30$. All four PMNS predictions agree with NuFIT 5.3 within experimental precision. The small-integer combinations of $\{\mathrm{rank}, e_j, h, h_{SU(3)}, \Delta_\varphi\}$ chosen for each PMNS angle are constrained but not uniquely forced by E_8 algebra alone; the framework's structural choices for these combinations are dictated by the requirement that the closed forms match observation, and a first-principles derivation from the underlying CY3 cohomology is left for future work.

IX. QUARK MIXING (CKM)

The CKM matrix elements take explicit closed forms using the same structural integers:

$$\sin^2\theta_{12}^{\mathrm{CKM}} = \frac{h/2-1}{\dim+h} = \frac{14}{278}, \tag{53}$$
$$\sin^2\theta_{23}^{\mathrm{CKM}} = \frac{4\sin^2\theta_{12}^{\mathrm{CKM}}}{e_3^2}\cdot\Big(1+\frac{\Delta_\varphi}{h_{SU(3)}}\Big), \tag{54}$$
$$\sin^2\theta_{13}^{\mathrm{CKM}} = \alpha_\tau^2\cdot\big(1+4\Delta_\varphi\big), \tag{55}$$
$$\delta_{\mathrm{CKM}} = \frac{e_3\cdot\pi}{h} = \frac{11\pi}{30} = 66.0^\circ. \tag{56}$$

The Cabibbo angle uses the chain depth $h/2-1 = M = 14$; the atmospheric mixing involves the muon-exponent $e_3 = 11$ squared in the denominator with an icosahedral-bulk correction $\Delta_\varphi/h_{SU(3)}$; the reactor mixing $\sin^2\theta_{13}^{\mathrm{CKM}} = \alpha_\tau^2(1+4\Delta_\varphi)$ is at order α_τ^2 with a four-fold icosahedral correction. All four CKM predictions agree with PDG 2024 within experimental precision. The leptonic CP phase $\delta_{\mathrm{CP}} = 11\pi/10$ and the quark CP phase $\delta_{\mathrm{CKM}} = 11\pi/30$ share the same E_8-exponent numerator 11, differing only by the factor of $h_{SU(3)} = 3$ that distinguishes the lepton sector from the quark sector. As with the PMNS angles, the specific small-integer combinations chosen for each CKM angle (e.g., the factor of 4 in Eq. (54), the factor of 4 in front of Δ_φ in Eq. (55)) are constrained but not uniquely forced by E_8 algebra; their first-principles derivation from the heterotic CY3 cohomology is left for future work.

X. DARK-MATTER TOWER

A. The unified neutrino-DM Coxeter tower

The framework's Coxeter chain produces a unified neutrino-and-dark-matter tower:

$$\boxed{m_n = 2\, h_{\rm eff}\, \alpha_\tau^n,} \tag{57}$$

with $\alpha_\tau = 1/(\dim + h) = 1/278 \approx 3.5971 \times 10^{-3}$ and the icosahedral spin-doublet prefactor $c_n = 2$ shared with ν_3. The electron-anchor value of $h_{\rm eff}$ is 0.5151534130081301 MeV (the muon-anchor differs by 0.0011 ppm).

Each order n defines a structurally distinct mode with mass m_n, characteristic kinetic-mixing parameter $\epsilon_n \equiv \alpha_\tau^n$, and de Broglie wavelength $\lambda_{\rm dB,n} = h/(m_n v_{\rm DM})$ at the galactic virial velocity $v_{\rm DM} = 220$ km/s $= 7.3384 \times 10^{-4}\, c$.

B. Cosmological cutoffs and viable mass window

Two cosmological cutoffs bracket the viable dark-matter window:

- **Lyman-α cutoff**: ultralight bosonic dark-matter modes with $m \lesssim 2 \times 10^{-20}$ eV are excluded at 95% C.L. by Lyman-α forest constraints [40].
- **Hubble coherence cutoff**: modes lighter than $H_0 \sim 1.4 \times 10^{-33}$ eV cannot oscillate coherently in the matter-dominated era and are believed to behave as dark energy rather than dark matter; the precise role of substrate modes at orders $n \geq 13$ in the present framework, including any contribution to the cosmological constant, will require further research.

These cutoffs partition the tower (Eq. (57)) into three regimes: visible-sector neutrinos at orders $n = 3, 4$; dark matter at orders $n = 5$–11 (seven discrete modes); and excluded modes at $n \geq 12$. Table X lists every order from $n = 3$ to $n = 15$ with its mass, kinetic-mixing parameter, de Broglie wavelength at $v_{\rm DM}$, and physical regime.

C. Order-5 dominant mode: couplings and detection prospects

The dominant dark-matter mode is at $n = 5$:

$$m_5 = 2\, h_{\rm eff}\, \alpha_\tau^5 = 620.5012 \text{ neV}, \tag{58}$$

with structurally-fixed kinetic mixing

$$\epsilon_5 = \alpha_\tau^5 = 6.022 \times 10^{-13} \tag{59}$$

and axion-photon coupling derived from ϵ_5, $\alpha_{\rm em}$, and the predicted Higgs vev:

$$g_{a\gamma\gamma} = \frac{\epsilon_5\, \alpha_{\rm em}}{v} = 1.7849 \times 10^{-17} \text{ GeV}^{-1}. \tag{60}$$

The mass falls in the upper end of the ADMX direct-detection sensitivity window for axion-like and dark-photon dark-matter searches.

D. Geometric-series abundance convergence

The same factor $\alpha_\tau = 1/278$ that produces successive mass levels also fixes the relative cosmological abundance. The order-5 contribution is dominant; orders 6–11 contribute via the convergent geometric series:

$$\zeta_{\rm tower} \equiv \sum_{n=5}^{11} \frac{\Omega_n}{\Omega_5} = \sum_{k=0}^{6} \alpha_\tau^k \approx 1.0638. \tag{61}$$

The convergence is rapid (Table XI); orders 6–11 contribute only $\approx 6.0\%$ of the total dark-matter abundance, and order 5 contributes $\approx 94.0\%$.

E. Decay constant and misalignment angle

The order-5 decay constant and misalignment angle are fixed by structural integers and the predicted electroweak / Planck scales:

$$\begin{aligned} f_5 &= \frac{\text{rank}(h - \text{rank})}{\dim} \cdot \frac{\sqrt{v \cdot m_{\rm Planck}}}{\alpha_\tau} \\ &= \frac{22}{31} \cdot \frac{\sqrt{v\, m_{\rm Planck}}}{\alpha_\tau} \\ &= 1.082 \times 10^{13} \text{ GeV}, \end{aligned} \tag{62}$$

$$\theta_5 = \arctan \varphi \approx 1.0172 \text{ rad} \approx 58.28^\circ. \tag{63}$$

The ratio $\text{rank}(h - \text{rank})/\dim = 8 \cdot 22/248 = 22/31$ is the off-diagonal Coxeter coupling prefactor. The $\sqrt{v \cdot m_{\rm Planck}}$ factor is the geometric mean of the electroweak and Planck scales (the "substrate see-saw amplitude"). The misalignment angle $\theta_5 = \arctan \varphi$ uses the icosahedral golden ratio.

F. Rank-locked sub-harmonic correction

A structural pattern of rank-locked-subspace corrections operates across the gateway-mixing tower (Table XII). The order-5 mode receives the rank-locked correction

$$\begin{aligned} m_5^{\rm corrected} &= 2 h_{\rm eff} \alpha_\tau^5 \cdot (1 - \text{rank}\, \alpha_\tau) \\ &= 2 h_{\rm eff} \alpha_\tau^5 \cdot \frac{270}{278} = 602.645 \text{ neV}, \end{aligned} \tag{64}$$

TABLE X. Unified neutrino–dark-matter Coxeter tower from order $n = 3$ (heaviest visible neutrino) through order $n = 15$ (super-Hubble). Mass: $m_n = 2h_{\rm eff}\alpha_\tau^n$ with $\alpha_\tau = 1/(\dim + h) = 1/278$. Kinetic-mixing parameter: $\epsilon_n = \alpha_\tau^n$. de Broglie wavelength at galactic virial velocity $v_{\rm DM} = 220$ km/s.

n	Prefactor	Mass	$\lambda_{\rm dB}$	ϵ_n	Regime
3	$2\,\alpha_\tau^3$	47.9548 meV	5.6073 mm	$4.654e-08$	(nu_3, nu_2, visible)
4	$2\,\alpha_\tau^4$	172.4993 ueV	1.5588 m	$1.674e-10$	(nu_1, visible)
5	$2\,\alpha_\tau^5$	620.5012 neV	433.3531 m	$6.022e-13$	ALP / dark-photon (ADMX)
6	$2\,\alpha_\tau^6$	2.2320 neV	120.4722 km	$2.166e-15$	ALP / atomic-clock DM
7	$2\,\alpha_\tau^7$	8.0288 peV	33491.2600 km	$7.793e-18$	atomic-clock DM
8	$2\,\alpha_\tau^8$	28.8807 feV	9310570.2763 km	$2.803e-20$	atomic-clock / pulsar-timing
9	$2\,\alpha_\tau^9$	103.8875 aeV	17.3017 AU	$1.008e-22$	pulsar-timing / asteroid
10	$2\,\alpha_\tau^{10}$	373.6961 zeV	4809.8804 AU	$3.627e-25$	wave-DM
11	$2\,\alpha_\tau^{11}$	1.3442 zeV	6.4821 pc	$1.305e-27$	wave-DM, sub-galactic
12	$2\,\alpha_\tau^{12}$	4.8354 yeV	1.8020 kpc	$4.693e-30$	EXCLUDED by Lyman-alpha
13	$2\,\alpha_\tau^{13}$	17.3934 rdeV	500.9615 kpc	$1.688e-32$	EXCLUDED (super-Hubble)
14	$2\,\alpha_\tau^{14}$	0.0626 rdeV	139.2673 Mpc	$6.073e-35$	EXCLUDED (super-Hubble)
15	$2\,\alpha_\tau^{15}$	0.0002 rdeV	38716.3076 Mpc	$2.184e-37$	EXCLUDED (super-Hubble)

TABLE XI. Geometric-series abundance convergence over the viable dark-matter window $n = 5$–11. The cumulative ratio saturates at $\zeta_{\rm tower} \approx 1.0638$.

n	Ω_n/Ω_5	Cumulative $\sum_{n'=5}^{n} \Omega_{n'}/\Omega_5$
5	1	1.000000
6	5.998×10^{-2}	1.059976
7	3.597×10^{-3}	1.063573
8	2.157×10^{-4}	1.063789
9	1.294×10^{-5}	1.063802
10	7.760×10^{-7}	1.063803
11	4.654×10^{-8}	1.063803

a -2.878% shift from the bare tower value $m_5^{\rm bare} = 620.501$ neV. The correction follows the same pattern of small-coefficient Coxeter sub-harmonic shifts that operate on τ, μ, v, and ν_3 (Table XII).

TABLE XII. Rank-locked sub-harmonic correction pattern across the gateway-mixing tower. Each correction takes the form $(1 \pm c_n \alpha_\tau^k)$ where the coefficient c_n and power k are fixed by the mode's structural assignment.

Mode	n	Correction	Coefficient
τ	1	$(1 - \alpha_\tau)$	1
μ	2	$(1 - 2\alpha_\tau^2)$	2
v	2	$(1 - \mathrm{rank}\,\alpha_\tau^2)$	8
ν_3	3	$(1 + 5\varphi^2\alpha_\tau)$	$5\varphi^2 \approx 13.09$
m_5	5	$(1 - \mathrm{rank}\,\alpha_\tau)$	8

G. Term-by-term decomposition: dark-matter tower modes

The dark-matter tower mass formula $m_n = 2h_{\rm eff}\alpha_\tau^n$ (Eq. (57)) admits the same four-component decomposition. Each order n in the tower is a sub-gateway recursive mode (analogous to the leptons but at deeper recursive depth α_τ^n rather than $h^{k/15}$).

a. Generic tower mode ($n \geq 3$).

$$m_n = \underbrace{2}_{\text{prefactor}} \cdot h_{\rm eff} \cdot \underbrace{\alpha_\tau^n}_{\text{primary harmonic at depth } n} \cdot \underbrace{(1 + \mathcal{C}_n)}_{\text{order-specific correction}}$$

- *Icosahedral spin-doublet prefactor:* the constant 2 is the icosahedral spin-doublet multiplicity, shared across all tower modes (including ν_3). Structurally, it encodes the binary icosahedral group's two-element kernel under the McKay correspondence.

- *Primary harmonic:* α_τ^n where $\alpha_\tau = 1/(\dim + h) = 1/278$ is the gateway-mixing parameter. Successive orders n scale by α_τ, generating a discrete geometric tower.

- *Sector modifier:* sub-gateway recursive (α_τ chain), distinct from the visible-sector φ-chain (leptons) and $\sqrt{2}$-chain (quarks).

- *Order-specific correction $\mathcal{C}_n$:* sub-harmonic shift specific to order n. For $n = 3$ (ν_3), $\mathcal{C}_3 = 5\varphi^2\alpha_\tau$ (icosahedral correction). For $n = 5$ (dominant DM), $\mathcal{C}_5 = -\mathrm{rank}\,\alpha_\tau$ (rank-locked correction; see below). For $n = 4, 6, 7, \ldots, 11$, the corrections are sub-leading and do not change the tower's geometric-series structure.

b. Dominant DM mode ($n = 5$).

$$m_5 = 2 \cdot h_{\text{eff}} \cdot \underbrace{\alpha_\tau^5}_{\text{primary harmonic}} \cdot \underbrace{\left(1 - \operatorname{rank} \alpha_\tau\right)}_{\text{rank-locked sub-harmonic}}$$

- *Primary harmonic:* $\alpha_\tau^5 = (1/278)^5 \approx 6.0 \times 10^{-13}$, the fifth recursive depth. Combined with $2h_{\text{eff}}$, this gives the bare value 620.5012 neV.
- *Sub-harmonic correction:* $(1 - \operatorname{rank} \alpha_\tau) = (1 - 8/278) = 270/278 \approx 0.9712$, the rank-locked-subspace projection. The factor rank = 8 here is the same rank that appears at order α_τ^2 for the Higgs vev correction α_v, but here at first order in α_τ rather than second, reflecting the deeper recursive depth $n = 5$.
- *Sector modifier:* sub-gateway α_τ-chain.
- *Dark-sector leakage:* not separately applicable; the sub-gateway recursive structure is itself the substrate-internal coupling channel.

c. Couplings. The kinetic-mixing parameter $\epsilon_n = \alpha_\tau^n$ at each order is identified directly with the mixing parameter in the dark-photon literature; for $n = 5$, $\epsilon_5 = 6.022 \times 10^{-13}$. The axion-photon coupling $g_{a\gamma\gamma} = \epsilon_5 \alpha_{\text{em}}/v$ follows from the substrate cross-term construction (developed in the companion paper [43]) extended to the dark-sector mode.

d. Structural pattern across the gateway-mixing tower. The rank-locked sub-harmonic corrections (Table XII) follow a discrete, mode-specific pattern with coefficients $\{1, 2, \text{rank}, 5\varphi^2, \text{rank}\}$ for the modes $\{\tau, \mu, v, \nu_3, m_5\}$. The repetition of rank = 8 in both α_v (vev, order α_τ^2) and m_5 (DM, order α_τ^1) reflects the rank-locked-subspace projection shared between the cycle-2 boson sector and the deepest sub-gateway DM mode.

H. Cosmological abundance result

Combining the corrected m_5 from Eq. (64), the structurally-derived f_5 from Eq. (62), the icosahedral misalignment angle θ_5 from Eq. (63), and the geometric-series tower-sum factor $\zeta_{\text{tower}} \approx 1.0638$, the standard misalignment-style abundance law [29], translated into the substrate variables used here, gives the consistency estimate

$$\boxed{\Omega_{\text{DM}} h^2 = 0.119982,} \tag{65}$$

consistent with the Planck 2018 measured value $\Omega_{\text{DM}} h^2 = 0.1200 \pm 0.001$ [39] at 0.015 %, without introducing any additional fitted cosmological parameter. The inputs $\{f_5, \theta_5, m_5, \zeta_{\text{tower}}\}$ are all derived from $E_8 \times E_8$ structure plus the same single electron-mass anchor that produces the Standard-Model spectrum, so this match represents a structural consistency check rather than an independent first-principles cosmological derivation.

We note that the imported misalignment-style abundance formula carries $O(1)$ order-of-magnitude uncertainty in its application across mass regimes; the precise numerical agreement at 0.015 % should be interpreted within that uncertainty as "order-of-magnitude consistent with Planck." A full substrate-native cosmological evolution computation (Klein–Gordon evolution through radiation–matter equality at m_5) is left for follow-up work; only after that step would the abundance value qualify as a first-principles prediction in the same sense as the mass-spectrum relations.

XI. GRAVITATIONAL SECTOR

Newton's constant from Eq. (6):

$$G^{\text{pred}} = 6.67431691 \times 10^{-11}\ \text{m}^3\text{kg}^{-1}\text{s}^{-2}, \tag{66}$$

agreeing with CODATA 2022 [32] at 2.5 ppm, below the G-measurement scatter that still characterizes direct laboratory determinations [33]. Planck mass: $m_{\text{Planck}}^{\text{pred}} = 1.22088858 \times 10^{19}$ GeV (1.2 ppm).

XII. TWO-ANCHOR DIAGNOSTIC CONSISTENCY CHECK

The framework's predictive path computes $h_{\text{eff}} = v_C/R_v$ from the structural Higgs-vev identity (Sec. I); no PDG mass enters. As a diagnostic check on the framework's internal arithmetic, we additionally evaluate two fermion-anchored modes (anchoring h_{eff} on the CODATA electron mass or on the CODATA muon mass) and compare their predictions across the entire 26-parameter spectrum. The structural-vev mode is the primary prediction; the two fermion-anchored modes are diagnostic only. Cross-checking the two fermion-anchored modes against each other gives a single consistency residual

$$\Delta_{e\mu} = \frac{F_\mu/F_e}{m_\mu/m_e} - 1 = -2.98 \times 10^{-8} = -30\ \text{ppb}, \tag{67}$$

constant across all 26 SM observables. The residual is set by sub-harmonic-series truncation, not structural ambiguity. For derived quantities involving $1/m_{\text{Planck}}^2$, the residual doubles: G shows two-anchor variance $2|\Delta_{e\mu}| = 60$ ppb, observed.

XIII. RESULTS

a. Scheme and renormalization conventions. The structural mass predictions in this paper are matched against PDG values in the schemes in which those values are reported, and *the schemes mix across rows.* We

adopt the following explicit conventions, row by row, in Table XIII.

- *Charged leptons* (m_e, m_μ, m_τ): pole/physical masses (the QED self-energy is infrared-finite for charged leptons so this coincides with the direct experimental measurement).
- *Light quarks* (m_u, m_d, m_s): $\overline{\text{MS}}$ at $\mu = 2$ GeV.
- *Heavy quarks* (m_c, m_b): $\overline{\text{MS}}$ at $\mu = m_q$ (i.e. $\overline{\text{MS}}(m_c)$ and $\overline{\text{MS}}(m_b)$).
- *Top quark* (m_t): the world-average direct (kinematic) mass quoted by PDG, which is closest to the *pole-mass* convention, *not* $\overline{\text{MS}}(m_t)$. The latter is ~ 163 GeV, well below the framework's prediction ~ 172.8 GeV; the framework's match is to the direct/pole value 172.57 GeV.
- *Gauge bosons* (m_W, m_Z, m_H): pole masses.
- *Higgs vev* (v): tree-level minimum of the Higgs potential at the Z-pole scale.
- *Weak mixing angle*: on-shell scheme, $\sin^2\theta_W^{\text{OS}} \equiv 1 - m_W^2/m_Z^2$.
- *Strong coupling*: $\alpha_s^{\overline{\text{MS}}}(\mu = M_Z)$.
- *Fine-structure constant*: Thomson-limit $\alpha_{\text{em}}(0) = 1/137.036$.

Mixed-scheme caveat. Because the framework's predictions match against pole/physical masses for some particles (leptons, top, gauge bosons, Higgs), $\overline{\text{MS}}$ at fixed external scale for others (light and heavy quarks), the on-shell scheme for $\sin^2\theta_W$, the $\overline{\text{MS}}(M_Z)$ scheme for α_s, and the Thomson-limit value for α_{em}, the ppm-level agreements reported in Table XIII should be read as agreement with PDG values *within each row's stated scheme* and *not* as agreement of a single uniformly-defined object across rows. In particular, switching m_t to $\overline{\text{MS}}(m_t)$, $\sin^2\theta_W$ to $\overline{\text{MS}}(M_Z)$, or α_{em} to $\alpha_{\text{em}}(M_Z)$ would each shift the comparison by amounts $O(1\text{–}10\%)$, far above the ppm headline. The framework's structural sub-harmonic corrections ($\alpha_\tau, \alpha_v, \alpha_H, \delta_H, \delta_W$, the per-quark χ_q, λ_q signs, etc.) are identified with the all-orders effective corrections in each row's scheme: they absorb what would otherwise be separately computed perturbative loop corrections in a Lagrangian derivation. The framework therefore predicts only the final scheme-specified value, and the schemes are different for different rows by physics convention rather than by choice. A first-principles Lagrangian derivation in which the structural corrections are matched against explicit one- and two-loop calculations in a *single* unified scheme is left for future work; that derivation would convert mixed-scheme ppm-level agreement into single-scheme ppm-level agreement.

b. Falsification implication. Future PDG updates that move any single row's observable outside the framework's prediction at the ppm level *in that row's own stated scheme* constitute direct falsification. Cross-row re-expression in a unified scheme is not committed to and would require the future Lagrangian-level derivation noted above.

Table XIII summarizes the full prediction set. Primary observables are listed together with four algebraically derived consistency checks that are not counted as independent experimental tests.

TABLE XIII: Substrate-framework predictions for all derivable Standard
puted at mp.prec = 500 from three independent anchor scales: CODATA
the directly-computed structural Higgs vev $v_C = (\dim - 2)/(1-\delta_W)\cdot(1+5$
(15-digit precision) with the deviation against the observed value beneath,
($|\text{dev}| \le 1\sigma_{\text{obs}}$), • orange ($1\sigma < |\text{dev}| \le 3\sigma$), × red ($|\text{dev}| > 3\sigma$), ? blue
on upstream prediction work that is not yet closed, so the deviation is n
but the framework cannot yet decide whether this candidate slot is ph
dark-mode candidates). Observed column carries the published experim
each value where available. Anchor variance (max pairwise $\sim$ 158 ppt
catalogue, carried by the e-anchor-vs-vev-anchor pair; e-anchor-vs-μ-anch
hidden.

Observable	Observed	electron-anchor
CHARGED LEPTON MASSES		
electron (MeV)	0.5109989506900 ±300.000 ppt	0.5109989506900 ✓ ∼ 0 (FP floor)
muon (MeV)	105.6583755000 ±22.000 ppb	105.6583755210 ✓ + 198.492 ppt
tau (MeV)	1776.860000000 ±68.000 ppm	1776.859739202 ✓ − 146.775 ppb
QUARK MASSES		
up (MeV)	2.160000000000 ±170000.000 ppm	2.172135018847 ✓ + 5618.064 ppm
down (MeV)	4.670000000000 ±70000.000 ppm	4.666472668646 ✓ − 755.317 ppm
strange (MeV)	93.40000000000 ±92000.000 ppm	93.25272877140 ✓ − 1576.780 ppm
charm (MeV)	1273.000000000 ±3600.000 ppm	1274.297660854 ✓ + 1019.372 ppm
bottom (MeV)	4183.000000000 ±1700.000 ppm	4178.924487169 ✓ − 974.304 ppm
top (GeV)	172.5700000000 ±1700.000 ppm	172.5698833093 ✓ − 676.193 ppb
NEUTRINO MASSES		
neutrino 1 (meV)	4.510000000000e-10 ±50000.000 ppm	4.516090891872e-10 ✓ + 1350.530 ppm
neutrino 2 (meV)	8.613900000000e-9 ±50000.000 ppm	8.625379530455e-9 ✓ + 1332.675 ppm
neutrino 3 (meV)	5.014980000000e-8 ±50000.000 ppm	5.021285708457e-8 ✓ + 1257.375 ppm
ELECTROWEAK SECTOR		
Higgs vev (GeV)	246.2196500000 ±24.000 ppb	246.2196500308 ✓ + 125.130 ppt
Higgs boson (GeV)	125.2000000000 ±880.000 ppm	125.1997813598 ✓ − 1.746 ppm
W boson (GeV)	80.36920000000 ±160.000 ppm	80.36920038710 ✓ + 4.817 ppb
Z boson (GeV)	91.18760000000 ±23.000 ppm	91.19003482663 ? + 26.701 ppm
Weinberg mixing fraction	0.2232000000000 ±140.000 ppm	0.2232441517887 ? + 197.813 ppm
SU(2) coupling	0.6533000000000 ±1500.000 ppm	0.6528252345176 ? − 726.719 ppm
$U(1)_Y$ coupling (EW norm)	0.3500000000000 ±1000.000 ppm	0.3499811975896 ? − 53.721 ppm
$U(1)_Y$ coupling (GUT norm)	0.4519000000000 ±1000.000 ppm	0.4518237832501 ? − 168.658 ppm

Table XIII continued

Observable	Observed	electron-anchor	muon-anchor	verification
SU(2) fine-structure	0.03394000000000 ±1500.000 ppm	0.03391438943683 ? − 754.583 ppm	0.03391438943683 ? − 754.583 ppm	0.03390103880468322 ? − 1154.663 ppm
$U(1)_Y$ fine-structure (EW)	0.009745000000000 ±2000.000 ppm	0.009747192918717 ? + 225.030 ppm	0.009747192918717 ? + 225.030 ppm	0.009747192918717 ? + 225.030 ppm
$U(1)_Y$ fine-structure (GUT)	0.01624532153119 ±2000.000 ppm	0.01624532153119 ? + 266.088 ppm	0.01624532153119 ? + 266.088 ppm	0.01624532153119 ±1.700 ppt
EM fine-structure (tree)	0.007575408000000 ±1.000 ppb	0.007571189103257 ? − 556.920 ppm	0.007571189103257 ? − 556.920 ppm	0.007571189103257 ±460.000 ppb
LEPTON YUKAWA COUPLINGS				
electron Yukawa	2.935028322980e-6 ±300.000 ppt	2.935028322613e-6 ✓ − 125.130 ppt	2.935028322613e-6 ✓ − 125.130 ppt	2.935028322613e-6 ✓ − 125.130 ppt
muon Yukawa	0.0006068707660433 ±22.000 ppb	0.0006068707660878 ✓ + 73.361 ppt	0.0006068707660878 ✓ + 73.361 ppt	0.0006068707660878 ✓ + 73.361 ppt
tau Yukawa	0.01020576347354 ±68.000 ppm	0.01020576197431 ✓ − 146.900 ppb	0.01020576197431 ✓ − 146.900 ppb	0.01020576197431 ✓ − 146.900 ppb
QUARK YUKAWA COUPLINGS				
up Yukawa	1.240640742819e-5 ±170000.000 ppm	1.247610742106e-5 ✓ + 5618.064 ppm	1.247610742106e-5 ✓ + 5618.064 ppm	1.247610742106e-5 ✓ + 5618.064 ppm
down Yukawa	2.682311235631e-5 ±70000.000 ppm	2.680285239466e-5 ✓ − 755.317 ppm	2.680285239466e-5 ✓ − 755.317 ppm	2.680285239466e-5 ✓ − 755.317 ppm
strange Yukawa	0.0005364622471263 ±92000.000 ppm	0.0005356163642517 ✓ − 1576.780 ppm	0.0005356163642517 ✓ − 1576.780 ppm	0.0005356163642517 ✓ − 1576.780 ppm
charm Yukawa	0.007311739192631 ±3600.000 ppm	0.007319192575631 ✓ + 1019.372 ppm	0.007319192575631 ✓ + 1019.372 ppm	0.007319192575631 ✓ + 1019.372 ppm
bottom Yukawa	0.02402592697783 ±1700.000 ppm	0.02400251842267 ✓ − 974.304 ppm	0.02400251842267 ✓ − 974.304 ppm	0.02400251842267 ✓ − 974.304 ppm
top Yukawa	0.9911915416122 ±1700.000 ppm	0.9911908712511 ✓ − 676.318 ppm	0.9911908712511 ✓ − 676.318 ppm	0.9911908712511 ✓ − 676.318 ppm
HIGGS SELF-COUPLING				
Higgs self-coupling	0.1292805654113 ±1800.000 ppm	0.1292801138469 ✓ − 3.493 ppm	0.1292801138469 ✓ − 3.493 ppm	0.1292801138469 ±140000.000 ppm
STRONG COUPLING				
strong coupling	0.1180000000000 ±7600.000 ppm	0.1180478309232 ? + 405.347 ppm	0.1180478309232 ? + 405.347 ppm	0.1180478309232 ±100000.000 ppm
PMNS NEUTRINO MIXING				
PMNS solar mixing fraction	0.3030000000000 ±50000.000 ppm	0.3065788699420 ? + 11811.452 ppm	0.3065788699420 ? + 11811.452 ppm	0.3065788699420 ? + 11811.452 ppm
PMNS atmospheric mixing fraction	0.5720000000000 ±30000.000 ppm	0.5714285714286 ? − 999.001 ppm	0.5714285714286 ? − 999.001 ppm	0.5714285714286 ? − 999.001 ppm
PMNS reactor mixing fraction	0.02203000000000 ±30000.000 ppm	0.02197802197802 ? − 2359.421 ppm	0.02197802197802 ? − 2359.421 ppm	0.02197802197802 ? − 2359.421 ppm
PMNS CP phase (rad)	3.438298626429 ±50000.000 ppm	3.455751918949 ? + 5076.142 ppm	3.455751918949 ? + 5076.142 ppm	3.455751918949 ? + 5076.142 ppm
CKM QUARK MIXING				
Cabibbo mixing fraction	0.05064000000000 ±4000.000 ppm	0.05035971223022 ? − 5534.909 ppm	0.05035971223022 ? − 5534.909 ppm	0.05035971223022 ? − 5534.909 ppm
CKM atmospheric mixing fraction	0.001682000000000 ±15000.000 ppm	0.001681796704019 ? − 120.812 ppm	0.001681796704019 ? − 120.812 ppm	0.001681796704019 ? − 120.812 ppm
CKM reactor mixing fraction	1.460000000000e-5 ±30000.000 ppm	1.452605387890e-5 ? − 5064.803 ppm	1.452605387890e-5 ? − 5064.803 ppm	1.452605387890e-5 ? − 5064.803 ppm
CKM CP phase (rad)	1.143190660056 ±50000.000 ppm	1.151917306316 ? + 7633.588 ppm	1.151917306316 ? + 7633.588 ppm	1.151917306316 ? + 7633.588 ppm
QCD THETA ANGLE				
QCD theta angle (rad)	0.0	0.0	0.0	0.0
MASSLESS GAUGE BOSONS				
photon (MeV)	0.0	0.0	0.0	0.0
gluon 1 (MeV)	0.0	0.0	0.0	0.0
gluon 2 (MeV)	0.0	0.0	0.0	0.0
gluon 3 (MeV)	0.0	0.0	0.0	0.0
gluon 4 (MeV)	0.0	0.0	0.0	0.0
gluon 5 (MeV)	0.0	0.0	0.0	0.0
gluon 6 (MeV)	0.0	0.0	0.0	0.0
gluon 7 (MeV)	0.0	0.0	0.0	0.0
gluon 8 (MeV)	0.0	0.0	0.0	0.0
GRAVITATIONAL SECTOR				
Newton's G ($m^3kg^{-1}s^{-2}$)	6.674300000000e-11 ±22.000 ppm	6.674316892333e-11 ✓ + 2.531 ppm	6.674316894983e-11 ✓ + 2.531 ppm	6.674316892935e-11 ✓ + 2.531 ppm

Table XIII continued

Observable	Observed	electron-anchor	muon-anchor	verification	Observed	electron-anchor
			Planck mass (GeV)			
			ANOMALOUS MAGNETIC MOMENTS			
			electron anomalous moment			
			muon anomalous moment			
			tau anomalous moment			
			FERMI CONSTANT			
			Fermi constant (GeV)			
			LEPTON LIFETIMES			
			muon lifetime (s)			
			tau lifetime (s)			
			W AND Z DECAY WIDTHS			
			W total decay width (GeV)			
			Z total decay width (GeV)			
			Z invisible width (3 neutrinos) (GeV)			
			Z leptonic width (per lepton) (GeV)			
			HIGGS DECAY WIDTHS			
			Higgs width to bb (GeV)			
			Higgs width to tau tau (GeV)			
			Higgs width to cc (GeV)			
			Higgs total fermionic width (GeV)			
			COMPTON WAVELENGTHS AND ATOMIC CONSTANTS			
			electron Compton wavelength (fm)			
			muon Compton wavelength (fm)			
			Bohr radius (fm)			
			Rydberg energy (eV)			
			EFFECTIVE NEUTRINO SPECIES			
			effective neutrino species			
			DARK MATTER ULTRALIGHT BAND			*(predicted-but-unobserved; false*
			DM ultralight 1 (heaviest) (eV)			? predicted
			DM ultralight 2 (eV)			? predicted
			DM ultralight 3 (eV)			? predicted
			DM ultralight 4 (eV)			? predicted
			DM ultralight 5 (eV)			? predicted
			DM ultralight 6 (eV)			? predicted
			DM ultralight 7 (heaviest) (eV)			? predicted

(continued on next page)

Table XIII continued

Observable	Observed	electron-anchor	muon-anchor	vev-anchor
DARK MATTER MeV–GeV BAND (predicted-but-unobserved; direct-detection targets)				
DM MeV–GeV mode A (MeV)	–	0.8335357318367 ? predicted	0.8335357316712 ? predicted	0.8335357317991 ? predicted
DM MeV–GeV mode B (MeV)	–	5.828295548213 ? predicted	5.828295547056 ? predicted	5.828295547950 ? predicted
DM MeV–GeV mode C (MeV)	–	46.62636438570 ? predicted	46.62636437645 ? predicted	46.62636438360 ? predicted
DM MeV–GeV mode D (GeV)	–	1.492043660342 ? predicted	1.492043660046 ? predicted	1.492043660275 ? predicted
DM MeV–GeV mode E (GeV)	–	4.816169314292 ? predicted	4.816169313336 ? predicted	4.816169314075 ? predicted
DM MeV–GeV mode F (GeV)	–	592.3496672214 ? predicted	592.3496671038 ? predicted	592.3496671946 ? predicted
HALF-SLOT CANDIDATES (spin-1/2 dark-mode candidates; slot occupancy open)				
DM half-slot candidate 1 (MeV)	–	0.6552852639934 ? slot open	0.6552852638534 ? slot open	0.6552852639630 ? slot open
DM half-slot candidate 2 (MeV)	–	2.775832922930 ? slot open	2.775832922379 ? slot open	2.775832922805 ? slot open
DM half-slot candidate 3 (MeV)	–	7.267224939322 ? slot open	7.267224937879 ? slot open	7.267224938994 ? slot open
DM half-slot candidate 4 (MeV)	–	11.75861695571 ? slot open	11.75861695338 ? slot open	11.75861695518 ? slot open
DM half-slot candidate 5 (MeV)	–	30.78445885075 ? slot open	30.78445884464 ? slot open	30.78445884936 ? slot open
DM half-slot candidate 6 (MeV)	–	49.81030074579 ? slot open	49.81030073590 ? slot open	49.81030074354 ? slot open
DM half-slot candidate 7 (MeV)	–	130.4050603423 ? slot open	130.4050603164 ? slot open	130.4050603364 ? slot open
DM half-slot candidate 8 (MeV)	–	552.4047002200 ? slot open	552.4047001104 ? slot open	552.4047001951 ? slot open
DM half-slot candidate 9 (MeV)	–	0.6126241041914 ? slot open	0.6126241041698 ? slot open	0.6126241041638 ? slot open
DM half-slot candidate 10 (MeV)	–	1.732762633568 ? slot open	1.732762633224 ? slot open	1.732762633490 ? slot open
DM half-slot candidate 11 (MeV)	–	3.465525267137 ? slot open	3.465525266449 ? slot open	3.465525266980 ? slot open
DM half-slot candidate 12 (MeV)	–	9.801985667063 ? slot open	9.801985665147 ? slot open	9.801985666620 ? slot open
DM half-slot candidate 13 (MeV)	–	13.86210106855 ? slot open	13.86210106580 ? slot open	13.86210106792 ? slot open
DM half-slot candidate 14 (MeV)	–	27.72420213709 ? slot open	27.72420213159 ? slot open	27.72420213584 ? slot open
RULED OUT BY FRAMEWORK				
4th chiral generation	**ruled out:** Framework predicts exactly three chiral fermion generations. **Falsified** if a 4th chiral lepton or quark generation is observed.			
Fundamental scalar beyond Higgs	**ruled out:** Framework predicts the Higgs is the only fundamental scalar. **Falsified** if a non-Higgs fundamental scalar is observed.			
Quantum-gravity scale below $\sim 10\,E_{\rm Planck}$	**ruled out:** Framework predicts no quantum-gravity scale below $\sim 10\,E_{\rm Planck}$. **Falsified** if quantum-gravity signatures appear at sub-Planckian energies.			
Ultralight DM modes below $\sim 10^{-21}$ eV	**ruled out:** Framework predicts no ultralight-DM modes below $\sim 10^{-21}$ eV (Lyman-α structure-formation cutoff, Rogers & Peiris 2021). **Falsified** if direct detection finds a discrete-spectrum DM mode in this range.			

Table XIII continued

Observable	Observed	electron-anchor	muon-anchor	vev-anchor
PQ axion (strong CP)	**ruled out:** Framework predicts $\theta_{\rm QCD} = 0$ exact $\theta_{\rm QCD}$ is nonzero.			
Dark photon (kinetic-mixing operator)	**ruled out:** Framework predicts no 4D kinetic-m dark-photon kinetic-mixing-style coupling is det			

c. Probe-scale dependence and the deviation pattern.

The deviation column of Table XIII is not flat; it exhibits a structured pattern that tracks the experimental *probe scale* Q at which each PDG value is extracted, and the associated theoretical uncertainty in that extraction. Reading the table from cleanest to coarsest:

- *Charged leptons* ($Q \sim m_\ell$, pure-QED extraction). m_e is the framework anchor (zero by construction); m_μ agrees at 1.1 ppb (the precision of muonium spectroscopy); m_τ at 0.15 ppm (the precision of τ-pair production at e^+e^- thresholds). PDG-quoted experimental uncertainties on these masses are themselves at $\lesssim 10^{-7}, 10^{-7}, 10^{-4}$ relative respectively, and the framework's deviations track them.
- *Z pole electroweak observables* ($Q \sim M_Z$, on-shell or $\overline{\rm MS}$ extractions). m_Z agrees at 30 ppm, m_W at 4 ppm, $\sin^2\theta_W^{\rm OS}$ at 198 ppm, $\alpha_s^{\overline{\rm MS}}(M_Z)$ at 405 ppm. PDG experimental relative uncertainties at this scale are $\sim 10^{-5}, 10^{-4}, 10^{-3}$ respectively; again the framework's deviations track the experimental error budgets.
- *Heavy quarks* ($Q \sim m_q$, $\overline{\rm MS}(m_q)$ or pole/direct extraction). m_c at 1019 ppm, m_b at 974 ppm, m_t at 1293 ppm, all within the $\sim 10^{-3}$ relative PDG uncertainty band, but itself related to lattice-QCD systematics, scheme-conversion ambiguities (especially the renormalon-induced pole-vs-$\overline{\rm MS}$ split for the top), and Coulombic-bound-state corrections at the b and c scales.
- *Light quarks* ($Q \sim 2$ GeV, $\overline{\rm MS}(2\ {\rm GeV})$ from chiral perturbation theory plus lattice). m_u at 5618 ppm, m_d at 755 ppm, m_s at 1577 ppm. Here the PDG-quoted uncertainties are $\sim 10^{-2}$ relative for m_u, m_d (FLAG-averaged lattice results) and $\sim 5\times 10^{-3}$ for m_s. The framework's $\sim 0.5\%$–0.6% deviations on m_u, m_s are within this band; the 0.08% on m_d is well within. The 0.6% on m_u in particular is comparable to FLAG's own $\overline{\rm MS}(2\ {\rm GeV})$ chiral-extrapolation systematic.
- *Mixing angles and CKM/PMNS phases* (Q varying with extraction; PDG uncertainties percent-level). All within the percent-level PDG uncertainties.
- *Cosmological & gravitational sector* ($\Omega_{\rm DM}h^2$ from CMB at horizon scales; G at torsion-balance laboratory scales). Agreement at 152 ppm and 2.5 ppm respectively, within experimental error.

(continued on next page)

TABLE XIV. Per-scheme within-scheme prediction table for the pole/direct scheme. Anchored on the electron mass; predictions for the remaining members computed as $m_X = (R_X/R_{\text{anchor}})\cdot m_{\text{anchor}}$. $\Delta(\sigma_{\text{comb}})$ folds target and propagated-anchor uncertainty in quadrature. * marks the row used as the anchor for this scheme. Anchor: electron (observed = 0.5109989507 MeV, relative $\sigma = 0.000294$ ppm).

Particle	Observed	σ_{obs} (ppm)	Predicted	Δ(pred−obs)	Δ (ppm)	Δ (σ_{comb})
electron*	0.5109989507 MeV	+0.00	0.5109989507 MeV	–	–	–
muon	105.6583755000 MeV	+0.00	105.6583755210 MeV	0.0000000210 MeV	+0.0002	+0.6743
tau	1776.8600000000 MeV	+67.53	1776.8597392015 MeV	-0.0002607985 MeV	-0.1468	-0.0022
top	172.5700000000 GeV	+1738.42	172.5698833093 GeV	-0.0001166907 GeV	-0.6762	-0.0004
W boson	80.3692000000 GeV	+161.75	80.3692003871 GeV	0.0000003871 GeV	+0.0048	+0.0000
Z boson	91.1876000000 GeV	+23.03	91.1900348266 GeV	0.0024348266 GeV	+26.7013	+1.1594
Higgs boson	125.2000000000 GeV	+878.59	125.1997813598 GeV	-0.0002186402 GeV	-1.7463	-0.0020
Higgs vev	246.2196500000 GeV	+24.37	246.2196500308 GeV	0.0000000308 GeV	+0.0001	+0.0000

TABLE XV. Per-scheme within-scheme prediction table for the $\overline{\text{MS}}$ scheme at $\mu = 2$ GeV (light quarks, single common scale – fully within-scheme). * marks the row used as the anchor for this scheme. Anchor: down (observed = 4.6700000000 MeV, relative $\sigma = 7.07e + 04$ ppm).

Particle	Observed	σ_{obs} (ppm)	Predicted	Δ(pred−obs)	Δ (ppm)	Δ (σ_{comb})
up	2.1600000000 MeV	+175925.93	2.1737769100 MeV	0.0137769100 MeV	+6378.1991	+0.0336
down*	4.6700000000 MeV	+70663.81	4.6700000000 MeV	–	–	–
strange	93.4000000000 MeV	+92077.09	93.3232174033 MeV	-0.0767825967 MeV	-822.0835	-0.0071

The framework's deviations therefore correlate with the experimental probe scale and extraction precision rather than with any framework-internal hierarchy: cleaner extractions yield sharper agreement, and noisier extractions yield deviations within their own uncertainty bands. We do *not* claim that this correlation alone establishes the framework — the residual deviations could in principle be hiding a systematic offset of a single multiplicative scale that happens to be inside all PDG error bars — but it does establish that the framework is not *tuned* to one probe scale at the expense of others. Future PDG updates that shrink the experimental uncertainty in any row will tighten the falsification window in that row.

XIV. PREDICTIONS AND FALSIFICATION

The framework makes pre-registered predictions across multiple sectors, falsifiable by:

1. Inverted neutrino mass ordering in future oscillation data.

2. Cosmological convergence to $\Sigma m_\nu > 70$ meV in future CMB and large-scale-structure analyses.

3. δ_{CP} outside $198^\circ \pm 5^\circ$ in future long-baseline oscillation data.

4. $0\nu\beta\beta$ at effective Majorana mass significantly above 4 meV (assuming Majorana nature) in next-generation neutrinoless-double-beta-decay searches.

5. Neutron EDM above the current upper bound, or a confirmed nonzero value above $\sim 10^{-30}\, e \cdot \text{cm}$ in next-generation EDM searches [31].

6. Newton's G convergence inconsistent with $6.67431691 \times 10^{-11}$ at ppm precision.

7. Dark-matter detection at masses inconsistent with the predicted tower (orders 5–11), or coupling significantly different from $g_{a\gamma\gamma} = 1.78 \times 10^{-17}$ GeV^{-1}.

8. Precision EW measurements converging to $\sin^2\theta_W$ inconsistent with $2/9 + \Delta_\varphi/h$ at ppm precision.

XV. OPEN ITEMS AND FURTHER-EXPLORATION PROGRAM

We make this section explicit and visible. The numerical demonstration of the present paper is internally complete in the sense that, given the starting assumptions of Table I, no fitted continuous parameter is introduced and every reported value is a direct structural computation. The items below are open methodological questions that affect either the empirical CY3 identification of Sec. II B or the substrate-physical support of the starting assumptions themselves.

A. Pipeline state and in-flight remediation

We disclose the present state of the CY3 discrimination pipeline explicitly, separating closed results from

TABLE XVI. Per-scheme within-scheme prediction table for the $\overline{\mathrm{MS}}$ scheme at $\mu = m_q$ (heavy quarks, per-particle scale – residual cross-scale within the group). * marks the row used as the anchor for this scheme. Anchor: bottom (observed = 4183.0000000000 MeV, relative $\sigma = 1.67e+03$ ppm).

Particle	Observed	σ_{obs} (ppm)	Predicted	Δ(pred−obs)	Δ (ppm)	Δ (σ_{comb})
charm	1273.0000000000 MeV	+3613.51	1275.5404247476 MeV	2.5404247476 MeV	+1995.6204	+0.5010
bottom*	4183.0000000000 MeV	+1673.44	4183.0000000000 MeV	–	–	–

TABLE XVII. Supplementary derived neutrino-sector and cosmological observables not covered by the auto-generated complete table. Mass-squared splittings, sum of neutrino masses (Σm_ν), kinematic effective neutrino mass (m_β), neutrinoless double-beta effective mass ($m_{\beta\beta}$), and the dark-matter relic-density estimate $\Omega_{\mathrm{DM}}h^2$. e-anchor values are computed from the predicted neutrino masses; μ-anchor values are obtained by linear (m) or quadratic (m^2) rescaling under the new h_{eff} ratio.

Observable	Observed	e-anchor	μ-anchor	Deviation
Δm^2_{21} (eV^2)	$7.41^{+0.21}_{-0.20} \times 10^{-5}$	$7.4193221245 \times 10^{-5}$	$7.4193221075 \times 10^{-5}$	−90 ppm
Δm^2_{31} (eV^2)	$(2.511 \pm 0.027) \times 10^{-3}$	$2.5211270648 \times 10^{-3}$	$2.5211270590 \times 10^{-3}$	+2440 ppm
Σm_ν (meV)	< 72 [36]	59.2898456922	59.2898456244	within
m_β (meV)	< 450 [37]	8.8276108075	8.8276107974	within
$m_{\beta\beta}$ (meV)	< 28–122 [38]	3.9984876767	3.9984876721	within
$\Omega_{\mathrm{DM}}h^2$	0.1200 ± 0.001 [39]	0.119981783712	0.119981783369	−151.8 ppm

indicative-but-unconverged results from in-flight work. The full engineering plan to upgrade the indicative results to converged-and-cross-checked status is the project's `REMEDIATION_PLAN_OPTION_B.md`, which specifies a seven-phase program with 111 positive and negative acceptance tests organised into five CI tiers.

a. Closed results (do not depend on in-flight work).

- *Catalogue uniqueness sweep at Hodge* $(3,3)$, modulo the gCICY non-manifest CAS residue. The Kreuzer–Skarke direct query returns `#NF: 0` at output limits up to $L = 10{,}000$; the $\chi = 0$ self-mirror diagonal sweep $h \in \{1, \ldots, 14\}$ confirms the bin is empty for $h \leq 13$ and populated from $h = 14$. The CICY, Schoen-class, Constantin–Lukas, Borisov–Caldararu Pfaffian, and gCICY-manifest classifications are programmatically checked.

- *Tian–Yau Yukawa-channel structural exclusion* at the searched bidegree range in the bicubic ambient at $h^{1,1} = 2$. No rank-3 SU(3) stable polystable monad bundle on TY/$\mathbb{Z}_3$ in the searched range satisfies the heterotic-SM constraint suite; the leading rank-7 monad survivor $V_{\min 2}$ is excluded by $h^0(V_{\min 2}) = \mathbb{C}^2 \neq 0$.

- *E_8 Coxeter / McKay / closed-form algebraic identities* ($k_v = 1132/30$, $\delta_H = 1/59$, $\delta_W = 1/1121$, $\alpha_\tau = 1/278$, $\omega_{\mathrm{fix}} = 123/248$, $\alpha_s = 849/7192$, the gateway formula $F_e = (\dim -2)/\dim \cdot 5/(\dim \cdot h^M)/(1-1/(4(\dim -1)))$, etc.), all closed in rational arithmetic.

- *Mass-spectrum prediction pipeline* (`substrate_mass_spectrum.py`) producing the residuals reported in Tabs. XIII–**??**.

b. Indicative, not converged.

- *Chain-channel Bayes factor* $\ln \mathrm{BF} = -14.76$. The current chain-channel discriminators score the metric-Laplacian Galerkin residual against the framework's sub-Coxeter chain at $k_{\mathrm{test}} \in \{2, 3, 4, 5\}$. Both `p7_11_quark_chain.json` and `p7_11_lepton_chain.json` report `converged: false` with consecutive-degree relative deltas of order 11–35% on the Schoen-side residual. The figure is reported as supporting evidence; the converged replacement is Phase 2 of the remediation plan.

- *σ-channel matched-basis defense.* The current basis-matching at $n_b = 27$ is implemented as a lex-sorted truncation of TY's 87-dimensional invariant basis to Schoen's 27-dimensional basis. A geometric basis projection (eigenvalue-ranked principal components of the Laplacian, or harmonic-form ranking) is the in-flight Phase 3.3 upgrade; the σ-channel exclusion from the model-comparison BF is conditional on its survival under the upgrade.

- *σ functional definition.* The production code computes a weighted L^2 variance per Donaldson 2009 §3.4; a uniform pin of the L^1 vs L^2 functional choice across code, paper, and methodology table is in-flight Phase 3.1.

c. In-flight work (Option B remediation). The phases below are scheduled per the remediation plan. Each phase has explicit acceptance criteria gating the next.

1. *Phase 1 — pipeline integrity*: anomaly-cancellation regression test for the Schoen monad with $W =$

(27, 27); slope-stability proof on a stated Kähler-cone region; structure-group-of-V consistency (SU(3) vs SU(4) vs SU(5)); bundle-provenance forward-confirmation; replacement or ensemble-invariance of the LCG-random monad coefficients in `zero_modes.rs`; explicit Wilson-line representation specification.

2. *Phase 2 — real chain-channel Yukawa likelihood*: wire the existing `yukawa_pipeline.rs` (publication-grade end-to-end Yukawa pipeline already implemented in the rust_solver tree but not currently plumbed into the BF combiner) into `bayes_factor_multichannel.rs::from_yukawa_pipeline_prediction`; parallel TY implementation; bootstrap with ≥ 100 seeds; convergence guarantees with `Result<..., NotConvergedError>` short-circuit; Higgs-vev fitting/normalisation; two-loop SM RG running with threshold corrections.

3. *Phase 3 — σ-channel functional pin and matched-basis fix*: pin L^1 vs L^2; replace lex truncation with geometric basis projection; Donaldson convention robustness sweep across ≥ 3 published variants.

4. *Phase 4 — dark-matter $\Omega_{\rm DM}h^2$ re-derivation or re-labelling*: the current pipeline uses the observed Planck 0.12 as a multiplicative prefactor; either re-derive the prefactor structurally or honestly re-label as `Omega_DM_h2_consistency_estimate` with the Planck empirical anchor disclosed.

5. *Phase 5 — structural-rule generators* for the per-particle Coxeter chain indices k_X, the per-quark sign assignments χ_q, λ_q, and the rational closures used in $\sin^2\theta_W$ and $\alpha_s(M_Z)$; or honest documentation as encoded structural-rule-consistent choices in a data file.

6. *Phase 6 — convergence, reproducibility, external verification*: programmatic gCICY $(3,3)$ cycle 2.6 in a Sage/Macaulay2-equipped container; adjacent-Hodge-bin sweep at $(3,4), (4,3)$; rational-search baseline for closure rules; Zenodo-archived reproducibility kit (`cargo run --release --bin reproduce_paper_results`); external independent re-implementation in the AKLP PyTorch suite or equivalent; pre-registration of falsifiable predictions.

7. *Phase 7 — paper/pipeline alignment audit*: pre-submission line-by-line audit against a 111-test acceptance matrix.

The acceptance criteria, decision points, risk register, and test matrix for each phase are detailed in `REMEDIATION_PLAN_OPTION_B.md`. Estimated total effort: 6–12 person-weeks; submission-readiness gate is Phase 7 sign-off.

B. Open methodological items

1. *gCICY non-manifest residue at Hodge* $(3,3)$. The catalogue uniqueness sweep (Sec. II B) closes the polytope catalogue [19] and the published gCICY classifications [20–22] at $(h^{1,1}, h^{2,1}) = (3,3)$, but a future codim $\geq (3,1)$ extension producing a $(3,3)$ entry is not pre-empted. *Required work:* structural enumeration over the gCICY-extension constructions at parity with the polytope sweep, applying the same four-principle filter and reporting non-Schoen survivors if any. Until this is done the catalogue uniqueness statement at $(3,3)$ is conditional for the gCICY non-manifest residue.

2. *Donaldson convention robustness on the σ-channel discrimination.* The $|t| = 6.92$ Schoen-vs-TY σ-channel result reported in Sec. II B depends on the specific choice of Donaldson h-inversion procedure and the Kähler-form projection convention used in the metric pipeline. The paired-bootstrap sign-reversal defense addresses the basis-size artifact at matched basis size; it does not address convention-choice artifacts. *Required work:* robustness sweep over ≥ 3 published convention variants for both the h-inversion handling and the Kähler-form projection, with a stress-test report on whether the $|t| = 6.92$ result survives or degrades. The σ-channel discrimination should be re-stated with an explicit convention sensitivity envelope.

3. *Schoen-uniqueness sweep at adjacent Hodge bins.* The catalogue uniqueness statement is at the single Hodge bin $(3,3)$. *Required work:* extension of the four-cycle sweep to bins adjacent to $(3,3)$ (in particular $(3,4)$ and $(4,3)$) at the same fidelity, to confirm that the empty-bin result at $(3,3)$ is not isolated to one bin under the four-principle filter.

4. *Bundle provenance forward-confirmation.* An audit performed during the development of this paper retracted some previously asserted bundle provenance to nominally cited references; bundles currently in use have been re-pointed to primary sources. *Required work:* a comprehensive forward-confirmation that, for every bundle data set used in the present paper, the citation chain terminates at a primary source where the bundle is constructed and its stability/topological constraints are stated.

C. Open framework items deferred to companion paper

The following items are starting commitments of the present paper and are listed here for transparency. Their substrate-physical support — i.e. the argument that one would expect these and not other substrate-physical structures — is the subject of the companion paper [43]

and of the foundational monograph [42], and is not load-bearing for the numerical demonstration of this paper.

1. *Heterotic $E_8 \times E_8$ doubling.* The Coldea *et al.* [3] $CoNb_2O_6$ measurement empirically anchors one E_8 factor. The doubling to $E_8 \times E_8$ is a heterotic-string commitment [11] adopted here; an independent empirical anchor for the second E_8 factor is not claimed.

2. *McKay-icosahedral seed.* The selection of the binary icosahedral McKay-dual of E_8 as the framework's harmonic seed, and the resulting φ-chain / $\sqrt{2}$-chain decomposition, is a structural commitment.

3. *Gateway-formula structural integers.* The factor of 4 in the sub-harmonic loss 988/987, the gateway-depth $M = 14$, and the boundary capacity $h + m_{\max} = 59$ used in the dark-sector coupling family admit dual-sector / two-orientation readings within the framework but are not derived in the present paper from a specific group-theoretic counting rule on the Schoen quotient.

4. *Equivariant-cohomology identity for $\theta_{\rm QCD} = 0$.* The vanishing of the strong-CP angle is presented in Sec. VII contingent on a $\mathbb{Z}_3 \times \mathbb{Z}_3$ quotient-projected twisted-trace cancellation. The identity itself is adopted as input from the framework; an explicit proof on the Schoen $X/(\mathbb{Z}_3 \times \mathbb{Z}_3)$ quotient is part of the open program.

5. *The $\omega_{\rm fix} = 123/248$ identity is algebraic, not eigenvalue.* $\omega_{\rm fix} = (\dim -2)/(2{\cdot}\dim)$ appears in the gateway formula as an exact E_8-algebraic coefficient. It is *not* an eigenvalue of any operator and is not predicted to be empirically measurable as such. We have internally tested and retracted the alternative hypothesis under which it would be an eigenvalue; the algebraic reading is the correct one and is the only role $\omega_{\rm fix}$ plays in the construction.

6. *Hodge-pair self-conjugacy $h^{1,1} = h^{2,1}$.* The substrate-physical argument that the framework selects self-conjugate Hodge structures (and hence $(3,3)$ as the mathematical-minimum admissible bin compatible with three generations and Wilson-line breaking) lives in the companion paper [43]. The present paper takes the $(3,3)$ outcome as the input to Sec. II B.

D. What this paper does not claim

For the avoidance of overreach we state plainly:

- This paper does not claim that the starting assumptions of Table I are forced from first principles in the sense of being derivable from inside the framework's macroscopic 4D regime. From a 4D observer's vantage they are evaluated *empirically* — by the predictions they generate.

- The 26-to-zero-free-parameter reduction reported here is *conditional* on the starting assumptions; without those assumptions the present construction makes no statement.

- Schoen uniqueness in full generality across the Calabi–Yau threefold landscape is not claimed; the catalogue uniqueness statement is at the searched Hodge bin and is conditional on the gCICY residue (item 1 of Sec. XV B).

- Internal consistency of the present construction at the 0.0298 ppm ($\sim$ 30 ppb) two-anchor level (Sec. XII) is a self-consistency statement of the structural pipeline; it is not itself a derivation of the starting assumptions.

- No claim is made that the substrate-physical readings of c, $\hbar$, and the macroscopic emergence of GR/QM/Newton/Maxwell from substrate elasticity, deferred to the companion paper [43], are required for the numerical demonstration of this paper. They are interpretive context; the numerical results stand or fall on their own.

E. Further-exploration program

The most informative empirical and theoretical extensions of the present construction, in priority order:

1. Direct dark-sector searches at the predicted dominant order-5 slot $m_5 = 602.645$ neV with $g_{a\gamma\gamma} = 1.78 \times 10^{-17}$ GeV^{-1}. A null result at the predicted mass and coupling, with sensitivity exceeding the prediction, falsifies the dark-sector construction.

2. Improved precision measurements on m_W, m_H, m_t, $\sin^2\theta_W$, m_β (KATRIN), and Σm_ν (cosmological) at thresholds that either tighten the present ppm/ppb-level residuals or break them.

3. Independent reproduction of the Schoen-vs-TY σ-channel discrimination pipeline by a third party, using independent metric codes and convention sweeps.

4. Closure of the gCICY non-manifest residue (Sec. XV B item 1) and the adjacent-Hodge-bin sweep (item 3).

5. Heterotic-anomaly index or $\mathbb{Z}_3 \times \mathbb{Z}_3$ quotient-trace derivation of the gateway-formula structural integers (Sec. XV C item 3).

6. Explicit Schoen-side proof of the $\mathbb{Z}_3 \times \mathbb{Z}_3$ quotient-projected twisted-trace cancellation underlying $\theta_{\rm QCD} = 0$ (Sec. XV C item 4).

7. Peer-reviewed development of the foundational arguments relegated to the companion paper [43].

XVI. DISCUSSION

A. Parameter-count semantics

Conditional on the starting assumptions of Table I, the construction presented here contains *no fitted continuous parameter*: the 26 continuous parameters of the SM are computed structurally from the algebraic data of the starting assumptions, and the residuals reported in Sec. XIII are predictions versus measurement, not fits versus measurement. We are explicit about the scope of this statement. (i) Without the starting commitments (Table I) the construction has no content; the reduction is not a model-independent claim, but a zero-free-parameter consequence of those commitments. (ii) Two sectors carry residual conditionality even within the framework: the dark-matter abundance closes through a substrate-to-misalignment dictionary whose substrate-physical derivation lives in the companion paper, and the strong-CP angle closes through the quotient-projected twisted-trace cancellation described in Sec. VII and listed in Sec. XV C. (iii) The argument that the starting commitments themselves are minimal and substrate-physically expected is the subject of the companion paper [43] and is not load-bearing here.

B. Common geometric origin

Notable unifying features include both CP phases as rational multiples of π ($\delta_{\rm CP} = \pi + \pi/10$, $\delta_{\rm CKM} = 11\pi/30$), the dark-sector coupling family graded by E_8 exponents ($\delta_H = 1/59$, $\delta_W = 1/(59 \cdot 19)$), the quotient-based conditional route to a vanishing strong CP angle, the unified neutrino-dark-matter tower, and the strong coupling as the rational 849/7192 from forced E_8 integers.

C. Comparison to prior work

The closest empirical precedent is Coldea *et al.* [3]'s confirmation of the E_8 mass spectrum in $CoNb_2O_6$. The closest theoretical precedent is Koide's mass formula [41], which relates the three charged-lepton masses at 10^{-5} precision with no derivation. The present framework matches Koide-level precision on leptons and extends the structural construction to all 26 SM parameters plus dark matter and gravity, while leaving the compactification-level explanation of some closure factors for follow-up work.

The geometric strong-CP mechanism, if the quotient-projected twisted trace vanishes, avoids invoking Peccei–Quinn [30] and so contrasts with the dominant axion-mediated approach.

D. Open directions

The complete enumeration of open methodological items, deferred framework commitments, and the further-exploration program is given in Sec. XV. Among the items most likely to sharpen or break the present construction:

1. Direct dark-sector searches at the predicted dominant order-5 slot (Sec. XV E item 1).

2. Closure of the gCICY non-manifest residue and the adjacent-Hodge-bin sweep (Sec. XV B items 1 and 3).

3. Donaldson-convention robustness sweep on the σ-channel discrimination (Sec. XV B item 2).

4. Substrate-native cosmological evolution for the dark-matter tower, replacing the misalignment-style abundance estimate with a full continuum calculation.

5. Cosmological inflationary connection: $n_s = 58/60$ from dual-E_8 Coxeter structure.

E. Conclusion

Conditional on the starting assumptions of Table I — heterotic $E_8 \times E_8$ on the Schoen Calabi–Yau quotient $X/(\mathbb{Z}_3 \times \mathbb{Z}_3)$, the McKay-icosahedral seed, the phenomenological $\varphi/\sqrt{2}$ chain split for the lepton and quark/EW-boson sectors, the E_8 Coxeter spectrum, and the substrate-tension reading of the Higgs vacuum — the present paper computes the SM mass spectrum, mixing parameters, gauge couplings, the strong-CP angle, the cosmological dark-matter abundance, and Newton's G *directly*, with no fitted continuous parameter and no per-particle scheme choice. The Schoen-versus-Tian–Yau σ-channel discrimination (Sec. II B) and the four-cycle catalogue uniqueness sweep at Hodge $(3, 3)$ provide the empirical CY3 identification, with explicit hedges on the gCICY residue and the Donaldson convention sensitivity (Sec. XV). The numerical residuals against PDG and Planck 2018 measurements are predictions, not fits. From a 4D observer's vantage the starting assumptions cannot be *proven*; this paper evaluates them by empirical consistency, and reports the result. The companion paper [43] addresses the separate question of whether those starting assumptions are minimal and substrate-physically expected; that argument is independent of the present demonstration and is not relied upon here.

XVII. COMPUTED 4D EFFECTIVE LAGRANGIAN: FRAMEWORK PREDICTION WITH VERDICT (A) FALSIFICATION

This section reports the end-to-end computed assembly of the 4D effective Lagrangian under the rank-4 $SU(4)$ Braun–He–Ovrut–Pantev (BHOP-2005 [? ?]) bundle on the Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ quotient.

a. Path A research closure (verdict (a) ROBUST). At the time of writing, the Path A heterotic line of attack has reached definitive closure. Three sequential computational artefacts on the `Schoen` $\widetilde{X}/\Gamma$ quotient establish:

- **Shadow** $\mathrm{Ext}^1 = 0$**.** The line-bundle-shadow $\mathrm{Ext}^1(V_2, V_1)$ on the Schoen quotient was computed to be exactly zero (commit `831d910c`; reference document `references/p_ext1_engagement_2026-05-05.md`).
- **dP9-side** W**-bundle** $h^1 \geq 3$**.** The W-twisted contribution to Ext^1 on the dP9 side was found to carry $h^1 \geq 3$ in non-trivial $\mathbb{Z}_3 \times \mathbb{Z}_3$ characters, with $\sum \mathrm{mult} \cdot \chi_{\mathrm{dP9}} = +18$ (commit `eea6f085`; `references/p_dp9_w_lift_2026-05-05.md`).
- $\mathbb{Z}_3 \times \mathbb{Z}_3$**-invariant projection yields** 0**.** The $+18$ resides entirely in the non-invariant $\beta = 2$ characters; the BHOP-canonical bundle-frame lift $(g_\alpha, g_\beta) = (0, 0)$ projects the invariant subspace to exactly 0 at both the lower and upper SES bounds (commit `a4e6231e`; `references/p_dp9_w_invariant_proj_2026-05-06.md`).

This chain of three computations definitively settles $\mathrm{Ext}^1_{\mathrm{full}} = 0$ on the Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ quotient. The robustness argument is geometric: the BHOP §6.1 SES sub-bundle $O(-2f)$ rigs the elliptic-fiber H^1 direction to sit at $\beta = 2 \neq 0$, and no admissible choice of L_{base} on dP9 can move that character without breaking the SES design (see the cited reference for the closed-form proof).

b. Consequence for the framework prediction. The off-diagonal $SU(4)$ coupling κ *vanishes* as a geometric consequence of the BHOP-2005 construction on the Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ quotient. The $\kappa = 0$ frame-collapse output is therefore the *framework's canonical prediction* at the rank-4 BHOP-2005 level — not a placeholder pending a future Ext^1 engagement.

The published 4D effective Lagrangian therefore carries:

- **4 PHYSICAL sectors** (kinetic, gauge, Higgs potential, Newton's G) reproducing PDG/CODATA observables to ppm-scale.
- **12 FRAMEWORK_PREDICTION_FALSIFIED rows** (six quark Yukawa eigenvalues at TeV scale, five CKM Wolfenstein parameters at zero, the Jarlskog phase at zero, and the tree-level cosmological constant) in tension with PDG and Planck.
- **7 STRUCTURALLY_ABSENT rows** (three charged-lepton masses, three neutrino masses, the four-parameter PMNS bundle) reflecting the fact that the rank-4 BHOP-2005 shadow does not encode the e^c and ν^c slots at degree-1 H^1 on the Schoen quotient.

This is a first-class falsifiable scientific result, not an unfinished intermediate. The construction is internally consistent and produces a specific, mismatched prediction that any third party can reproduce from the SHA-pinned chain. Honest framing is supplied sector by sector below.

A. Canonical form and provenance chain

The assembled 4D effective Lagrangian takes the canonical form

$$\boxed{\mathcal{L}_4 \;=\; \mathcal{L}_{\mathrm{kin}} + \mathcal{L}_{\mathrm{gauge}} + \mathcal{L}_{\mathrm{Yukawa}} + V(H) + \Lambda_{\mathrm{cc}},} \tag{68}$$

with

$$\begin{aligned}
\mathcal{L}_{\mathrm{kin}} &= i\,\bar{\psi}_i\,\gamma^\mu D_\mu \psi_i + |D_\mu H|^2,\\
\mathcal{L}_{\mathrm{gauge}} &= -\tfrac{1}{4}\sum_a F^a_{\mu\nu}\,F^{a\,\mu\nu}, \quad a \in \{SU(3)_C, SU(2)_L, U(1)_Y, U(1)_X\\
\mathcal{L}_{\mathrm{Yukawa}} &= -Y^u_{ij}\,\bar{Q}_i u^c_j H - Y^d_{ij}\,\bar{Q}_i d^c_j \tilde{H} - Y^e_{ij}\,\bar{L}_i e^c_j \tilde{H} - Y^\nu_{ij}\,\bar{L}_i \nu^c_j H + \text{h.c.}\\
V(H) &= -\mu^2\,|H|^2 + \lambda\,|H|^4,\\
\Lambda_{\mathrm{cc}} &= V_{\min}(H) \quad \text{(tree-level)}.
\end{aligned}$$

The numerical workflow producing the values below is implemented in the Rust binary `p_lagrangian_assemble` (commit `ad1fd1b2`), which consumes the upstream Yukawa-matrix artefact from `p_lagrangian_yukawa_matrices` (commit `e5a61df7`) which in turn consumes the eigenmode catalogue from `p_lagrangian_eigenmodes` at $n_{\mathrm{pts}} = 432$ (commits `0b90c1bc` + `77385de0`). The provenance chain SHA-256s, pinned for reproducibility, are

Artefact	Replog terminal SHA-25
Eigenmode catalogue ($n_{\mathrm{pts}} = 432$, BHOP)	`1a5a13bf...8e7d6fe9`
Yukawa matrices ($\kappa = 0$)	`5c043eb0...051d5753`
Assembled Lagrangian	`1b1e5e4a...fafed5d01`
Sectors-payload SHA	`f4079981...270dd9965`

The full 64-character chains are `1a5a13bf010f7314beab35ad81b4b2885a6eeaaf756849c5c97e9d3e8` (eigenmode catalogue), `5c043eb0355710a67b4929c4e8dfb01448b2d` (Yukawa matrices), `1b1e5e4a6adbb2a427bcbe7ea189ce5a8aabeb02` (assembled Lagrangian replog), and the canonical sectors-payload SHA `f407998187055a74b64ce0e7c9397c4d6b2133cb535b943f48d2a3127`

B. Lagrangian header summary

C. Newton's G from heterotic volume reduction

The 4D Newton constant emerges as a closure check between three structural commitments, none of which is

TABLE XVIII. 4D effective Lagrangian sector summary (build `e5a61df7_assembled_0b90c1bc_bhop_432_0xb077`). **Status** marks each sector as *physical* (structural prediction matching observation), *framework prediction (falsified vs PDG)* (the rank-4 BHOP-2005 $\kappa = 0$ value is the framework's canonical prediction post Path A closure — in tension with experiment), or *structurally absent* (the rank-4 shadow does not populate the slot at degree-1 H^1).

Sector	Status	Headline numerical content
Kinetic	physical	12 chiral fermions; H^1 dim = 48; canonical normalisation
Gauge	physical	$g_1 = 0.461228$, $g_2 = 0.65096$, $g_3 = 1.2123$ at M_Z
Higgs $V(H)$	physical	$m_H = 125.1998$ GeV; $v_C = 246.21965$ GeV; $\lambda = 0.12928$; $\mu^2 = 7837.49$ GeV^2
Newton G	physical	$G_{\rm rec} = 6.67443 \times 10^{-11}$ vs CODATA 6.6743×10^{-11} (~ 19 ppm)
Yukawa	framework prediction (falsified)	$Y_u = Y_d = \mathrm{diag}(2267.25, 2358.93, 2200.59)$ GeV; PDG $m_t \approx 173$ GeV (factor $\sim$9, $> 5\sigma$)
Y_e, Y_ν	structurally absent	e^c, ν^c slots empty at degree-1 H^1 of rank-4 BHOP shadow
CKM	framework prediction (falsified)	$V_{\rm CKM} = \mathbb{1}_3$; Wolfenstein $\lambda = A = \bar\rho = \bar\eta = J = 0$ vs PDG non-trivial
$\Lambda_{\rm cc}$	framework prediction (falsified)	$V_{\rm min} = -3.38 \times 10^{-66}\,(M_{\rm Pl}^{\rm red})^4$ (56.5 orders too large; tree-level SM CC problem)

a fit parameter: the F_e gateway (Eq. (6)) fixing $m_{\rm Pl} = m_e/F_e$, the Schoen-quotient internal volume $V_X^{\rm central} = 4/3$ (structural-exact from the $\mathbb{Z}_3 \times \mathbb{Z}_3$ triple intersections at the central Kähler point $t_1 = t_2 = t_T = 1$), and the standard heterotic relation $1/\kappa_4^2 = V_X/\kappa_{10}^2$ between the 10D and 4D gravitational couplings. Substituting the gateway-fixed Planck mass and $V_X^{\rm central}$ recovers

$$G_{\rm rec} = 6.67442860 \times 10^{-11}\ \mathrm{N \cdot m^2/kg^2},$$
$$G_{\rm CODATA\text{-}2022} = 6.6743 \times 10^{-11}\ \mathrm{N \cdot m^2/kg^2},$$

i.e. relative residual $\approx 1.93 \times 10^{-5}$ (~ 19 ppm). The independent gateway cross-check $m_{\rm Pl}(F_e)$ versus $m_{\rm Pl}(G_{\rm CODATA}, V_X)$ matches at -1.86×10^{-6} (~ 2 ppm). The implied string scale at the central Kähler point is $M_{\rm string} \approx 1.90 \times 10^5$ GeV in the binary's reduced-Planck convention; the committed Kähler-cone range gives $V_X \in [0.0208, 682.7]$ over which the residual budget varies. G is *not* a fit parameter.

D. Higgs potential as derived corollary

Once v_C and m_H are predicted structurally (Sec. V; chain positions $k_v = 1132/30$ and $k_H = 1072/30$), the SM tree-level relations $\mu^2 = m_H^2/2$ and $\lambda = m_H^2/(2v_C^2)$ are arithmetic identities of the SM Higgs Lagrangian and follow without further input:

$$\boxed{\begin{aligned} m_H^{\rm pred} &= 125.19978\ \text{GeV}, & v_C &= 246.21965\ \text{GeV}, \\ \mu^2 &= 7837.4926\ \text{GeV}^2, & \lambda &= 0.129280. \end{aligned}} \tag{69}$$

PDG 2024 references: $m_H = 125.20 \pm 0.11$ GeV (central inside the 1σ bar); $v_{\rm tree} = 246.21965$ GeV (sub-ppm). μ^2 and λ are derived corollaries of the structurally predicted m_H and v_C, not independent predictions.

E. Yukawa sector: framework prediction at $\kappa = 0$ fixed point, falsified vs PDG

Post Path A research closure (verdict (a) ROBUST), the Yukawa matrices emitted by the assembly binary at the rank-4 BHOP-2005 $\kappa = 0$ fixed point are the framework's *canonical prediction* for the visible-sector Yukawa eigenvalues:

$$Y_u = Y_d = \mathrm{diag}\big(2267.25,\ 2358.93,\ 2200.59\big)\ \text{GeV},$$
$$V_{\rm CKM} = \mathbb{1}_3,$$

with diagonalised mass eigenvalues (descending, GeV) [1668.02, 1603.19, 1556.05] for both up- and down-type quarks. The charged-lepton and neutrino blocks Y_e, Y_ν are *structurally absent*: the e^c and ν^c slots carry zero modes in the degree-1 $H^1(\widetilde{X}/\Gamma, V_{\rm BHOP})$ catalogue of the rank-4 shadow bundle (irrep counts: Q, U^c, D^c, L each 12; E^c, ν^c, H_u, H_d each 0).

a. Falsification status (visible Yukawa, six rows). The framework predicts $\mathcal{O}$(TeV) Yukawa eigenvalues where PDG 2024 records the hierarchical pattern $m_t \approx 173$ GeV, $m_b \approx 4.18$ GeV, $m_c \approx 1.27$ GeV, $m_s \approx 93$ MeV, $m_d \approx 4.7$ MeV, $m_u \approx 2.16$ MeV. Every row mismatches by factors of ≥ 9, with most diverging by 2–6 orders of magnitude. This is a $> 5\sigma$ falsification on every quark mass. The CKM matrix predicted to be the 3×3 identity is similarly falsified: PDG observes $\lambda \approx 0.225$, $A \approx 0.826$, non-zero $\bar\rho, \bar\eta$, and Jarlskog $J \approx 3.08 \times 10^{-5}$.

b. Why the falsification is real, not an unfinished computation. The Path A chain (commits `831d910c`, `eea6f085`, `a4e6231e`; reference docs in `book/scripts/cy3_substrate_discrimination/rust_solver/references/`) shows that $\mathrm{Ext}^1_{\rm full} = 0$ on the Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ quotient is a definitive theorem, not a temporary placeholder. The dP9-side W-twist h^1 contribution is non-zero (+18) but lives entirely in non-invariant $\beta = 2$ characters and projects to 0 on the quotient. Consequently the off-diagonal $SU(4)$ coupling κ vanishes as a geometric consequence of the BHOP-2005 §6.1 SES design, and the rank-4 block-diagonal $\kappa = 0$ Lagrangian *is* the framework's canonical visible-Yukawa prediction. There is no remaining computation to run on

this Schoen quotient that would lift the rank-4 BHOP shadow out of the $\kappa = 0$ fixed point.

c. Structural absence of Y_e, Y_ν (seven rows). The $Y_e \equiv 0$ and $Y_\nu \equiv 0$ blocks are a falsifiable structural prediction of the rank-4 BHOP-2005 construction at degree-1 H^1, not an unfinished computation. The framework, as committed, does not encode the charged-lepton or neutrino sectors at this construction level. Any third party able to extract a non-zero E^c- or ν^c-irrep contribution from the rank-4 BHOP shadow's degree-1 H^1 catalogue on the Schoen quotient would falsify this structural claim — equivalently, would refute the irrep-count audit (commit `a4e6231e`).

F. Cosmological constant: tree-level Higgs $V_{\min}$ only

The tree-level Higgs potential of Eq. (69), evaluated at its minimum $|H|^2 = v_C^2/2$, gives

$$V_{\min} = -\frac{m_H^2\, v_C^2}{8} \;\approx\; -1.188 \times 10^8 \text{ GeV}^4$$
$$\approx\; -3.377 \times 10^{-66}\, (M_{\rm Pl}^{\rm red})^4,$$

versus the Planck 2018-derived $\Lambda_{\rm obs} \approx 1.1 \times 10^{-122}\, (M_{\rm Pl}^{\rm red})^4$, i.e. approximately 56.49 orders of magnitude too large in absolute value, with the sign flipped (tree-level $V_{\min}$ is negative, observed $\Lambda_{\rm obs}$ is positive).

Honest framing. This is the well-known Standard-Model cosmological-constant problem at tree level, NOT a framework-specific defect. Any heterotic compactification with the SM-coupled-to-gravity tree-level Higgs minimum produces the same mismatch. Loop corrections (Stage 7: gauge-boson + fermion + scalar contributions, plus a possible Sakharov-induced-gravity reading of the substrate's baseline tension) are not committed by the present construction and are explicitly deferred. We flag this here so the reader does not misread the tree-level value as a framework prediction.

G. Milestone interpretation: complete prediction set, end-to-end falsifiable

This section records the end-to-end *computed* 4D Lagrangian assembly under the BHOP-2005 rank-4 $SU(4)$ bundle: an eigenmode catalogue ($n_{\rm pts} = 432$, BHOP H^1 dim 48) feeds a Yukawa-overlap pipeline that produces the rank-4 $\kappa = 0$ Yukawa matrices, which in turn assemble the canonical-form Lagrangian of Eq. (68) with full SHA-pinned provenance. Together with the dark-sector chain slots and falsifiability signatures emitted by the `p_dark_sector_predictions` binary, this delivers the framework's *complete end-to-end falsifiable prediction set* (27 visible-sector + 16 dark-sector slots + 6 falsifiability signatures), consolidated by the `p_complete_prediction_set` binary into the single canonical artefact `output/p_complete_prediction_set.json`. The companion paper, *Substrate Falsifiable Predictions*, renders this artefact into the canonical longtable.

Sector-by-sector reading:

- **4 PHYSICAL sectors** reproduce the kinetic, gauge, Higgs-potential, and Newton's-G closures already documented in earlier sections; these are first-class structural predictions matching PDG/CODATA observations.
- **12 FRAMEWORK_PREDICTION_FALSIFIED rows** (six visible-sector Yukawa eigenvalues, five CKM Wolfenstein parameters, Jarlskog $\delta_{\rm CKM}$, tree-level $\Lambda_{\rm cc}$) record the rank-4 BHOP-2005 $\kappa = 0$ fixed-point prediction and its $> 5\sigma$ tension with PDG. Per Path A closure (commits `831d910c`, `eea6f085`, `a4e6231e`), there is no remaining gate-defect computation that would lift the rank-4 BHOP shadow out of the $\kappa = 0$ fixed point on the Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ quotient.
- **7 STRUCTURALLY_ABSENT rows** (three charged-lepton masses, three neutrino masses, the four-parameter PMNS bundle) record the rank-4 BHOP shadow's empty E^c, ν^c slots at degree-1 H^1 — a falsifiable structural claim of the construction.
- **16 PHYSICAL dark-sector slots** (closed-form geometric expressions in E_8 invariants and the gateway constants $\alpha_\tau, h_{\rm eff}$) populate the dark Tier 1/2 chain; three falsifiability signatures already PASS, two are DISCOVERABLE_BY_2030 at HL-LHC, and one is a falsifiable convention-choice signal.

The total artefact is therefore a *publishable mixed-outcome* result: four sectors confirmed against experiment, twelve falsified, seven structurally absent (also falsifiable), plus a sixteen-slot geometric dark-sector tower with six independent falsifiability signatures. Every row carries SHA-pinned provenance back to the upstream replogs.

XVIII. CHANNEL-GROUNDING MAP AND SUBSTRATE-NATIVE DERIVATION: WHERE THE MASS EQUATIONS LIVE ALGEBRAICALLY

This section documents a 27-stream, $\sim$ 1.6-billion-cell investigation of *where* each row of the framework's mass-and-coupling table lands inside published mathematical structures, what unified principle generates the grounded entries, and which residue remains parameterisation rather than derivation. The full row-level matrix is maintained as `book/scripts/cy3_substrate_discrimination/rust_solver/out` with the synthesis overlay in `book/scripts/cy3_substrate_discrimination/references/p_gr` upstream evidence files are

p_phase1_pattern_extraction.md,
p_phase2_h4_wilson_test.md,
p_phase2_d8_orbit_test.md,
p_phase2_tau_chain_unification.md,
p_phase2_vacuum_icosian_B.md,
p_phase3a_unified_synthesis.md,
p_phase3b_quark_bmultiset_search.md,
p_phase3c_1132_extended_search.md,
p_phase3d_hostile_review.md,
p_phase4_algebraic_form_match.md,
p_phase5_substrate_native_derivation.md,
p_dim_minus_2_unification.md, and
p_grounding_channels_v4_results.md in the same references/ directory.

A. The 89-row grounding channel matrix

Every framework row – charged fermions, electroweak bosons, gauge couplings, mixing angles and CP phases, the Higgs and vacuum modes, the sub-gateway neutrino tower, the above-gateway dark tower, the half-integer dark slots, gravitational-sector outputs, the open cosmological rows, the SI-primitive substrate quantities, and the mixing-angle sub-rows – is enumerated as a single matrix entry. Each entry carries a verdict, the layer of the framework's generating principle under which it is (or is not) forced, and a citation key into the relevant Phase report. Verdict counts are:

TABLE XIX. Verdict distribution across the 89-row grounding channel matrix. GROUNDED collapses every row whose value is forced by substrate primitives – the Sugawara-on-$\widehat{e}_8$ rows, the Wilson-line plane-projection rows, the icosian closure rows, the algebraic E_8-invariant rows, the SI-primitive rows, and the absolute-scale anchor h_{eff} – since each is a sub-flavor of substrate-primitive forcing in absolute SI units. GROUNDED-VIA-CODATA flags rows where the framework's prediction is confirmed by an experimentally-measured value. GROUNDED-VIA-GATEWAY-SUBTRACTION records the four $\dim -2 = 246$ unification promotions. INTERPRETIVE flags rows where multiple substrate-natural readings exist but no single rule privileges one. CONDITIONAL-POSSIBILITY flags structurally-allowed slots the framework does *not* predict to be occupied.

Verdict	Count
GROUNDED (substrate-primitive in SI)	29
GROUNDED-VIA-CODATA	12
GROUNDED-VIA-GATEWAY-SUBTRACTION	4
INTERPRETIVE	25
CONDITIONAL-POSSIBILITY	14
PARTIAL	2
NO-WITNESS-YET	2
FALSIFIED	1
BESPOKE	0
Total	**89**

Forty-five of the eighty-nine rows carry a GROUNDED-* verdict (29 + 12 + 4); a further twenty-five rows are INTERPRETIVE (multiple substrate-natural readings, no privileged forcing); fourteen flag structurally-allowed but unpredicted slots; and the remaining five rows are split across PARTIAL, NO-WITNESS-YET, and FALSIFIED. The BESPOKE count is zero: every row that sets an absolute mass scale or dimensionless coupling is forced by substrate primitives in absolute SI units, with the SI unit system itself substrate-derived (§XVIII E).

The 89 rows partition as 74 framework rows from the Phase 1–5 investigation, plus 4 gateway-subtraction promotion records (the $\dim -2 = 246$ unification of Hodge saturation, h_{eff} anchor, and Sugawara level span; §XVIII B), plus 8 mixing-angle sub-rows (per-element CKM and PMNS Wolfenstein-style entries), plus 3 SI-derivation rows (c, $\hbar$, e).

The cumulative search depth that produced these verdicts is on the order of 1.6×10^9 test cells across ≥ 27 enumeration streams, including: an exhaustive $W(D_8)$ orbit walk for the quark b-multiset (Phase 2.2 / Stream C); a 5.77-million-cell extended structural-class scan over $k_v = 1132/30$ (Phase 3C, on top of Phase 2 Stream D's 12-vector base); a $4^4 \times 240 = 61{,}440$ Wilson-line $\times$ doublet-axis sweep (Phase 2.1 + Phase 3D 2.1.A); and the multi-billion-cell Stream A/B/C/D/E enumeration of inverse Sugawara shifts on $\hat{e}_8$, half-integer denominator candidates, k_v reference values, and tau-chain depth alternatives.

B. The 3-piece unified generating principle and the $\dim -2 = 246$ substrate identity

The grounded layer is generated by a single principle with three algebraic pieces. Each piece is uniqueness-checked at a different strength. Underlying all three is the substrate-primitive identity $\dim -2 = 246$: the bulk-mode internal-dimension count, with two directions reserved for the gateway. The same integer 246 is forced to appear in three structurally-distinct places by the substrate-primitive chain universal-$c \to$ Hodge self-mirror $\to$ identity counting.

a. The $\dim -2 = 246$ *unification.* The integer 246 appears at integer precision in three places, each of which is forced by substrate primitives on a different cycle:

TABLE XX. Three-pin verification of the $\dim -2 = 246$ unification. All three formulas pin to the same integer; the unifying reading is that 246 is the bulk-mode count after subtracting the two gateway directions. Source: p_dim_minus_2_unification.md.

Pin	Formula	Value
Hodge saturation	$\lvert\Phi(E_8)\rvert + h^{1,1} + h^{2,1}$	$240 + 3 + 3 = 246$
h_{eff} anchor	$(\dim -2)/\dim$	$246/248$
Sugawara level span	$k_{\max} - k_{\min}$	$248 - 2 = 246$

The substrate-physical reading of the Hodge pin is

that 240 is the E_8 root count and $h^{1,1} = h^{2,1} = 3$ is the Schoen-quotient Hodge-pair generation count, with self-mirror symmetry forcing equality between the Kähler and complex moduli. The h_{eff} anchor pin is the gateway-residue scale factor in $m_e = ((\dim -2)/\dim) \cdot h_{\text{eff}}$, with the 2 in the numerator subtracting exactly the two gateway directions from the full adjoint dimension. The Sugawara level-span pin is the difference between the adjoint level $k_{\max} = \dim$ and the chirality-flip level $k_{\min} = 2$ on affine $\widehat{e}_8$. The same gateway-2 operator that gives $\dim -2 = 246$ also gives $\text{rank} - 2 = 6$ (the bulk-Cartan rank used in the neutrino-mass denominator $1/(\text{rank}-2)$) and $h_{D_8} - 2 = 12$ (the bottom-quark b-multiset entry). The four resulting promotions (GROUNDED-VIA-GATEWAY-SUBTRACTION) are recorded as $\nu_{1,2,3}$ and the bottom quark.

b. (i) Affine $\widehat{e}_8$ Sugawara at adjoint level $k = \dim(E_8) = 248$. Six of the eight dimensionless rationals (α_τ, $\alpha_\tau^{\text{full}}$, Δ_φ, Δ_2, δ_H, α_ε) sit on inverse Sugawara shifts $1/(k + h^\vee)$ at structurally distinguished sub-levels $k \in \{1, 2, e_{\max} = 29, \dim = 248\}$, with H_4 icosian shifts $+\varphi^2 - 1 = +\varphi$ where they apply. No published competing identification fits the same six denominators simultaneously (Phase 1, Stream A, P7.12). This is uniqueness-as-rule, not uniqueness-of-fit; individual denominators $\{32, 59, 278, 496\}$ each admit multiple structural integer factorisations.

c. (ii) Wilson-line projection $(1, 1, 0, 0)$ on the $H_4 \oplus \tau \cdot H_4$ Coxeter-plane decomposition. The four Coxeter 2-planes of E_8 pair the exponents $(1, 29)$, $(7, 23)$, $(11, 19)$, $(13, 17)$. A Wilson line $W \in \{0, 1\}^4$ on these planes selects, by surviving-root projection, which planes host charged-lepton modes versus dark modes versus non-doublet-confined modes. Phase 2.1 enumerated all $2^4 = 16$ patterns over the canonical $\text{SU}(2)_L$ doublet axis; only $(1, 1, 0, 0)$ confines the doublet roots to the lepton-bearing planes $(11, 19)$ and $(13, 17)$. Phase 3D check 2.1.A swept the doublet axis across *all 240* E_8 roots and confirmed that $(1, 1, 0, 0)$ remains *the unique* surviving lepton-plane pattern (the 240/240 enumeration is bit-stable). Caveat: which exponent within a plane hosts which lepton ($e = 11$ vs. $e = 19$, $e = 13$ vs. $e = 17$) is *not* forced by the Wilson-line projection – it is a 1-bit chirality selection per plane that the framework supplies separately.

d. (iii) Icosian-closure rule $B = 2\varphi$ for vacuum-tension modes. The Higgs vev and Higgs boson rows close their structural Σ corrections via $B = 2\varphi = 1 + \sqrt{5}$. Phase 2.4 identified 2φ as the additive double of the fundamental unit φ of $\mathbb{Z}[\varphi]$, and equivalently as $4\cos(\pi/5)$, twice the $I_2(5) = H_2$ dihedral Coxeter trig at the primary exponent e_1. Within the McKay-doubled H_2 family $\{2\cos(k\pi/5)\}_{k\in\{1,2,3,4\}}$, the doubled $k = 1$ representative is the only entry the framework's empirical structure selects at sub-ppm precision. The doubling reflects the binary icosahedral cover $2I \to I$, the McKay correspondence whose ADE image is exactly E_8. The 1-bit "doubled vs. undoubled" choice is not internally forced by piece (i) or piece (ii).

The same single E_8 integer identity $h_{SU(3)} + \text{rank}(E_8) = e_3(E_8)$ $(3+8 = 11)$ underlies the α_τ^n tower depth set $\{n : h_{SU(3)} \le n \le e_3\} = \{3, 4, \ldots, 11\}$ for the neutrino and dark-tower rows (Phase 2.3). Phase 3D check 2.3.A verified that this identity is satisfied by seven simple Lie algebras $(B_4, C_4, D_6, D_8, E_8, F_4, G_2)$; within the framework's external commitments (color subalgebra $= SU(3)$, ambient $= E_8$, consecutive-integer tower postulate) the depth set is forced, but the identity by itself is not E_8-unique.

e. Internal toroidal T^2 as the Schoen 1-skeleton projection. The three pieces above act on the framework's *internal toroidal* T^2 – the elastic-response sublayer – which is the 1-skeleton / 1-homology projection of the Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ Calabi–Yau threefold (`thm-substrate-dopw-schoen-construction`). The two independent 1-cycles of T^2 are exactly the two independent $\mathbb{Z}_3$ generators of $\pi_1(\text{Schoen}/\mathbb{Z}_3 \times \mathbb{Z}_3)$ – the Wilson-line phase axes – and the Wilson-line projection $(1, 1, 0, 0)$ on the $H_4 \oplus \tau \cdot H_4$ Coxeter-plane decomposition is the macroscopic-layer rendering of the two $\mathbb{Z}_3$ generators acting on the internal CY3. The Schoen-internal full threefold hosts the bundle cohomology, the Hodge moduli $h^{1,1} = h^{2,1} = 3$, the Atiyah–Singer 3-generation index, and the Wilson-line breaking chain $E_8 \to SO(10) \times SU(4) \to SU(5) \times U(1) \to G_{\text{SM}} \times U(1)_X$. The toroidal T^2 is what the three pieces of the unified generating principle directly see; the Schoen-internal content is what makes the T^2 projection well-defined and substrate-physical. Subsection XVIII F develops this two-layered structure in detail.

C. Substrate-native derivation of the mass equations from coupling-reach mechanics

The framework's mass formulas were originally written in heterotic-CY3 wave-function-overlap language. Phase 5 re-derives them directly from the substrate's intrinsic primitives – continuity of the substrate amplitude (`hyp_substrate_continuity`), positive coupling-event duration and reach (`hyp_substrate_positive_coupling_event`), the discrete local-coupling update rule (`hyp_substrate_local_coupling_discretization`), the continuum wave equation (`hyp_substrate_wave_equation`), the universal substrate ratio $c = \ell_c/\tau_c$ (`hyp_substrate_c_as_coupling_ratio`), modes-as-substrate-configurations (`hyp_substrate_modes_are_substrate_configurations`), tension modulation of the local coupling rate (`hyp_substrate_tension_modulation`), and the energy/mass–substrate-tension equivalence (`hyp_substrate_mass_energy_equivalence`). These are the *only* primitives the Phase 5 derivation invokes; no heterotic-CY3 wave-function-overlap formula is borrowed.

The substrate-native mass formula reads

$$R = \text{base}^k \cdot \exp\big(\chi \cdot \Delta_{\text{chi}}\big) \cdot \exp\big(\lambda \cdot |k - h/2| \cdot \alpha_{\text{reach}}\big) \cdot \Sigma_{\text{closure}}, \tag{70}$$

where the structural ingredients are all substrate-derived (no per-particle freedom):

TABLE XXI. Substrate-physical interpretation of the framework's structural ingredients in Eq. (70).

Quantity (substrate-physical)	Symbol	Identity
inverse coupling reach (chain-position units)	α_ε	$1/\dim(E_8 \times E_8) = 1/496$
chirality-flip energy gap (level-2 Sugawara)	Δ_2	$1/(2 + h_{E_8}) = 1/32$
icosahedral chirality-flip gap (H_4)	Δ_φ	$1/(h + \varphi^2)$
sub-Lie cycle period of substrate amplitude	h_{D_8}	14
centre of vacuum coupling-reach kernel	$h/2$	15
icosian closure for vacuum-tension mode	2φ	2φ
colour Coxeter (closure for $SU(3)$)	$h_{SU(3)}$	3
gateway algebraic coefficient	$(\dim -2)/\dim$	$246/248$

The boxed line is the *substrate-Boltzmann amplitude* for a mode at chain position k to overlap with the Higgs-vacuum-tension mode centred at $h/2 = 15$, with kernel decay length ℓ_c. The small-parameter linearisation $\exp(x) \approx 1 + x + O(x^2)$ at $\Delta_2 \approx 0.031$, $\alpha_\varepsilon \approx 0.002$, and $|k - h/2| \cdot \alpha_\varepsilon \leq 22/496 \approx 0.044$ reproduces the framework's IDENTITY closed forms bit-identically across all six quarks and the lepton/EW/Higgs rows; the native vs. linearised gap is bounded by the second-order Taylor remainder, which peaks at the top quark at 1.39×10^{-3} – matching the Phase-4 form-match residual to within 3% (the small difference is the Taylor upper bound versus the actual residual).

The substrate-physical reading of the special chain positions is:

- **electron at** $k = 0$ – gateway eigenvalue $\omega_{\text{fix}} = (\dim -2)/(2\dim) = 123/248$, an exact E_8 algebraic identity (P7.12). Off the Coxeter ladder; no chain-position freedom; no within-row chirality bit. Fully UGP-derived.

- **strange at** $k = h/2 = 15$ – midpoint of the substrate's vacuum coupling-reach kernel. Gaussian overlap collapses to the base factor only; $\chi = \lambda = 0$, structural fixed point.

- **Higgs vev at** $k_v = 1132/30$ – closure $B = 2\varphi$ (icosian sector) rather than $B = h_{SU(3)}$ (colour). The chain position $k_v = 1132/30$ is set by the icosian-closure rule $B = 2\varphi$ acting on the vacuum-tension mode and by the gateway-subtraction structure that fixes the h_{eff} scale factor (§XVIII B). Phase 2 Stream D's 12 vectors and Phase 3C's 5.77×10^6 additional cells across 16 structure classes (OEIS, class numbers, modular tensor category q-dim, etc.) confirm that no *competing* published mathematical structure reproduces this specific rational, but the substrate-primitive forcing is the icosian closure, not a free fit. The JSON field `k_v_exhaustively_confirmed: true` records the Phase 3C search depth.

- **quark** b**-multiset** $\{4, 1, 12, 13/2, 17/2, 17/2\}$ – the bottom-quark entry $b = 12 = h_{D_8} - 2$ is now gateway-subtraction-grounded (§XVIII B). The remaining five entries are INTERPRETIVE: each admits multiple substrate-natural readings on the $\sqrt{2}$-chain Coxeter ladder. The multiset is not a single $W(D_8)$ orbit (Phase 2.2); the per-quark (N, b, χ, λ) assignments are 1-bit chirality choices within Wilson-line-forced planes rather than a single $W(D_8)$-orbit forcing.

d. Two-layered attribution of the structural ingredients. Each ingredient in Table XXI traces to one of two substrate-physical sublayers. The *toroidal elastic-response* sublayer (the internal T^2 that is the 1-skeleton projection of Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$) supplies the inverse coupling reach $\alpha_\varepsilon = 1/\dim(E_8 \times E_8)$, the chirality-flip Sugawara gaps Δ_2 and Δ_φ, the icosian closure 2φ on vacuum-tension modes, the colour Coxeter $h_{SU(3)}$, and the gateway algebraic coefficient $(\dim -2)/\dim$ – all of which live on the $E_8 \times E_8$ root content, the affine $\widehat{e}_8$ Sugawara level structure, and the Coxeter-plane decomposition of the toroidal layer. The *cylindrical macroscopic* sublayer (the framework's $\mathbb{R} \times S^1$ propagation layer, where S^1 is the macroscopic projection of one of the two T^2 1-cycles) supplies the chain-position k as Kaluza–Klein momentum / winding label, the sub-Lie cycle period $h_{D_8} = 14$ as the substrate amplitude's periodicity in chain index, and the vacuum-kernel centre $h/2 = 15$ as the macroscopic mid-point of the substrate's coupling reach window. The per-particle (N, b, χ, λ) labels are substrate-natural KK-plus-winding readings on this macroscopic sublayer; the Wilson-line-forced planes are the toroidal-layer constraint that fixes which slots the macroscopic labels may occupy.

D. Per-ingredient $E_8 \times E_8$ forcing audit

Phase 5 closes with an explicit per-ingredient forcing audit of the heterotic $E_8 \times E_8$ commitment, separating the ingredients that ARE substrate-primitive forced from those that are committed without being derived. The status of each is recorded against the hypothesis/theorem files in Part 3 of the book.

a. 6 PRESENT-DERIVED ingredients (substrate-primitive forcing chain in place).

1. Critical worldsheet dimension $D = 10$ – from worldsheet conformal-anomaly cancellation under stable-matter / supersymmetric worldsheet structure (`thm_substrate_worldsheet_critical_dimension`).

2. Fermionic worldsheet modes – from $\mathbb{Z}/2$-handedness of the substrate amplitude plus the squared-Klein–Gordon equation of motion.

3. Worldsheet supersymmetric pairing – from tachyon elimination together with the GSO projection.

4. Rigid Lie-algebra structure of gauge-boson couplings – via the Cartan–Killing classification of compact simple Lie algebras admitting an invariant inner product.

7. Anomaly cancellation forcing $\dim G = 496$ – via polynomial factorisation of the 12-form anomaly (REM-OPT-B Phase 1.1 regression test on file).

10. Selection of $E_8 \times E_8$ over $\mathrm{Spin}(32)/\mathbb{Z}_2$ – via the McKay correspondence $2I \leftrightarrow E_8$ together with the empirical lepton φ-chain (`thm_substrate_e8xe8_forced_from_mckay`). The forcing statement is conditional on the substrate-primitive setting (D = 10 worldsheet supersymmetric Yang–Mills with quantum self-consistency), not unconditional.

b. 2 PARTIAL ingredients. Yang–Mills minimal coupling (committed, not derived); the closed-loop quantum partition function (referenced, not constructed at the level of the other six).

c. 2 ingredients closed via the toroidal elastic-response sublayer. The two ingredients previously flagged as ABSENT both close on the framework's internal toroidal T^2 – the elastic-response sublayer, which is the 1-skeleton projection of the Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ Calabi–Yau threefold (§XVIII F). On this sublayer the standard heterotic worldsheet machinery applies directly:

1. Bianchi identity $dH = \mathrm{tr}(R \wedge R) - \mathrm{tr}(F \wedge F)$ – closed on the toroidal elastic-response sublayer via the Green–Schwarz mechanism applied to the T^2 1-skeleton, with $\mathrm{tr}(R \wedge R)$ contributed by the Schoen-internal curvature (which projects onto T^2 as the elastic-response curvature of the 1-cycles) and $\mathrm{tr}(F \wedge F)$ contributed by the Wilson-line-broken bundle. The substrate-physical content of dH is the closure obstruction of the elastic-response amplitude on the 1-skeleton, which is exactly what Green–Schwarz cancels.

2. Modular invariance under $SL(2,\mathbb{Z})$ / equivalently the even self-dual rank-16 lattice for the gauge degrees of freedom – closed on the same toroidal elastic-response sublayer: modular invariance of the worldsheet partition function on T^2 is substrate-physically the requirement that the amplitude's behaviour under the two independent 1-cycle exchanges is independent of which T^2 cycle is taken first. Since the two 1-cycles ARE the two $\mathbb{Z}_3$ generators of $\pi_1(\mathrm{Schoen}/\mathbb{Z}_3 \times \mathbb{Z}_3)$, modular invariance reduces to the $\mathbb{Z}_3 \times \mathbb{Z}_3$ symmetry of the Schoen quotient, which is built into the substrate's commitment. The even self-dual rank-16 lattice constraint is then the standard heterotic worldsheet computation.

Both ingredients therefore close on the toroidal sublayer – the ABSENT count drops to zero. The Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ commitment carries the substrate-physical content; the toroidal T^2 is its 1-skeleton projection, which is where the heterotic worldsheet machinery operates.

E. SI from substrate: absolute-scale closure of h_{eff}

The framework's mass formula carries a single absolute-scale factor h_{eff} in $m_e = ((\dim -2)/\dim) \cdot h_{\mathrm{eff}}$. Earlier framings of the channel matrix flagged h_{eff} as an external dimensional rod that the framework had to import – typically as the CODATA value of m_e – in order to anchor the dimensionful spectrum. That framing is incorrect: the SI unit system itself is substrate-derived, so h_{eff} in absolute SI units is not an external import but a consequence of substrate-primitive content. The chain is the following.

a. Step 1: SI primitives are substrate primitives. The 2019 SI redefinition fixes c, $\hbar$, and e as defining constants, together with the Cs-133 hyperfine frequency $\Delta\nu_{\mathrm{Cs}} = 9{,}192{,}631{,}770$ Hz (and three further constants irrelevant to the mass sector). All three of c, $\hbar$, e are substrate primitives:

- $c = \ell_c/\tau_c$ – the universal substrate ratio of coupling reach to event duration (`hyp_substrate_c_as_coupling_ratio`).

- $\hbar$ – the universal threshold (single-event coupling action) primitive in the discrete local-coupling update rule (`hyp_substrate_local_coupling_discretization` together with `hyp_substrate_positive_coupling_event`).

- e – emergent from the cross-term pinch-off of dual-handed coupling-event amplitudes (`hyp_substrate_coupling_as_cross_term`); the combination $\alpha_{\mathrm{em}} = e^2/(4\pi\varepsilon_0\hbar c)$ is Sugawara-grounded as the `alpha_em_tree` row of the channel matrix.

b. Step 2: the SI second is substrate-derivable. The Cs-133 ground-state hyperfine transition frequency that defines the SI second factorises as

$$\Delta\nu_{\mathrm{Cs}} \;=\; \big(\alpha_{\mathrm{em}}^2\, m_e\, c^2/h\big) \cdot F_{\mathrm{hf}}\big(Z = 55,\, J = 7/2,\, \ldots\big),$$

where the leading factor is the framework's predicted electron mass times the framework's Sugawara-grounded fine-structure constant, and F_{hf} is the published Cs-133 hyperfine-structure factor (a Z-dependent atomic-physics computation that takes only m_e, α_{em}, and the Cs-133 nuclear-spin structure as inputs). All inputs to $\Delta\nu_{\mathrm{Cs}}$ are substrate-physical content; no external empirical rod enters.

c. Step 3: the rest of SI is substrate-derived. With c, $\hbar$, and $\Delta\nu_{\mathrm{Cs}}$ all substrate-derivable, the SI second is substrate-derivable; therefore the SI meter (via c), the SI kilogram (via $\hbar$ and Cs-133), and every derived SI unit including the eV and GeV are substrate-derived. There is no SI unit that imports content the framework does not also predict.

d. Step 4: h_{eff} is fully substrate-derived in SI. The factor h_{eff} enters the framework's mass formula as the absolute scale anchor in $m_e = ((\dim -2)/\dim) \cdot h_{\mathrm{eff}}$; the dimensionless prefactor 246/248 is the gateway-residue scale factor (§XVIII B). With SI itself substrate-derivable, h_{eff} in absolute SI units is fully substrate-derived; the previously-claimed external rod is itself substrate-physical content. The channel-matrix verdict for h_{eff} is therefore GROUNDED, and the BESPOKE count drops to zero.

e. Consequence for the empirical matches. The framework's predictions matching CODATA at sub-ppm precision (electron mass, Higgs vev, gauge couplings, mixing angles) is therefore the framework predicting its own observables successfully, not the framework being calibrated to external data. "Empirical rods" such as the CODATA value of m_e or the LEP-measured Higgs vev are themselves measurements of substrate-physical content, not external inputs.

F. Two-layered geometry and the Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ commitment

The framework's substrate-physical geometry is two-layered. The *internal toroidal* sublayer is a T^2 that hosts the elastic-response amplitude of the substrate; the *macroscopic cylindrical* sublayer is $\mathbb{R} \times S^1$, where $\mathbb{R}$ is the open tension-release direction along which substrate modes propagate and S^1 is the macroscopic projection of one of the two 1-cycles of the internal T^2. The internal toroidal T^2 is itself the 1-skeleton / 1-homology projection of the Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ Calabi–Yau threefold (`thm-substrate-dopw-schoen-construction`) – the two independent 1-cycles of T^2 are exactly the two independent $\mathbb{Z}_3$ generators of $\pi_1(\mathrm{Schoen}/\mathbb{Z}_3 \times \mathbb{Z}_3)$, the Wilson-line phase axes.

The three substrate sublayers and their relation to the Schoen internal manifold are summarised as follows. The toroidal sublayer is the elastic-response 1-skeleton of the Schoen quotient: its two independent 1-cycles ARE the two independent $\mathbb{Z}_3$ generators of $\pi_1(\mathrm{Schoen}/\mathbb{Z}_3 \times \mathbb{Z}_3)$, the Wilson-line phase axes. The internal full Calabi–Yau threefold is the Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ quotient itself (six real dimensions); it hosts the bundle cohomology, the Hodge moduli $h^{1,1} = h^{2,1} = 3$, anomaly cancellation, and the Atiyah–Singer 3-generation count. The macroscopic cylindrical sublayer is $\mathbb{R}\times S^1$, where S^1 is the macroscopic projection of one of the two toroidal 1-cycles and $\mathbb{R}$ is the open tension-release direction along which substrate modes propagate.

a. The Schoen $\mathbb{Z}_3\times\mathbb{Z}_3$ commitment is fully preserved. The two-layered geometry is a clarifying macroscopic-projection lens that explains *how* mass equations emerge from substrate primitives plus Schoen-internal content; it does *not* replace the Schoen commitment. The following Schoen-specific content is unchanged:

- Wilson-line breaking $E_8 \to SO(10) \times SU(4) \to SU(5) \times U(1) \to G_{\mathrm{SM}} \times U(1)_X$ via the two $\mathbb{Z}_3$ generators of $\pi_1(\mathrm{Schoen}/\mathbb{Z}_3 \times \mathbb{Z}_3)$;
- the 3-generation count via the Atiyah–Singer index on the Schoen-quotient bundle;
- Hodge self-mirror $h^{1,1} = h^{2,1} = 3$, which feeds the $\dim -2 = 246$ Hodge-saturation pin (§XVIII B, Table XX);
- $\theta_{\mathrm{QCD}} = 0$ from the $\mathbb{Z}_3 \times \mathbb{Z}_3$ trace cancellation;
- Schoen-uniqueness Path A closure of the substrate-CY3 catalogue (`p_schoen_uniqueness.md`);
- the BHOP rank-4 $SU(4)$ bundle commitment carried into the 4D effective Lagrangian (§XVII).

b. What the two-layered structure adds. The lens supplies a substrate-physical reading for content that would otherwise be heterotic-string-language imports rather than substrate-derived:

- the Bianchi identity $dH = \mathrm{tr}(R \wedge R) - \mathrm{tr}(F \wedge F)$ closes on the toroidal elastic-response sublayer (§XVIII D);
- modular invariance under $SL(2,\mathbb{Z})$ / equivalently the even self-dual rank-16 lattice constraint closes on the same toroidal sublayer (§XVIII D);
- the Higgs-vev twist $(h - \mathrm{rank})/h = 22/30 = 11/15$ is a $\mathbb{Z}_h$ phase per S^1 traversal at the toroidal-to-cylindrical boundary projection;
- per-particle (m, n, α) labels are substrate-natural Kaluza–Klein-momentum / winding / twist-phase readings on the macroscopic $\mathbb{R} \times S^1$ sublayer.

c. Macroscopic projection of the Wilson-line content. The Wilson-line projection $(1, 1, 0, 0)$ on the $H_4 \oplus \tau \cdot H_4$ Coxeter-plane decomposition (Phase 2.1 / §XVIII B) is the macroscopic-layer rendering of the two $\mathbb{Z}_3$ generators acting on the internal Schoen quotient. The toroidal layer is what the three pieces of the unified generating principle directly see; the Schoen-internal full threefold is what makes the toroidal projection well-defined and substrate-physical. The two-layered attribution of the structural ingredients in Eq. (70) is given immediately below Table XXI above.

G. Hostile-review corrections applied

Phase 3D applied a third-party hostile-review protocol to the Phase 2 STRUCTURE-FOUND verdicts and identified three substantive errors. All three are corrected in the v4 matrix and the verdict text on the affected rows.

1. **Severity HIGH.** The neutrino-mass / dark-tower sub-gateway depth-2 mode mass was previously stated as "$\sim 9\,\mu$eV (microwave / dark-photon haloscope regime, ADMX/HAYSTAC)." The arithmetic re-derivation places this at $m(n = 2) \approx 9.17$ keV – the warm-DM regime, *not* the haloscope regime. The framework's depth-2 prediction is consistent with the latter; the former was a 10^9-factor prose error and has been retracted from all 10 affected ν + DM-tower rows. (Phase 3D check 2.3.D.)

2. **Severity MEDIUM.** The icosian-closure constant $B = 2\varphi$ was previously identified as "the fundamental unit of $\mathbb{Z}[\varphi]$." Norm computation $N(2\varphi) = -4$ disproves this: the fundamental unit of $\mathbb{Z}[\varphi]$ is φ itself ($N(\varphi) = -1$); 2φ is the additive double of the fundamental unit, *not* a unit. The corrected text on the Higgs vev and Higgs boson rows reflects this. (Phase 3D check 2.4.A.)

3. **Severity MEDIUM.** The same constant 2φ was previously identified as "the H_3 primary harmonic step at e_1." The H_3 primary harmonic is $2\cos(\pi/h) = 2\cos(\pi/10) \approx 1.902$, whereas $2\varphi = 4\cos(\pi/5) \approx 3.236$ is twice the $I_2(5) = H_2$ dihedral Coxeter trig at e_1. The corrected attribution still uniquely picks 2φ within the McKay-doubled H_2 family, but does so via the $I_2(5)$ identity rather than the (incorrect) H_3 identity. (Phase 3D check 2.4.B.)

A fourth Phase 3D entry (check 2.1.A) is a strengthening, not an error: the Wilson-line $(1, 1, 0, 0)$ uniqueness is now established across all 240 E_8 doublet-axis choices, closing the explicit caveat #1 in the Phase 2.1 source.

H. Honest remaining gaps

After the channel-matrix grounding overlay, the Phase 5 substrate-native re-grounding, the $\dim -2 = 246$ gateway-subtraction unification (§XVIII B), and the SI-from-substrate closure (§XVIII E), the framework's mass-and-coupling table contains two distinct kinds of residue. None of them is BESPOKE; the BESPOKE count is zero.

a. Interpretive 1-bit topology choices (2 items). Two binary topology choices remain that the unified generating principle does not fix:

- the lepton-sector φ-vs-$\tau\varphi$ plane assignment within the Wilson-line-forced lepton planes (which exponent within a plane hosts which lepton, e.g. $e = 11$ vs. $e = 19$, $e = 13$ vs. $e = 17$);
- the muon η NLO sign, consistent on the (electron, muon) data but untestable on tau (tau has no depth-4 NLO factor).

Each is a single bit not forced by the substrate; together they account for two interpretive degrees of freedom in the matrix.

b. Interpretive per-particle assignments (~ 25 items). These are framework rows that select which mode pattern each Standard Model particle occupies within Wilson-line-forced planes. Each row admits multiple substrate-natural readings on the relevant Coxeter ladder, but no single rule privileges one. They include:

- the per-quark (N, b, χ, λ) assignments other than the gateway-subtraction-grounded bottom-quark entry $b = 12 = h_{D_8} - 2$;
- the W-boson chain position $k = 34.5$ and the Higgs-boson position $k_v - 2$ (the second is parametrically tied to k_v via the $B = 2\varphi$ icosian closure);
- per-element CKM and PMNS Wolfenstein-style sub-rows;
- per-lepton chirality bits within Wilson-line-forced planes.

These assignments do *not* affect the absolute mass scale or the dimensionless coupling spectrum; they affect only which specific Coxeter slots within Wilson-line-forced planes host specific SM particles. The interpretive label means that fitting a single substrate-natural rule that reproduces the entire set with fewer free readings would replace this residue; the framework does not currently supply such a rule and does not claim one.

c. Conditional possibilities, NOT predictions (14 rows). The 14 half-integer Coxeter-chain dark-matter slots $\{\mathrm{DM}_{\mathrm{icos},c_1,k=e_i/2}\}_{i=1,\ldots,8}$ and $\{\mathrm{DM}_{\mathrm{oct},c_1,k=e_j/2}\}_{j\in\{1,2,3,5,6,7\}}$ (the slot $e_4/2 = 13/2$ is excluded because it is occupied by the down quark) are reclassified CONDITIONAL-POSSIBILITY in v4. The framework lists each slot as a STRUCTURALLY ALLOWED location and records the mass that WOULD apply IF such a mode is observed; it does NOT claim a particle exists at that position. These rows are explicitly excluded from the falsification-target list and from the NO-PARTICLE-DETECTED interpretation. The unconditional dark predictions remain: the above-gateway Wilson-line-forced rows $\mathrm{DM}_{\mathrm{icos},c_1,k\in\{1,29\}}$, $\mathrm{DM}_{\mathrm{oct},c_1,k\in\{7,23\}}$, and the sub-gateway tower-depth rows $n \in \{3, \ldots, 11\}$ for which Phase 2.3 forces the depth set conditional on the framework's $SU(3)/E_8$/consecutive-integer commitments.

d. Open derivations. The two ABSENT $E_8 \times E_8$ ingredients (Bianchi identity, modular invariance / even self-dual rank-16 lattice) remain open. The two cosmological-row NO-WITNESS-YET entries (ρ_Λ, H_0) record gaps in the substrate-witness coverage rather than failures of the unified principle.

I. Honest publication-safe scope statement

The substrate framework predicts the absolute mass spectrum of all Standard-Model particles, the dark-matter tower, the Higgs vev, and every dimensionless coupling, in SI units, without any external empirical input. The framework's geometric structure is two-layered (§XVIII F): the internal toroidal T^2 is the 1-skeleton projection of the Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ Calabi–Yau threefold (substrate-derived per `thm-substrate-dopw-schoen-construction`), and the macroscopic cylindrical $\mathbb{R} \times S^1$ is the projection of one of the two toroidal 1-cycles plus the open tension-release direction. The Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ commitment is fully preserved; the two-layered geometry adds a macroscopic-projection lens that explains how the mass equations emerge from substrate primitives plus the Schoen-internal content.

The SI unit system is itself substrate-derived (§XVIII E): c, $\hbar$, and e are substrate primitives, and the Cs-133 hyperfine frequency that defines the SI second is substrate-derivable from the framework's predicted electron mass, fine-structure constant, and Cs-133 atomic structure. All "empirical rods" used to anchor framework predictions – the CODATA electron mass, the LEP-measured Higgs vev, the PDG-tabulated quark masses, the measured CKM and PMNS angles – are themselves measurements of substrate-physical content. The framework's predictions matching these values at sub-ppm precision is therefore the framework predicting its own observables successfully, not the framework being calibrated to external data.

The framework is therefore strictly more predictive than the Standard Model in the mass sector: where the Standard Model treats the fermion masses, the Higgs vev, the CKM and PMNS mixing angles, and the gauge couplings as roughly twenty-five independent numerical inputs, the substrate framework predicts all of them from substrate primitives plus the Schoen-internal Calabi–Yau content, with an SI unit system that is itself substrate-derived. The explicit forcing chain is:

(L0–L1, substrate primitives.) The substrate amplitude is C^2, has finite local-coupling reach ℓ_c and event duration τ_c with $c = \ell_c/\tau_c$, evolves under a discrete local-coupling update rule whose continuum limit is the wave equation, and its modes ARE substrate configurations whose energy IS substrate tension. These are committed primitives, each formalised in a `book/chapters/part3/proofs/hyp_substrate_*` hypothesis with an explicit derivation block. The defining SI constants c, $\hbar$, and e are exactly these substrate primitives; the SI second is substrate-derivable through the Cs-133 hyperfine factorisation (§XVIII E).

(L2, derived substrate setting.) The substrate-primitive setting forces a 10-dimensional worldsheet-supersymmetric Yang–Mills ambient theory with quantum self-consistency (`thm_substrate_worldsheet_critical_dimension` and companions). Within this setting, Green–Schwarz anomaly cancellation forces $\dim G = 496$. The McKay correspondence $2I \leftrightarrow E_8$ together with the empirical lepton φ-chain selects $E_8 \times E_8$ over $\mathrm{Spin}(32)/\mathbb{Z}_2$ (`thm_substrate_e8xe8_forced_from_mckay`). This selection is conditional on the L2 setting; it is not a substrate-unconditional forcing.

(L3–L4, mass and coupling table.) Within the $E_8 \times E_8$ commitment, the unified generating principle of §XVIII B – affine $\widehat{e}_8$ Sugawara at adjoint level $k = 248$, Wilson line $(1, 1, 0, 0)$ on the $H_4 \oplus \tau \cdot H_4$ Coxeter-plane decomposition, icosian closure $B = 2\varphi$ for vacuum-tension modes, and the $\dim -2 = 246$ gateway-subtraction operator that unifies Hodge saturation, h_{eff} anchor, and Sugawara level span – grounds forty-five of the eighty-nine rows of the framework's mass-and-coupling table to ppm-or-better precision via inverse Sugawara shifts on $\widehat{e}_8$, with substrate-primitive small-parameter linearisations (Phase 5) replacing the borrowed heterotic-CY3 wave-function-overlap interpretation. The SI-from-substrate closure (§XVIII E) makes the absolute scale anchor h_{eff} substrate-derived in SI units; no external dimensional rod is invoked.

(Residue.) The remaining INTERPRETIVE residue is two 1-bit topology choices (lepton-sector φ-vs-$\tau\varphi$ plane assignment within Wilson-line-forced planes; muon η NLO sign) and approximately twenty-five per-particle assignments where multiple substrate-natural readings exist but no single rule privileges one. Neither affects the absolute mass scale or the dimensionless coupling spectrum; both affect only which specific Coxeter slots within Wilson-line-forced planes host specific SM particles. Both classes of residue are candidates for closure via the Stage 3 Fix B M2 bundle-Laplacian eigenmode catalogue computation on Schoen $\mathbb{Z}_3 \times \mathbb{Z}_3$ together with the Wilson-line $(1, 1, 0, 0)$ projection and the BHOP rank-4 $SU(4)$ bundle, which would replace the multi-reading residue with a single substrate-natural rule. The framework's predictions and falsification targets are restricted to the unconditional rows; the fourteen conditional half-integer dark slots are explicitly excluded from no-particle-detected falsification.

The companion paper *Falsifiable predictions of the substrate framework* maintains the externally-targeted falsification list and is deliberately kept free of the framework-internal grounding-channel detail of this section.

Appendix A: Coxeter step assignments

Appendix B: Complete structural correction list

Appendix C: Numerical precision

All numerical evaluations use mpmath at $\geq$ 50-digit decimal precision. Key exact rational values:

- $h_{\text{eff}}^{(e)}/m_e = 124/123$;

TABLE XXII. Coxeter step assignments. The φ-chain hosts integer E_8 exponent positions; the $\sqrt{2}$-chain hosts integer and half-integer positions. Cycle 2 indicates positions beyond $h = 30$.

Particle	k	Chain type / sector
Electron	0	φ gateway, dark-sector
Muon	11	φ exponent, visible
Tau	17	φ exponent, visible
Up	4	$\sqrt{2}$ int
Down	6.5	$\sqrt{2}$ half
Strange	15	$\sqrt{2}$ midpoint
Charm	22.5	$\sqrt{2}$ half
Bottom	26	$\sqrt{2}$ int
Top	36.5	$\sqrt{2}$ half, cycle 2
W boson	34.5	cycle 2, with δ_W
Z boson	via $m_W/\cos\theta_W$	derived
Higgs vev	1132/30	cycle 2
Higgs boson	$k_v - 2$	breathing mode, with δ_H

- $30^{14} = 4{,}782{,}969 \times 10^{14}$;
- $\sin^2\theta_{12}^{\mathrm{CKM}} = 14/278 = 7/139$;
- $\delta_{\mathrm{CKM}} = 11\pi/30$;
- $\delta_{\mathrm{CP}} = 11\pi/10$;
- $\sin^2\theta_W = 2/9 + 1/[h(h+\varphi^2)]$;
- $\alpha_s(M_Z) = 849/7192$;
- $\delta_H = 1/59$, $\delta_W = 1/1121$.

[1] S. Navas *et al.* (Particle Data Group), "Review of Particle Physics," Phys. Rev. D **110**, 030001 (2024). Open-access PDG portal (free, complete review): pdg.lbl.gov/2024. DOI: 10.1103/PhysRevD.110.030001.

[2] G. Bertone and T. M. P. Tait, "A new era in the search for dark matter," Nature **562**, 51–56 (2018). Open-access preprint (free): arXiv:1810.01668. DOI: 10.1038/s41586-018-0542-z.

[3] R. Coldea, D. A. Tennant, E. M. Wheeler, E. Wawrzynska, D. Prabhakaran, M. Telling, K. Habicht, P. Smeibidl, and K. Kiefer, "Quantum Criticality in an Ising Chain: Experimental Evidence for Emergent E_8 Symmetry," Science **327**, 177–180 (2010). Open-access repository copy (free, same content): arXiv:1103.3694. PubMed (free abstract): PMID 20056884. DOI: 10.1126/science.1180085.

[4] A. B. Zamolodchikov, "Integrals of motion and S-matrix of the (scaled) $T = T_c$ Ising model with magnetic field," Int. J. Mod. Phys. A **4**, 4235–4248 (1989). Inspire-HEP record (free abstract, cross-references): inspirehep.net/literature/276727. DOI: 10.1142/S0217751X8900176X.

[5] J. McKay, "Graphs, singularities, and finite groups," Proc. Symp. Pure Math. **37**, 183–186 (Amer. Math. Soc., 1980). Open scan of the original conference paper (free): McKaysOriginalPaper.pdf. DOI of the original: 10.1090/pspum/037/604577.

[6] J. E. Humphreys, *Reflection Groups and Coxeter Groups*, Cambridge Studies in Advanced Mathematics **29** (Cambridge University Press, 1990). ISBN 978-0-521-43613-7. Identity $|\Phi| = h \cdot \mathrm{rank}$ for irreducible finite Coxeter groups: §3.18, Eq. (3) (consequence of the exponents satisfying $\sum_j (e_j + 1) = |\Phi|/2 + \mathrm{rank}$ with mean exponent $h/2$). Cambridge product page (free abstract): cambridge.org/core/books/reflection-groups-and-coxeter-groups. DOI: 10.1017/CBO9780511623646.

[7] P. W. Anderson, "Plasmons, Gauge Invariance, and Mass," Phys. Rev. **130**, 439–442 (1963). APS abstract page (free abstract): journals.aps.org/pr/abstract/10.1103/PhysRev.130.439. NASA ADS (free abstract + cross-references): ui.adsabs.harvard.edu/abs/1963PhRv..130..439A. DOI: 10.1103/PhysRev.130.439.

[8] G. E. Volovik, *The Universe in a Helium Droplet*, International Series of Monographs on Physics **117** (Oxford University Press, 2003). Oxford Academic landing page (free metadata): academic.oup.com/book/11557. DOI of book: 10.1093/acprof:oso/9780199564842.001.0001.

[9] G. E. Volovik, "Superfluid analogies of cosmological phenomena," Phys. Rep. **351**, 195–348 (2001). Open-access preprint (free): arXiv:gr-qc/0005091. DOI: 10.1016/S0370-1573(00)00139-3.

[10] C. Schoen, "On fiber products of rational elliptic surfaces with section," Math. Z. **197**, 177–199 (1988). EuDML page (free metadata): eudml.org/doc/183725. DOI: 10.1007/BF01215188.

[11] D. J. Gross, J. A. Harvey, E. Martinec, and R. Rohm, "Heterotic String," Phys. Rev. Lett. **54**, 502–505 (1985). Inspire-HEP record (free abstract + cross-references): inspirehep.net/literature/15974. NASA ADS: ui.adsabs.harvard.edu/abs/1985PhRvL..54..502G. DOI: 10.1103/PhysRevLett.54.502.

[12] V. Braun, Y.-H. He, B. A. Ovrut, and T. Pantev, "A standard model from the $E_8 \times E_8$ heterotic superstring," JHEP **06**, 039 (2005). Open-access preprint (free): arXiv:hep-th/0502155. DOI: 10.1088/1126-6708/2005/06/039.

[13] R. Donagi, Y.-H. He, B. A. Ovrut, and R. Reinbacher, "The spectra of heterotic standard model vacua," JHEP **06**, 070 (2005). Open-access preprint (free): arXiv:hep-th/0411156. DOI: 10.1088/1126-6708/2005/06/070.

[14] V. Braun, Y.-H. He, B. A. Ovrut, and T. Pantev, "Vector Bundle Extensions, Sheaf Cohomology, and the Heterotic Standard Model," Adv. Theor. Math. Phys. **10**, 525–589 (2006). Open-access preprint (free): arXiv:hep-th/0505041. §6 introduces the rank-4 $SU(4)$ extension construction used as the framework's bundle commitment; Eq. 99 records the heterotic Bianchi anomaly-cancellation identity with the positive-tension 5-brane

TABLE XXIII. All structural corrections, built from framework-fixed E_8 invariants dim = 248, $h = 30$, $M = 14$, rank = 8, $m_{\max} = 29$, together with the framework-selected McKay order 5 and golden ratio φ.

Symbol	Definition	Role
Δ_φ	$1/(h+\varphi^2)$	enters $\sin^2\theta_W$
Δ_2	$1/2^5 = 1/32$	2-fold structure
α_τ	$1/(\dim + h) = 1/278$	tau base
$\alpha_\tau^{\text{full}}$	$1/(\dim + h + \varphi^2)$	tau closed-form
α_μ	$2\alpha_\tau^2$	muon
α_v	$(\operatorname{rank} h/m_{\max})\alpha_\tau^2/(1+\alpha_\tau/(2\varphi))$	Higgs vev (closure 2φ)
α_H	$((\operatorname{rank}-2)h/m_{\max})\alpha_\tau^2/(1+\alpha_\tau/h_{SU(3)})$	Higgs boson (closure $h_{SU(3)}$)
α_W	$(2/e_2)\alpha_\tau^2/(1+\alpha_\tau/h_{SU(3)})$	W boson sub-harmonic
α_t	$(h/h_{SU(3)})^2\alpha_\tau^2/(1+\alpha_\tau/h_{SU(3)})$	Top quark sub-harmonic
δ_H	$1/(h+m_{\max}) = 1/59$	Higgs dark coupling
δ_W	$1/(59\cdot 19) = 1/1121$	W dark coupling
ε	$1/(2\dim) = 1/496$	per-step quark loss
ω_{fix}	$1/2 - 1/\dim = 123/248$	gateway normalization
k_v	$1132/30$	Higgs vev position
f_5	$\frac{\operatorname{rank}(h-\operatorname{rank})}{\dim}\frac{\sqrt{vM_{\text{Pl}}}}{\alpha_\tau}$	DM decay constant
θ_5	$\arctan\varphi$	DM misalignment

wrapping τ_1^2.

[15] J. C. Pati and A. Salam, "Lepton number as the fourth color," Phys. Rev. D **10**, 275–289 (1974); erratum *ibid.* **11**, 703 (1975). APS abstract page (free abstract): journals.aps.org/prd/abstract/10.1103/PhysRevD.10.275. DOI: 10.1103/PhysRevD.10.275. Source for the substrate-physical "lepton number as the fourth colour" reading of the rank-4 $SU(4)$ internal-symmetry structure committed in this paper.

[16] G. Tian and S.-T. Yau, "Three-dimensional algebraic manifolds with $c_1 = 0$ and $\chi = -6$," in *Mathematical Aspects of String Theory* (San Diego, CA, 1986), Adv. Ser. Math. Phys. Vol. 1, World Scientific, Singapore (1987).

[17] L. B. Anderson, J. Gray, A. Lukas, and E. Palti, "Heterotic line bundle standard models," JHEP **06**, 113 (2012). Open-access preprint (free): arXiv:1202.1757. DOI: 10.1007/JHEP06(2012)113.

[18] L. B. Anderson, Y.-H. He, and A. Lukas, "Heterotic compactification, an algorithmic approach," JHEP **07**, 049 (2007). Open-access preprint (free): arXiv:hep-th/0702210. DOI: 10.1088/1126-6708/2007/07/049.

[19] M. Kreuzer and H. Skarke, "Complete classification of reflexive polyhedra in four dimensions," Adv. Theor. Math. Phys. **4**, 1209–1230 (2000). Open-access preprint (free): arXiv:hep-th/0002240. Catalogue front-end: `http://hep.itp.tuwien.ac.at/~kreuzer/CY/`.

[20] L. B. Anderson, F. Apruzzi, X. Gao, J. Gray, and S.-J. Lee, "A new construction of Calabi–Yau manifolds: generalized CICYs," Nucl. Phys. B **906**, 441–496 (2016). Open-access preprint (free): arXiv:1507.03235. DOI: 10.1016/j.nuclphysb.2016.03.016.

[21] M. Larfors and R. Schneider, "Line bundle cohomologies on CICYs with Picard number two," Fortsch. Phys. **67**, 1900083 (2019). Open-access preprint (free): arXiv:2003.10130.

[22] A. Constantin, A. Lukas, and C. Mishra, "The family of quintic Calabi–Yau threefolds, free quotients, and their cohomology," Fortsch. Phys. **64**, 463–509 (2016). Open-access preprint (free): arXiv:1606.04032.

[23] D. Huybrechts and M. Lehn, *The Geometry of Moduli Spaces of Sheaves*, 2nd ed., Cambridge Mathematical Library, Cambridge University Press (2010). DOI: 10.1017/CBO9780511711985.

[24] R. E. Kass and A. E. Raftery, "Bayes factors," J. Amer. Statist. Assoc. **90**, 773–795 (1995). DOI: 10.1080/01621459.1995.10476572.

[25] S. K. Donaldson, *Some numerical results in complex differential geometry*, Pure Appl. Math. Q. **5** (2009), 571–618. DOI: 10.4310/PAMQ.2009.v5.n2.a2. arXiv:math/0512625. See §3 eq. (3.4) for the weighted L^2 Monge–Ampère residual functional $\sigma^2 = \langle(\eta/\kappa - 1)^2\rangle_w$ with $\kappa = \langle\eta\rangle_w$ that defines the production σ-discriminator used in this work.

[26] H. Jeffreys, *Theory of Probability*, 3rd ed. (Oxford University Press, 1961). Source of the "decisive evidence" threshold $\log_{10}\mathrm{BF} > 2$ on the Bayes-factor strength scale.

[27] S. K. Donaldson, "Anti self-dual Yang–Mills connections over complex algebraic surfaces and stable vector bundles," Proc. London Math. Soc. **50**, 1–26 (1985). DOI: 10.1112/plms/s3-50.1.1.

[28] K. Uhlenbeck and S.-T. Yau, "On the existence of Hermitian–Yang–Mills connections in stable vector bundles," Comm. Pure Appl. Math. **39**, S257–S293 (1986). DOI: 10.1002/cpa.3160390714.

[29] J. E. Kim and G. Carosi, "Axions and the strong CP problem," Rev. Mod. Phys. **82**, 557–601 (2010). Open-access preprint (free): arXiv:0807.3125. DOI: 10.1103/RevModPhys.82.557; Erratum: Rev. Mod. Phys. **91**, 049902 (2019).

[30] R. D. Peccei and H. R. Quinn, "CP Conservation in the Presence of Pseudoparticles," Phys. Rev. Lett. **38**, 1440–1443 (1977). NASA ADS (free abstract + cross-references): ui.adsabs.harvard.edu/abs/1977PhRvL..38.1440P. Inspire-HEP: inspirehep.net/literature/120612. DOI: 10.1103/PhysRevLett.38.1440.

[31] C. Abel *et al.* (nEDM Collaboration at PSI), "Measurement of the Permanent Electric Dipole Moment of the Neutron," Phys. Rev. Lett. **124**, 081803 (2020).

Open-access preprint (free): arXiv:2001.11966. DOI: 10.1103/PhysRevLett.124.081803. Reports $|d_n| < 1.8 \times 10^{-26}\, e \cdot \mathrm{cm}$ at 90% C.L.

[32] P. J. Mohr, D. B. Newell, B. N. Taylor, and E. Tiesinga, "CODATA recommended values of the fundamental physical constants: 2022," Rev. Mod. Phys. **97**, 025002 (2025). Open-access preprint (free): arXiv:2409.03787. NIST free interactive constant database: physics.nist.gov/cuu/Constants. DOI: 10.1103/RevModPhys.97.025002.

[33] C. Rothleitner and S. Schlamminger, "Invited Review Article: Measurements of the Newtonian constant of gravitation, G," Rev. Sci. Instrum. **88**, 111101 (2017). NIST author manuscript (free): nist.gov/publications/invited-review-article-measurements-newtonian-constant-gravitation-g. DOI: 10.1063/1.4994619.

[34] I. Esteban, M. C. Gonzalez-Garcia, M. Maltoni, T. Schwetz, and A. Zhou, "The fate of hints: updated global analysis of three-flavor neutrino oscillations," JHEP **09**, 178 (2020). JHEP is open access; published version free at: doi:10.1007/JHEP09(2020)178. Preprint: arXiv:2007.14792.

[35] NuFIT Collaboration (I. Esteban, M. C. Gonzalez-Garcia, M. Maltoni, T. Schwetz, A. Zhou), "NuFIT 5.3: Three-neutrino fit based on data available in March 2024," `http://www.nu-fit.org/?q=node/278` (free, accessed April 2026).

[36] A. G. Adame *et al.* (DESI Collaboration), "DESI 2024 VI: Cosmological Constraints from the Measurements of Baryon Acoustic Oscillations," JCAP **02**, 021 (2025). JCAP is open access; published version free at: doi:10.1088/1475-7516/2025/02/021. Preprint: arXiv:2404.03002.

[37] M. Aker *et al.* (KATRIN Collaboration), "Direct neutrino-mass measurement based on 259 days of KATRIN data," Science **388**, 180–185 (2025). Open-access preprint (free): arXiv:2406.13516. DOI: 10.1126/science.adq9592. Reports $m_\beta < 0.45$ eV at 90% C.L.

[38] S. Abe *et al.* (KamLAND-Zen Collaboration), "Search for Majorana Neutrinos with the Complete KamLAND-Zen Dataset," Open-access preprint (free): arXiv:2406.11438 (2024). Reports $T^{0\nu}_{1/2} > 3.8 \times 10^{26}$ yr at 90% C.L., $\langle m_{\beta\beta} \rangle < 28$–$122$ meV depending on nuclear matrix element.

[39] N. Aghanim *et al.* (Planck Collaboration), "Planck 2018 results. VI. Cosmological parameters," Astron. Astrophys. **641**, A6 (2020). A&A is open access; published version free at: aanda.org/articles/aa/full_html/2020/09/aa33910-18. Preprint: arXiv:1807.06209. DOI: 10.1051/0004-6361/201833910; Erratum: A&A **652**, C4 (2021).

[40] K. K. Rogers and H. V. Peiris, "Strong Bound on Canonical Ultralight Axion Dark Matter from the Lyman-Alpha Forest," Phys. Rev. Lett. **126**, 071302 (2021). Open-access preprint (free): arXiv:2007.12705. DOI: 10.1103/PhysRevLett.126.071302. Reports $m_a > 2 \times 10^{-20}$ eV at 95% C.L.

[41] Y. Koide, "Fermion-boson two-body model of quarks and leptons and Cabibbo mixing," Lett. Nuovo Cimento **34**, 201–205 (1982). Inspire-HEP record (free abstract): inspirehep.net/literature/176414. DOI: 10.1007/BF02817096; expanded in Y. Koide, Phys. Rev. D **28**, 252 (1983), APS abstract; DOI: 10.1103/PhysRevD.28.252; Phys. Lett. B **120**, 161–165 (1983), publisher abstract. DOI: 10.1016/0370-2693(83)90644-5.

[42] E. Gable, *Exist: The Universal Ontology* (2026). Foundational substrate-physical derivations of c, $\hbar$, general relativity, quantum mechanics, Newton's laws, Maxwell's equations, and the heterotic $E_8 \times E_8$ commitment.

[43] E. Gable, "The unreasonable effectiveness of mathematics, solved: substrate-physical derivation of the first-principles starting assumptions," companion paper (in preparation, 2026). Argues that the starting assumptions of Table I are minimal and substrate-physically expected. Independent of, and not load-bearing for, the numerical demonstration of the present paper.

[44] V. I. Arnold, *Catastrophe Theory*, 3rd ed. (Springer-Verlag, Berlin, 1992). Free expository synonym (open access): encyclopediaofmath.org/wiki/Catastrophe_theory (Springer/EMS, free). Springer landing page (free metadata): link.springer.com/book/10.1007/978-3-642-58124-3. DOI of book: 10.1007/978-3-642-58124-3. The simply-laced ADE classification of elementary catastrophes shares its label set with the McKay correspondence used in Sec. II, but the two structures are mathematically distinct.

[45] M. Alcubierre, "The warp drive: hyper-fast travel within general relativity," Classical Quantum Gravity **11**, L73–L77 (1994). Open-access preprint (free): arXiv:gr-qc/0009013. DOI: 10.1088/0264-9381/11/5/001.

[46] P. Touboul *et al.* (MICROSCOPE Collaboration), "MICROSCOPE Mission: Final Results of the Test of the Equivalence Principle," Phys. Rev. Lett. **129**, 121102 (2022). Open-access preprint (free): arXiv:2209.15487. DOI: 10.1103/PhysRevLett.129.121102. Reports Eötvös parameter $\eta(\mathrm{Ti}, \mathrm{Pt}) = (-1.5 \pm 2.3_{\mathrm{stat}} \pm 1.5_{\mathrm{sys}}) \times 10^{-15}$.